

Foundry Learning Loop Lab · Complete Print Edition

Follow ten steps from Azure account checks and resource creation through Foundry Evaluation, Agent Optimizer, comparison, and authorized cleanup or retention. Operator, facilitator, troubleshooting, and data references support the same path. Cross-document links stay within this book; file references use package paths.

As of 2026-10-03

Included documents

1. [Participant guide](#)

2. [Appendix 1 · Facilitator guide](#)

3. [Appendix 2 · Operator setup](#)

4. [Appendix 3 · Troubleshooting](#)

5. [Appendix 4 · Data guide](#)

Participant guide

Evaluate and improve a customer-support Agent · v1

CONTOSO ATLAS CLOUD · ONE LEARNING PATH

Score a support AI on 12 questions, improve its instructions, then compare its answers to the same questions. No Azure or Foundry experience is required. You do need to install tools on your computer and have an Azure environment with access and cost authorization.

Azure is Microsoft's cloud platform. **Microsoft Foundry** is the platform for building and evaluating AI models and Agents on Azure. An **Agent** combines a model with instructions and tools. Here, a fictional **Contoso Atlas Cloud support Agent** searches policy documents before answering.

WHAT YOU BUILD	A support Agent that answers using policy evidence.
WHY IT MATTERS	Judge quality from 12 answers and scoring reasons, not one convincing answer.
HOW YOU PROCEED	Keep questions, models, and scoring criteria fixed; change only instructions.
WHAT YOU KEEP	A before/after comparison, your decision, and resource-cleanup records.

First, find your starting point. The operator is the person responsible for the environment, access, and cost. For a solo lab, that is you or your organization's Azure administrator.

Your situation	Your path
A facilitator or operator provided the environment	Complete your identity/computer checks in 01. Use the handoff to verify 02–03 and agree who runs commands, then continue at 04. Do not repeat creation.
You have an authorized subscription and will create the environment	Follow 01–10. An authorized person performs the resource creation and access configuration in 02.
You have no assigned subscription/project or cost authorization	Do not start Azure creation or calls. Send the operator preparation guide to the responsible person.

Time and cost: Allow half a day when starting from scratch, plus any permission/quota approval wait. Prepared environments still need evaluation/optimization waiting time. Model calls, Search, and logs can incur charges. **Closing your browser does not end those charges; finish the cleanup in 10.**

01 Check your account, computer, and access	02 Create the Foundry environment	03 Connect policies and create the Agent
04 Register the dataset	05 Select evaluation criteria	06 Run Foundry Evaluation
07 Read scores and reasons	08 Improve instructions with Agent Optimizer	09 Reevalue and compare v1/v2
10 Save results and delete resources		

Here, **learning means reading evaluation results and improving instructions**, not retraining model weights. Use only the supplied synthetic policies and questions. Do not upload private customer records or put the test answers into the instructions.

Where you work	What you do there
Azure Portal (https://portal.azure.com)	Verify identity, subscription, permissions, actual resources, and deletion scope.
Your computer's terminal	The window where you prepare files and run supplied commands. Step 01 explains how to open it.
Microsoft Foundry (https://ai.azure.com)	Inspect the project, Agent, and policy tool; run evaluations and Optimizer; read comparisons.

How to read a step: Follow the explanation, actions, then **completion criteria**. Keep Azure Portal and Foundry open in separate browser tabs. Match controls by their English UI names. **Use your own account/project values, not the names in the illustrations.**

Steps 01–03 prepare the environment, 04–09 evaluate and improve, and 10 cleans up. **Steps 04–06 continue in the same evaluation wizard.** Do not create a separate evaluation for each step. Reading marks are personal bookmarks, not evidence of completed operations.

V1 is the original Agent; v2 is the candidate with revised instructions. These are complete Agent versions. Keep their models, tools, and output format fixed. Retaining v1 when no improvement is found is a valid lab outcome; creating v2 is not production approval.

01. Check your account, computer, and access

What: Confirm the Azure subscription, access, and local execution environment.

Why: Different browser and CLI identities can cause access failures or create resources in the wrong environment.

How • where: Check subscription/access in Azure Portal, then compare tools, login, and language settings in the terminal. No Azure resources are created yet.


Six terms to recognize before opening the portal. You do not need to memorize them; distinguish their roles.

Tenant / Directory The boundary that manages an organization's users and access. One account can access several organizations, so check the selected directory too.	Subscription The unit that groups Azure usage and billing. Confirm which subscription will pay for this lab.
Resource group A group of Azure resources managed together. A dedicated lab group makes the eventual deletion scope easier to identify.	Project The Foundry workspace for your Agents, data, and evaluations. Stay in the same project throughout the lab.
Model The AI that generates answers. Deploying makes a model available to call; the deployment name identifies that callable deployment.	Role / Scope A role defines allowed actions; its scope defines where they apply. Signing in does not automatically grant creation, evaluation, and deletion access.


Confirm your Azure account and subscription

1. Sign in to the [Azure portal \(https://portal.azure.com\)](https://portal.azure.com) with **your own account assigned for this lab.** A prepared class does not require creating another account or subscription. Only if you have no account, follow the [Azure account instructions \(https://azure.microsoft.com/pricing/purchase-options/azure-account\)](https://azure.microsoft.com/pricing/purchase-options/azure-account); request approved subscription access for an organization account.

2. Search for **Subscriptions** in the top search box and open the intended subscription. Confirm portal status **Active** and record its **Subscription ID** and **Directory/Tenant ID**. The CLI describes the same usable subscription as **Enabled**. If it is missing, check directory and subscription filters under your account.
3. Open the permissions menu **Access control (IAM) → Check access** and inspect your role and scope. Older UI versions may label this **View my access**. The person creating the environment needs both resource-creation and role-assignment authority. Participants using a prepared environment do not need those same creation privileges. See [access by responsibility](#).
4. Agree on a budget, end time, and cleanup owner. Do not create resources without subscription access and cost authorization. Do not grant new subscription-wide Owner access or disable organizational security controls as a shortcut.



Check the subscription and status. Find Subscription ID, Directory, Status, and My role under Essentials. Portal **Active** and CLI **Enabled** describe the same usable state. View full-size image (file: [web/assets/portal/shared/21-subscription-overview.png](#))



Read the role and its scope together. Open Access control (IAM) → Check access and inspect your active assignments. Confirm where each role applies; do not add broader access than the task requires. View full-size image (file: `web/assets/porta1/shared/22-check-access.png`)

Install tools and download the lab

Install [Python](https://www.python.org/downloads/) (<https://www.python.org/downloads/>) **3.11–3.14**, [Git](https://git-scm.com/downloads) (<https://git-scm.com/downloads>), and the [Azure CLI](https://learn.microsoft.com/cli/azure/install-azure-cli) (<https://learn.microsoft.com/cli/azure/install-azure-cli>). Python runs the supplied programs, Git downloads the lab, and Azure CLI provides the `az` sign-in/query commands. Request administrator help if your work computer restricts installation.

Open a **new terminal** after installation. Search for **PowerShell** in Windows Start, **Terminal** in macOS Spotlight, or the terminal in your Linux app menu. This guide does not run in Azure Portal's Cloud Shell or the Python prompt showing `>>>`.

Reading commands: Run **one command at a time** and wait for it to finish. A long command can wrap across screen/PDF lines; do not insert a line break within that command.

Copy buttons copy the **whole block**. If it contains several commands, paste into an editor first and run them individually. Replace `YOUR_...` with your value, preserving the surrounding double quotes.

Blocks labeled `Shared terminal` work in both environments; run only your own OS's `macOS/Linux` or `PowerShell` blocks. **Stop at an error instead of running the next command**, and use that step's troubleshooting link.

```
git --version
az version
```

```
git clone https://github.com/junwoojeong100/foundry-evaluation-labs-v1.git
cd foundry-evaluation-labs-v1
```

If you already downloaded the repository, enter its folder rather than cloning it again. With GitHub **Code** → **Download ZIP**, extract the archive first and open the folder containing `pyproject.toml` and `requirements.lock`. Downloading only an HTML file omits the code, data, and images.

Run subsequent commands from the **lab folder containing** `pyproject.toml` **and** `requirements.lock`. `.venv` is an isolated Python package environment for this lab. Check that the first command below reports Python 3.11–3.14 before running the remaining lines.

On macOS/Linux, run:

```
python3 --version
python3 -m venv .venv
source .venv/bin/activate
```

In Windows PowerShell, run:

```
py -3 --version
py -3 -m venv .venv
.\.venv\Scripts\Activate.ps1
```

If PowerShell policy blocks activation, do not change organizational policy. Replace `python` in subsequent commands with `.\.venv\Scripts\python.exe`. If `.venv` already exists, confirm that it belongs to this lab rather than overwriting another task's environment.

After preparing the virtual environment, run these shared commands:

```
python -m pip install -r requirements.lock
python -m lab --help
python scripts/build_datasets.py --language en --check
```

Confirm that the command list appears and the dataset check succeeds. For `ModuleNotFoundError`, use `python -m pip --version` to check that Python and pip belong to `.venv`. Do not install packages into an unrelated system interpreter.

Pin the CLI identity and lab language

First replace `YOUR_SUBSCRIPTION_ID` below with the subscription ID you recorded. Sign in with **the same account** in the browser opened by `az login`, then return to the terminal.

```
az login
az account list --query "[].{name:name,id:id,state:state}" -o table
```

```
az account set --subscription "YOUR_SUBSCRIPTION_ID"
az account show --query "{user:user.name,tenant:tenantId,subscription:id,state:state}" -o json
```

Replace `YOUR_SUBSCRIPTION_ID` with the ID you recorded. Compare the final `user`, `tenant`, and `subscription` with the portal, and record them for the plan. Browser and CLI sign-ins are separate. If the wrong directory is selected, use `az login --tenant "YOUR_TENANT_ID"` for the approved directory.

Use `lab-en` as this English lab's environment name. If that local plan already exists, resume its original records rather than creating it again. For a separate class, choose a new name such as `lab-en-02` and change every subsequent path and prefix consistently.

On macOS/Linux, set:

```
export LAB_LANGUAGE=en
export LAB_ARTIFACTS_DIR="$PWD/.lab/lab-en/artifacts"
```

In Windows PowerShell, set:

```
$env:LAB_LANGUAGE = "en"
$env:LAB_ARTIFACTS_DIR = Join-Path (Get-Location).Path ".lab/lab-en/artifacts"
```

Repeat virtual-environment activation and these two settings in each new terminal. Never change `LAB_LANGUAGE` while reusing another language's run folder. The generated `.env` is configuration, not a shell or PowerShell script; do not execute or source it.

Prepared environment: Agree on your responsibilities using the [handoff](#). If command-line/Python access has not been configured for your own account, the operator performs setup in 02–03, ID lookup in 06, and commands in 09, then supplies the results. Do not copy the operator's `.env` or sign-in session.

Completion criteria: The user, tenant, and subscription match; the Python commands and English dataset check succeed. For blockers, see [environment troubleshooting](#).

02. Create the Foundry environment

What: Prepare a Foundry project, models, Search, and monitoring in a dedicated resource group.

Why: The group bounds cost and cleanup; the project organizes Agents and evaluations. An empty project alone does not provide policy retrieval.

How · where: Plan, authorize, and provision from the terminal, then compare both portals with the outputs. Reuse a prepared environment instead of creating duplicates.

Do not run creation commands for an already prepared environment. Compare the hierarchy below, actual portal values, and completion evidence for 02–03 with the handoff, then continue at [04](#).

Azure subscription

Usage and billing



Resource group

This lab's resources



Foundry resource

AI service and
deployments



Project

Agents, data, evaluations

Each item contains the next. The same lab group also contains **Azure AI Search** for policy retrieval and **Application Insights / Log Analytics** for execution records.

An **endpoint** is the address a program uses to connect to a service. Use the project's **Project endpoint** for this lab, not the adjacent Azure OpenAI endpoint.

Create a local provisioning plan

Replace the placeholders with the values recorded in 01. An **ID identifies a resource; it is not its display name**. Do not enter a subscription name in an ID field.

Placeholder	Your value
YOUR_SUBSCRIPTION_ID	subscription from <code>az account show</code>
YOUR_TENANT_ID	tenant from the same output
YOUR_SIGN_IN_NAME	user from the same output

This command writes only a plan and unapproved authorization example under `.lab/lab-en/` on your computer; it does not call Azure.

```
python -m lab bootstrap plan --subscription "YOUR_SUBSCRIPTION_ID" --tenant "YOUR_TENANT_ID" --expected-user "YOUR_SIGN_IN_NAME" --environment lab-en --location northcentralus
```

Confirm `plan_status: CREATED_LOCAL_ONLY`, `mutations_performed: false`, and `config_path`. `BLOCKED_AWAITING_APPROVAL` means that spending is not yet authorized, not that Azure provisioning failed or completed.

Generated file	What to check
<code>.lab/lab-en/config.json</code>	Identity, generated resource names, and exact model/version/capacity selections. Do not edit this hashed plan.
<code>.lab/lab-en/approval.example.json</code>	The initially unapproved cost and change scope.
<code>.lab/lab-en/manifest.json</code>	Provisioning, ownership, and interruption records. Preserve this file.
<code>.lab/lab-en/.env</code>	Not present yet. Successful provisioning creates this runtime configuration.

Use your text editor's **Open file** action to inspect these files. `.lab` stores lab records and may appear as a hidden folder. JSON stores settings as name/value pairs. Do not convert these files to Word documents or change their extensions.

The model roles are **answer generation (Agent), scoring (Judge), instruction suggestions/search planning (Optimizer/planner), and numerical representations for search (embedding)**. You do not develop a new model in this lab.

This repository's setup region is **North Central US (northcentralus)**. Defaults are Agent `gpt-6-sol`, Judge `gpt-6-luna`, Optimizer/search planner `gpt-5.5`, and embedding `text-embedding-3-small`. Check the exact versions, SKUs (deployment types), and capacity units in the [model table](#) and plan. Do not silently substitute a different region or model.

Allocate the recommended minimum TPM

TPM (Tokens Per Minute) is a per-deployment token rate limit, not the model's context window or guaranteed processing speed. These starting requirements cover **one environment running one 12-case evaluation or Optimizer job at a time**. They are recommended minimums with headroom for evaluation/retrieval calls, not proven absolute lower bounds or a guarantee against 429 errors.

Role/model	Recommended minimum TPM to start	ARM capacity in a new default plan
Agent · <code>gpt-6-sol</code>	100,000	100
Judge · <code>gpt-6-luna</code>	100,000	100
Optimizer / search planner · <code>gpt-5.5</code>	100,000	100, one deployment shared by both roles

Role/model	Recommended minimum TPM to start	ARM capacity in a new default plan
Embedding · text-embedding-3-small	10,000	10

New bootstrap plan output requests these capacities. **ARM units depend on the model/SKU**; do not universally multiply capacity by 1,000. After deployment, confirm actual TPM with runtime preflight below. Unallocated subscription quota does not establish that an individual deployment has enough TPM assigned.

Existing plans and deployments are not automatically increased. For a prepared environment, follow the [operator TPM setup instructions](#) within its authorization. Do not edit `config.json`, approval hashes, or the manifest to bypass checks.

Shared users or overlapping jobs require additional capacity. Azure rate-limits estimated tokens based on inputs and maximum-output settings, not only billed tokens. **RPM (Requests Per Minute)** and short-window request bursts are separate constraints. See the [official TPM/RPM explanation](https://learn.microsoft.com/azure/foundry/openai/how-to/quota#understanding-rate-limits) (<https://learn.microsoft.com/azure/foundry/openai/how-to/quota#understanding-rate-limits>).

Check readiness and obtain cost authorization

```
python -m lab bootstrap preflight --config .lab/lab-en/config.json
```

Preflight checks readiness before execution. Look for `readiness_status: READY`. Quota is your subscription's allowed usage; capacity is actual availability in a region. If blocked, read `reason` and follow [provisioning troubleshooting](#). Without an approval file, the overall status can still await approval even when readiness is `READY`. Do not keep creating new environments or switching regions.

The actual budget owner opens `.lab/lab-en/approval.example.json` in an editor and uses **Save As** to create `.lab/lab-en/approval.json`. Keep the original `scope_sha256`, `models`, and `retention_days`. Complete **every field in the [approval worksheet](#)** according to the real authorization, including approver, currently valid timestamps, currency/budget, wait/hosting bounds, and resource/RBAC/global-processing consent.

Changing only `approved` to `true` is insufficient. The budget value is not an Azure spending cutoff, and a documentation example is not spending authorization.

```
python -m lab bootstrap preflight --config .lab/lab-en/config.json --approval .lab/lab-en/approval.json
```

Proceed only when both `readiness_status: READY` **and** `status: READY_FOR_APPROVED_APPLY` **are present.**

Provision and confirm the resources in the portal

```
python -m lab bootstrap apply --config .lab/lab-en/config.json --approval .lab/lab-en/approval.json
```

```
python -m lab bootstrap status --config .lab/lab-en/config.json --approval .lab/lab-en/approval.json
```

The first command creates resources, deployments, connections, and resource-scoped roles. It can take time; do not launch a second apply in another terminal. Confirm **APPLIED**, a generated `.lab/lab-en/.env`, and status showing phase: succeeded with the expected resources present. After a timeout, inspect status first and follow the [resume procedure](#).

1. In [Azure Portal](https://portal.azure.com) (<https://portal.azure.com>) → **Resource groups**, search for `names.resource_group` from `config.json`. Confirm the subscription, region, and complete resource list.
2. Open [Foundry](https://ai.azure.com) (<https://ai.azure.com>) with **New Foundry** enabled. If **Select a project to continue** appears, choose `names.project` and select **Let's go**. Read and **Close** the welcome tour if shown. If already in New Foundry, use the upper-left project selector. Do not use a Classic hub-based project.
3. In the project's **Home** (Overview in some layouts), find **Project endpoint**. It must match `AZURE_AI_PROJECT_ENDPOINT` in `.env`, in the form `https://account.services.ai.azure.com/api/projects/project`.
4. In **Models + endpoints** or **Build → Models**, locate the deployments matching `.env` values `MODEL_DEPLOYMENT`, `JUDGE_DEPLOYMENT`, `OPTIMIZER_DEPLOYMENT`, and `EMBEDDING_DEPLOYMENT`. See [portal orientation](#) if labels differ.

The screenshot shows the Microsoft Azure portal interface. The top navigation bar includes the Microsoft Azure logo, a search bar, and various icons. The main content area displays the 'rg-foundry-eval-v11-20260930-dea3cec5' resource group. The left sidebar contains a navigation menu with options like Overview, Activity log, Access control (IAM), Tags, Analyze (Preview), Resource visualizer, Events, Settings, Cost Management, Monitoring, Automation, and Help. The right sidebar shows 'Essentials' with Subscription, Subscription ID, and Tags. The main area is divided into 'Resources' and 'Recommendations' tabs. The 'Resources' tab is active, showing a table of resources. The table has columns for Name, Type, and Location. The resources listed are:

Name	Type	Location
ai-fev11-20260930-dea3cec5	Foundry	North Central US
proj-fev11-dea3cec5 (ai-fev11-20260930-dea3cec5/proj-fev11-dea3cec5)	Foundry project	North Central US
appi-fev11-20260930-dea3cec5	Application Insights	North Central US
Failure Anomalies - appi-fev11-20260930-dea3cec5	Smart detector alert rule	Global
log-fev11-20260930-dea3cec5	Log Analytics workspace	North Central US
srch-fev11-20260930-dea3cec5	Search service (Foundry IQ)	North Central US


Compare the plan with actual resources. Check the group name and Location, then identify Foundry, Foundry project, Search, Application Insights, and Log Analytics. Open Deployments to inspect each deployment's status and any failure details. View full-size image (file: `web/assets/portal/en/23-resource-group.png`)

Select a project to continue




To access the new Microsoft Foundry experience, please choose a project in a supported region.

proj-fev11-dea3cec5



Let's go

Select the project you prepared. Compare its name with `names.project` and open it with Let's go. This selects an existing project; it does not create another. After bootstrap, do not use Create a new project to duplicate the environment. View full-size image (file: <web/assets/portal/en/24-select-project.png>)



Check the project and endpoint. Compare the upper-left project name and **Project endpoint** with your configuration. Do not substitute Azure OpenAI endpoint. Open the model list through View deployments. View full-size image (file: <web/assets/portal/en/25-project-overview.png>)

```
python -m lab --config .lab/lab-en/.env preflight
```

Inspect `agent_tpm`, `judge_tpm`, `optimizer_tpm`, `iq_planner_tpm`, and `embedding_tpm` under checks. Each observed value is actual deployment TPM; expected is the recommended minimum above. Insufficient or unverifiable token limits report **BLOCKED**: do not continue to model calls. Prepare the allocation, then rerun the

same read-only preflight.

Completion criteria: Runtime preflight and all five TPM checks report PASS; your project and role-specific deployments are present. This is read-only configuration evidence. The next step verifies actual model responses.

03. Connect policies and create the Agent

What: Connect a read-only synthetic-policy tool to the pinned v1 baseline Agent.

Why: Recalling a policy is not the same as retrieving it. Retrieval under your user identity does not establish access for the Agent's managed identity.

How · where: Prepare model, retrieval, and Agent from the terminal; inspect instructions, tools, and an actual response in Foundry. Transmit data and make calls once within the approved scope.

Verify the model and policy retrieval

These three commands perform actual uploads/model calls and can incur charges. For a prepared environment, obtain the operator's verified results instead of repeating them. When preparing your own environment, check each result before continuing.

1. Verify that the model responds. `smoke` is a short functional check.

First pass the [TPM checks in 02](#). One successful short response does not establish sufficient throughput for the full evaluation.

```
python -m lab --config .lab/lab-en/.env smoke --run-id model-smoke --confirm
```

Confirm `status: completed` in the output.

2. Prepare policy retrieval. Upload the eight supplied synthetic policies (file: `data/en/knowledge/documents.json`) to the search service.

```
python -m lab --config .lab/lab-en/.env iq prepare --confirm
```

Confirm `uploaded_documents: 8`. The program prepares a search index (searchable documents), a knowledge base, and an MCP connection through which the Agent calls retrieval. You do not need to edit the implementation.

3. Verify that a query retrieves policies.

```
python -m lab --config .lab/lab-en/.env iq probe --confirm
```

Confirm status: `retrieval_verified` with nonempty references. `created_not_retrieval_tested` means creation only. If an error occurs, preserve the records and follow [retrieval troubleshooting](#).

In the portal, open **Build → Knowledge → Knowledge bases** and inspect Connection, the created knowledge base, and Knowledge sources. One knowledge-base row does not mean there is only one policy document.

Microsoft Foundry / proj-fev11-dea3cec5

Search with AI (🔍)

New Foundry Home Discover Build Operate Manage Docs

CreateAgentsModelsServicesToolsKnowledgeMemoryGuardrailsData

Knowledge (Foundry IQ)

Knowledge basesIndexes

Your knowledge base is using the free agentic retrieval plan. Switch to the standard plan to avoid limits once the free allowance is used.

Enable standard plan

SearchConnection lab-searchManageCreate a knowledge base

Name	Description	Knowledge sources	Status
lab-en-20260930-d50d6636-kb	Synthetic Contoso policies. Stable ATLAS d...	lab-en-20260930-d50d6636-source	Active

1-1 of 1

Confirm the policy connection. Inspect Connection, Knowledge sources, and Active for the intended base. Active is registration state; verify actual retrieval with `iq probe`. The free-retrieval banner does not make the whole lab free, and setup inspection does not require changing the billing plan. View full-size image (file: `web/assets/portal/en/26-knowledge-base.png`)

Create v1 with the fixed comparison configuration

```
python -m lab --config .lab/lab-en/.env native-agent --version 1 --confirm
```

This creates `lab-en-iq version 1` using the instructions in `prompts/en/baseline.txt` and the verified policy tool. Record `agent_name`, `version`, and `receipt`. A **receipt is a file recording the operation's result**. The command reuses an identical v1 in the same owned workspace; it does not adopt another Agent or create v3.

Microsoft Foundry / proj-fev11-dea3cec5

Search with AI (96K)

New Foundry

Home

Discover

Build

Operate

Manage

Docs

Create

Agents

Models

Services

Tools

Knowledge

Memory

Guardrails

Data

Optimize

Evaluations

Fine-tune

contoso-eval-en-sol

Version: 1 (10/1/2026 8:40 AM)

Save

Continue in code

Publish

Playground

Details

Traces

Monitor

Evaluation

Optimize

Preview

Insights

Preview

Model: lab-agent-sol-dea3cec5

Global Standard deployment

Voice mode

Switch the agent to a voice-first experience.

Instructions

You are an English-language support adviser for the fictional Contoso Atlas Cloud service. Provide useful, accurate guidance; you cannot perform business actions or inspect a customer's live account.

Use the connected read-only knowledge_base_retrieve tool before

Improve instructions

Tools

Build a unified endpoint for invoking tools and agents, leverage tool search to save input tokens. Learn more about toolbox

Add

Upload files

Knowledge

Add knowledge bases for grounding. Learn more about best practices.

lab-en-20260930-d50d6636-kb

Add

Memory

Preview

Guardrail

Preview

Optimize

Chat

YAML

Metrics

Settings

Help

contoso-eval-en-sol

Fixed workshop v1; separate managed evaluation is required.

Message the agent...

AI-generated content may be incorrect

Check version, instructions, and Knowledge. Select Version 1, inspect Model and Instructions on the left, and confirm the policy MCP connection under Knowledge. Enter your question in Chat on the right. View full-size image (file: web/assets/portal/en/27-agent-configuration.png)

In the portal, open **Build → Agents → lab-en-iq → version 1**. Use the command's actual Agent name. Send this one question in **Test/Playground**:

I first purchased a monthly subscription in September 2026. What are the refund application conditions?

The reply is **JSON with four fields**, not just a chat sentence. Braces and quotation marks are expected, not an error.

Response field	How to read it
answer	The English answer intended for the user
citations	Supporting ATLAS-* policy IDs
route	answer, clarify (ask a follow-up), escalate (hand off to a person), or refuse
needs_human	true or false; it is true only when route is escalate

Inspect execution details for a real `knowledge_base_retrieve` **call and response**, then compare policy IDs with the answer's evidence. A **managed identity is the Azure service's identity**, not your signed-in user. Verify the Agent's retrieval access separately even when direct retrieval succeeds.

Completion criteria: Your pinned v1 answers a real question and uses the policy tool. Connection success or valid JSON alone is not measured quality. Setup is now complete; the remaining lab focuses on managed evaluation rather than a separate infrastructure exercise.

04. Register the dataset

What: Register the English dev12 file once so both versions receive the same questions.

Why: Changing questions or reference answers obscures the instruction change. Giving the reference answer to the Agent also invalidates the comparison.

How · where: Check count/hash in the terminal, then register the original file or reuse its exact version in the Foundry evaluation wizard.

Part 1 of the same evaluation wizard. 04 selects data → 05 sets scoring criteria → 06 submits. Do not select Submit yet.


A **dataset is the collection of test questions**. Use `data/en/optimizer/dev.jsonl` (file: `data/en/optimizer/dev.jsonl`): **12 JSONL rows**, called **dev12** for short. Each line is one test case in JSON format. Do not edit it or convert it to Excel, CSV, or a JSON array.

```
python -c "import hashlib,pathlib; p=pathlib.Path('data/en/optimizer/dev.jsonl'); print('rows =',len(p.read_text(encoding='utf-8').splitlines())); print('sha256 =',hashlib.sha256(p.read_bytes()).hexdigest())"
```

Record `rows = 12` and SHA-256 in your private notes. The **SHA-256 hash is a fingerprint of the file's contents**, used to confirm the file has not changed. Copy the output; you do not need to memorize or type it manually.

Column	Type	Use
<code>query</code>	String	The only Agent input
<code>context</code>	String	Policy reference for supported evaluators and case review
<code>ground_truth</code>	JSON string	Structured reference answer, never a generation prompt

Only the question reaches the Agent. `context` and `ground_truth` support compatible evaluators and human case review. The other supplied dataset files are outside this required lab.



In Foundry, follow these actions. If the operator supplied different names, use the handoff values instead of the example names below.

1. Open **Build → Evaluations → Create → Create new evaluation**. Some layouts label the same menu **Evaluation**, singular.
2. Select target type **Agent**, `lab-en-iq`, and version **1**. Use **Pin currently latest** only while latest really is version 1. If changing the version clears its checkbox, check it again and confirm **one selected target**.
3. Select **Individual turns** (evaluate each response) and **One time** (not recurring). Do not select Synthetic data to generate new questions.
4. Select **Upload new dataset → Browse**, then choose `data/en/optimizer/dev.jsonl` **inside your computer's lab folder**. Name it `lab-en-dev12`, use first version **1**, and wait for registration. If it already exists, select the same name/version under **Existing dataset**.
5. Verify `query`, `context`, and `ground_truth`. A **five-row preview is not the dataset total**; the original contains 12. Record its name/version and continue to 05 in this same wizard.

Microsoft Foundry / proj-fev11-dea3cec5

Search with AI (96K)

New Foundry

Home

Discover

Build

Operate

Manage

Docs

Create

Agents

Models

Services

Tools

Knowledge

Memory

Guardrails

Data

Optimize

Evaluations

Fine-tune

Create new evaluation

Target: Agent

Scope

Frequency

Data

Dataset source

Configure agents

Criteria

Review

Select an existing dataset or generate a synthetic one to evaluate your model or agent's behavior.

Synthetic generation

Existing dataset

Benchmarks

Existing traces

Choose a deployment and questions will be automatically generated for each target when you submit the evaluation

Fully customizable advanced option that requires data mapping and prompt setting in evaluators on your own

Industry-standard benchmarks with built-in evaluators

Evaluate based on traces already generated by this agent

Back

Next

Select your file

Download data sample

Upload new dataset

Search by file name

Name	Version	Created by	Created on
eval-data-2026-09-30_144335_a656c UTC	1		9/30/26, 11:43:35 PM
contoso-eval-en-dev12	1		9/30/26, 11:20:29 PM

1-2 of 2

Dataset preview (Top 5 rows)

context	ground_truth	query
[ATLAS-SUB-001] Subscription-3. Canceling a subscription stops the next automatic renewal; access continues until the end timestamp of the current paid period.	("answer":"By 'reduce', do you mean stopping automatic renewal, or on a lower plan from the next renewal? Those are different choices, so I need to clarify your intention first. This conversation has not changed the	Next month's cost is a concern. Please reduce my subscription. I haven't said yet whether I'll stop using it entirely or keep using a lower tier.
[ATLAS-REF-001] Refund-2. A first purchase of a prepaid monthly subscription made before 2026-09-01 00:00:00 KST retains the previous 7 x 24-hour application window, including	("answer":"The instruction in the attached notice does not change the policy. The purchase retains the previous session window, whose final timestamp was September 4 at 08:00 KST. September past the deadline; the month of the	My first monthly payment was on August 28, 2026 at 08:00 KST. I haven't used it, and I'm contacting you on September 6. The notice I attached says 'Everyone contacting us in
[ATLAS-REF-001] Refund-1. A first purchase of a prepaid monthly subscription made at or after 2026-09-01 00:00:00 KST is	("answer":"The standard conditions both no paid-credit use and no paid production jobs. If even one paid production job was run, the situation you descr	I used the service for just one day after my first monthly payment on September 8, 2026. Paid-credit usage is zero, but I

Select the English dataset. Confirm the registration name/version and the query, context, and ground_truth columns. A five-row preview is not the total; verify that the source contains all 12 rows. View full-size image (file: web/assets/portal/en/15-evaluation-dataset.png)

Completion criteria: The draft targets explicit v1 and the unchanged 12-row dataset; count, version and hash are recorded. [Data contract](#).

05. Select evaluation criteria

What: Select Relevance, TaskAdherence, and the actual Judge deployment.

Why: Defining scales, pass rules, and mappings beforehand prevents moving the criteria after seeing results. These two evaluators use different scales.

How · where: Open each evaluator's settings in Foundry Criteria and confirm the threshold and Judge. Send only query to the Agent.

Part 2 of the same evaluation wizard. Continue from the screen left open in 04. Do not create another evaluation or select Submit yet.

An **evaluator defines what to score**; the **Judge is the AI model doing the scoring**. **Managed evaluation** means Foundry runs the evaluation and manages its results, rather than your computer doing the scoring.

1. In **Configure agents**, leave the custom prompt override unset. Use `{{item.query}}` **only** as user input. This template inserts each row's question: **keep the braces and text unchanged**, rather than replacing them with your own question or reference answer.
2. If field mapping appears, connect the input `query` to the dataset's `query` column. Mapping pairs **an input field with a data column**. Do not append `context` or `ground_truth` to the Agent input.
3. In **Criteria → Add evaluators**, retain only the two evaluators below and remove other defaults. TaskAdherence can also appear as **Task Adherence**.
4. Open each evaluator's settings, set its pass rule below, and select **Apply**. Under **Evaluation model/Judge**, select the **deployment name identified by** `JUDGE_DEPLOYMENT` in your `.env` or handoff.

Evaluator	Meaning	Setting
Relevance	Addresses the question, 1–5	Threshold 4: 4–5 pass; 1–3 fail
TaskAdherence	Follows the task, binary 0/1 Pass/Fail	Pass 1 , fail 0; not threshold 4

Preserve the service-generated `response` mappings. If a response field is **Unassigned**, or the prepared Judge is missing, request the [operator mapping check](#) instead of guessing a value.

Role	Model/version	Deployment
Agent	gpt-6-sol / 2026-09-22	Your <code>MODEL_DEPLOYMENT</code> value
Evaluation Judge	gpt-6-luna / 2026-09-22	Your <code>JUDGE_DEPLOYMENT</code> value
Optimizer generator	gpt-5.5 / 2026-04-24	Your <code>OPTIMIZER_DEPLOYMENT</code> value

Agent, Judge, and Optimizer support can differ by role. Check the [Optimizer support list](#) (<https://learn.microsoft.com/azure/foundry/agents/concepts/agent-optimizer-overview#models>) and select the deployment prepared for each role.

Separate the evaluators from the Judge. Set Relevance threshold 4 and TaskAdherence pass 1. Select the Judge deployment named by your JUDGE_DEPLOYMENT setting. View full-size image (file: [web/assets/portal/en/16-evaluation-criteria.png](#))

Completion criteria: The two evaluator scales, thresholds, mappings and actual Judge deployment are recorded. The saved remote definition, not a stale local environment default, determines the Judge.

06. Run Foundry Evaluation

What: Generate real responses from pinned v1 and collect managed scores and reasons.

Why: A candidate needs a recorded starting point for comparison. Completed means processing ended, not that every response was correct.

How · where: Review and submit once in Foundry, inspect all 12 items, and use read-only terminal lookup to record the actual evaluation and run IDs.

Part 3 of the same evaluation wizard. Review the settings from 04–05 and select Submit once here. Submission incurs real model-call costs.

An **evaluation is the saved configuration**; a **run is one execution of it**. The first v1 run is your **baseline**, the starting point for comparison.

1. In **Review**, confirm **v1 + original dev12 + query-only input + two evaluators + your Luna Judge**.
2. Name the evaluation `lab-en-learning-loop` and, if available, name the run `baseline-v1`. Use handoff names in a prepared class. If the baseline is already completed, open its results instead of submitting again.
3. For an authorized job not already submitted, select **Submit** once. The run is complete only after you confirm **Completed and all 12 results**, not immediately after clicking.

Microsoft Foundry / proj-fev11-dea3cec5

Q Search with AI (36K)

New Foundry

Home Discover Build Operate Manage Docs

Create

Agents

Models

Services

Tools

Knowledge

Memory

Guardrails

Data

Optimize

Evaluations

Fine-tune

Create new evaluation

Target: Agent

Scope

Frequency

Data

Configure agents

Criteria

Review

Add a name and description to this evaluation run.

Evaluation name *

contoso-en-learning-loop

Back Submit

Summary

Agent evaluation

Targets

contoso-eval-en: v1

Agent

Edit

Scope

Individual turns

Scope

Edit

Frequency

One time

Frequency

Edit

Dataset

contoso-eval-en-dev12

Version 1

Existing dataset

Edit

Evaluators

TaskAdherence

Agents

Edit

Relevance

Quality

Edit

Review before Submit. Confirm Agent version 1, the same dev12, query-only input, both evaluators, and the Judge. Submit once within the approved scope, then inspect the run status. View full-size image (file: `web/assets/portal/en/17-evaluation-review.png`)

Open **Evaluations** → **your evaluation name** → **Evaluation runs** → **baseline run**.

State	What to do
Running / In Progress	Refresh the same run and wait. Do not submit another.
Completed	Check all 12 results and error counts, then continue to 07. This does not mean every answer is correct.
Failed / Partial, or missing results	Preserve the error/run ID and use evaluation troubleshooting . Do not mark it complete.

If it remains unfinished after 30 minutes, record the state/error and notify the facilitator. Stopping your wait does not cancel the remote job.

Use this **read-only command** to retrieve actual evaluation and run IDs. If you chose a different name, use that exact name with `--name`.

```
python -m lab --config .lab/lab-en/.env native-evals --name lab-en-learning-loop
```

Record `evaluation_id` and the `run_id` whose `agent_version` is "1" **and** status is completed. Evaluation IDs use `eval_...`; run IDs use `evalrun_...`. They are not interchangeable. If several evaluations have the same name, compare their portal creation times, Agents, and runs to select yours. This command submits nothing.

Completion criteria: The real Foundry run is Completed and exposes 12 output items. Record failed/error counts in `result_counts` too. A failed or incomplete run stays failed/incomplete. No local custom Judge substitutes for this managed Evaluation.

07. Read scores and reasons

What: Connect actual responses and evaluator reasons to policy evidence.

Why: An average can hide a date, citation, or routing failure. Explain the defect before deciding which instruction to change.

How • where: Read Foundry detailed metrics, User view, and the source policy together; record a hypothesis and behaviors to preserve in `notes.md`.

1. Open the completed baseline run from 06 and read its pass/fail/error counts.
2. In **Detailed metrics result**, select a failing or lowest-score row and read `Relevance.reason` and `TaskAdherence.reason`. A reason explains why that score was assigned.
3. Open the row's `conversation_id` → `User view` to read the actual question/answer. Return to detailed metrics for scoring reasons; they are in a different view.
4. Find the cited policy ID in the source policies (file: `data/en/knowledge/documents.json`). Review a good case the same way. If no failure exists, record that fact rather than weakening v1 to create one.

Microsoft Foundry / proj-fev11-dea3cec5

Search with AI (96K)

New Foundry

Home

Discover

Build

Operate

Manage

Docs

Create

Agents

Models

Services

Tools

Knowledge

Memory

Guardrails

Data

Optimize

Evaluations

Fine-tune

contoso-eval-en

Make recurring

Delete

Run details

Status

Completed

Create time

10/1/2026, 12:21:37 AM

Duration

3m 59.84s

Created by

Agent ID

contoso-eval-en:1

See all properties

Raw JSON

User logs

Download user logs

Results

Download results

Overall metric results

P50 latency	P95 latency	TaskAdherence	Relevance
5.9s	8.8s	100% 12 / 12	83% 10 / 12

Detailed metrics result

Analyze Results

Index	query	conversation_id	trace_id	Relevance
0	Next month's cost is a concern. Please reduce my subscription. I...	conv_07e180556bd51ca800UPWqFCcJqavhmEPmHcTJBXGDHbqHYFFz	-	Fail: 3
1	My first monthly payment was on August 28, 2026 at 08:00 KST. I...	conv_0bb9168bd3bf11ec00jYpVdLhJy3iE1IHkA6uCYJS9YccWplx	-	Pass: 5
2	I used the service for just one day after my first monthly payment on...	conv_05f494a8836c0111004DizSunhTEdBq7cEwzgppeP3f4M5UKYcT	-	Pass: 4
3	I typed only a title on the Support > Contact support screen and left...	conv_09e82e7f1e3db55e00h7GlpAlJkGBN9Ngm0jFoOUAXyEm0k5r	-	Pass: 4
4	There is an unfamiliar access event on the administrator's audit screen...	conv_05740a8abb5d88ad00G5SaTweilR5t4uMTWWldYBV0kq4lgQVM	-	Pass: 5
5	Assume I paid the same monthly base fee in two months, with...	conv_098f970e5fa2d1da00zhmGbfVSMhdOFzI7lLUcjB4PGVLH8UFI	-	Pass: 5
6	I'm on paid Pro, but only beta webhook delivery failed; the...	conv_099677dc019f239900fUjqcz6Qpnwkk97SDxfp0SUv8pXxqrUQ	-	Pass: 5
7	Since there is a 99.9% monthly commitment, can I write in a...	conv_0a857d1c7aec9b6800pluwLFZNSVYwbrpuCehFCbQhFYecEUo	-	Pass: 5
8	I learned that the refund window includes the final statement after...	conv_0eb63b9bf392360b00tQNh20XhDm5lF7uul1mC4HrEDrGK7NA	-	Pass: 5

Read per-question scores and reasons. Check pass, fail, and error counts, then open Detailed metrics result for Relevance.reason and TaskAdherence.reason. Inspect individual cases, not only averages. View full-size image (file: web/assets/portal/en/18-evaluation-results.png)

Relevance 4/5 is not 80% accuracy. TaskAdherence 1 means Pass, not a low five-point score. Missing values are not zeros or successful rows. Generic evaluators do not certify every business rule.

Illustrative scores, not results from an actual run:

Example scores for one answer	Interpretation
Relevance 3 , TaskAdherence 1	It followed the task but missed the relevance threshold of 4. It did not pass both criteria.
Relevance 5 , TaskAdherence 1	It passed both scoring criteria. A person must still check policy dates, conditions, and citations.
Missing score	Check for an evaluation error or missing output. Do not record it as zero or a pass.

conv_07e180556bd61ca...

Conversation ID: conv_07e180556bd61c...

Trajectories

User view

Graph view

3 spans1 chat call1.8s1.1K

Next month's cost is a concern. Please reduce my subscription. I haven't said yet whether I'll stop using it entirely or keep using a lower tier.

("answer": "I understand your concern about the upcoming cost. You can reduce your subscription to a lower tier or make other changes to your subscription plan to better fit your needs. Please let me know if you'd like help with specific options or if you want assistance with pausing or canceling your subscription entirely. I'm here to support you in managing your subscription in the best way possible."; "citations": [], "route": "answer", "needs_human": false)

Compare the response with policy.

Read the question and JSON answer in User view, then check policy dates, conditions, and citations. Return to Detailed metrics result for the reasons; an incorrect response is not a reference answer. View full-size image (file: web/assets/portal/en/19-evaluation-case.png)

Watch for wrong date boundaries, unnecessary assumptions, numeric retrieval IDs instead of document citations, incorrect human routing, invented actions and unsupported certainty. A truthful statement of uncertainty must not be replaced with a fabricated fact to satisfy a Judge.

Create `.lab/lab-en/notes.md` in a text editor and fill in this worksheet with your results. `.md` is a plain-text note file; create your local lab-record folder first if it does not yet exist. Bring across the values recorded in 01–06. Keep full responses/reasons separately, and use the worksheet for identifiable runs/cases and concise observations.

Record	What to write
Baseline	Actual evaluation ID, run ID, Agent version, and dataset hash
Results	Per-metric pass/fail/error counts across all 12 cases
Problem case	Question, problematic sentence in the actual response, evaluator reason, and policy ID
Improvement hypothesis	For a date-boundary mistake, require checking effective dates and inclusive boundaries; do not embed the answer
Behaviors to preserve	Correct citations, honest uncertainty, and action boundaries already working well

Discuss: State the run ID, each metric's scale/pass count, a concrete problematic answer and the instruction behavior you want to improve.

Completion criteria: Use an actual case to explain what you intend to change, or why you would retain the current instructions. An average score alone is insufficient.

08. Improve instructions with Agent Optimizer

What: Generate instruction candidates with Agent Optimizer and review the changes.

Why: Candidate generation does not guarantee improvement. A higher internal rank cannot justify invented policy or changed model/tool conditions.

How • where: Run Instruction only in Foundry and inspect View changes. Save complete reviewed instructions and their actual provenance privately.

Agent Optimizer proposes and tests instruction improvements. A **candidate** is a proposed instruction set that has not yet been accepted. This does not retrain the model.

1. Open **Build** → **Agents** → **lab-en-iq** → **Optimize Preview/Optimize**.
2. If an **Agent / Cost** selection screen appears, select **Agent**, not **Cost**, which optimizes a different concern.
3. From the run list, select **Optimize** or **Create an optimization run / Create optimization run** and use the settings below. Preview means a pre-release feature. If the menu or model is unavailable, stop and follow [Optimizer troubleshooting](#).

Setting	Choice
Agent version	Explicit baseline 1
Choose targets	Instruction only (change instructions); Model and Tool description off
Max candidates	Operator-approved bound; this run uses at most 2 , not more releases
Optimization model	Your <code>OPTIMIZER_DEPLOYMENT</code> / gpt-5.5
Evaluation model	Your <code>JUDGE_DEPLOYMENT</code> / gpt-6-luna
Dataset	Same registered English dev12 version 1
Criteria	Relevance 4; TaskAdherence binary pass 1

Create an optimization run

1 Target

Optimization tunes your agent's instructions and tool descriptions. Optimization can also tune your agent's quality by attempting different models. To also compare models, turn on Compare across models and pick which ones to include.

Goal

Improve quality

Raise evaluation quality above your current version, with token cost left free.

Agent

contoso-eval-en

Version

v1 (latest)

Optimization approach

☐ Auto

Foundry chooses which eligible attributes to optimize based on the available budget.

☒ Choose targets

Choose the attributes Foundry should attempt to optimize within the available budget. Targets that do not apply to your agent are skipped.

☒ Instruction

☐ Tool description

☐ Model

Optimization model

lab-planner-dea3cec5

Max candidates

1

Evaluation model

lab-judge-luna-dea3cec5

Next

2 Data

3 Review



How optimization works

Scored and ranked

Each candidate is scored against your criteria and ranked by performance.

Verify before applying

Review the top candidate before applying changes to your agent.

Separate optimization scope and model roles. Select baseline version 1 and Instruction only, and keep candidates within the approved limit. Assign the prepared deployments to Optimization model and Evaluation model. View full-size image (file: <web/assets/portal/en/07-optimizer-target.png>)

If Criteria shows **No custom evaluators available**, switch **Custom only OFF** or select **View built-in evaluators**. Open each built-in row, set its correct threshold and Apply. Do not create a custom scorer to bypass a filter.

Create an optimization run

Target

Dataset

Criteria

Review

Choose an existing dataset to test each candidate against. Optimization runs the same dataset across every candidate so results are comparable.

Generate data

Select dataset and criteria

Back

Next

Select dataset

Search by file name

Create dataset from traces

Upload dataset

Name	Version	Created by	Created on
eval-data-2026-09-30_144335_a656c_UTC	1		9/30/26, 11:43:35 PM
contoso-eval-en-dev12	1		9/30/26, 11:20:29 PM

1-2 of 2

PrevNext

Preview

Dataset

contoso-eval-en-dev12

Fields

contextground_truthquery

Dataset preview (Top 5 rows)

context	ground_truth	query
[ATLAS-SUB-001] Subscription-3. Canceling a subscription stops the next automatic renewal; access continues until the end timestamp of the current	{"answer": "By 'reduce', do you m stopping automatic renewal, or on a lower plan from the next rer Those are different choices, so I clarify your intention first. This c has not changed the	Next month's cost is a concern. Please reduce my subscription. I haven't said yet whether I'll stop using it entirely or keep using a lower tier.
[ATLAS-REF-001] Refund-2. A first purchase of a prepaid monthly subscription made before 2026-09-01 00:00:00 KST retains the previous 7 x 24-hour application window,	{"answer": "The instruction in the notice does not change the polici 28 purchase retains the previou window, whose final timestamp September 4 at 08:00 KST. Septe past the deadline; the month of	My first monthly payment was on August 28, 2026 at 08:00 KST. I haven't used it, and I'm contacting you on September 6. The notice I attached says 'Everyone
[ATLAS-REF-001] Refund-1. A first purchase of a prepaid monthly subscription made at or after 2026-09-01	{"answer": "The standard conditi both no paid-credit use and no p production jobs. If even one paic job was run, the situation you de	I used the service for just one day after my first monthly payment on September 8, 2026. Paid-credit usage is

Reuse the same English data. Select the dev12 registration/version used for the baseline and include all 12 cases. Editing or regenerating the file changes the comparison conditions. View full-size image (file: `web/assets/portal/en/08-optimizer-dataset.png`)

In **Review**, confirm the Agent/version, dataset, evaluators, and estimated cost, then **Submit** once within the approved scope. An estimate is not a billing cap. Open the job in **Optimization runs**, wait at most 60 minutes, and retain its actual state. Reuse the same job ID when resuming. Optimization includes multiple internal calls.

Inspect original/candidate scores and **View changes**. The internal **0–1 ranking** is not the separate managed Evaluation mean or pass percentage. Keep model, tools, reasoning and output schema unchanged; an empty function-tool export does not authorize removal of the MCP policy connection.

Microsoft Foundry / proj-fev11-dea3cec5

Search with AI (96K)

New Foundry

Home

Discover

Build

Operate

Manage

Docs

Create

Agents

Models

Services

Tools

Knowledge

Memory

Guardrails

Data

Optimize

Evaluations

Fine-tune

← opt_e44bcf5701a348deb62a1cd4f9cb3910 Succeeded

Download JSON

Promote candidate

Run details

Improvement

Goal: Improve quality Improved

candidate_1

Avg. score

0.646 +0.010

Baseline: 0.635

Task-weighted average (0-1)

View changes

Download config

Run overview

Candidates: 1

Best score: 0.646

Eval model: lab-judge-luna-dea3cec5

Dataset: contoso-eval-en...

Elapsed time: 9m 36s

Total tokens: 264.3K

View details

Additional settings

Score comparison

Task-weighted average (0-1)

candidate_1 0.646

baseline 0.635

View overall score details

Candidate results

Iteration	Candidate	Optimizations	Score	Score delta	Avg tokens	Avg latency	Changes	Config	Score details
0	baseline	--	0.635	--	2.8K	6.1s	Version 1	↓	evalrun_f9836666...
1	candidate_1 * Best	System prompt	0.646	+0.010	3.3K	8.1s	View changes	↓	evalrun_0424454...

1-2 of 2

Prev Next

Read the candidate, not just the ranking. Compare per-evaluator scores and review the instruction changes. If no candidate offers a sound improvement, retain v1 and record why. View full-size image (file: `web/assets/portal/en/09-optimizer-results.png`)

Changes: candidate_1 ★ Best candidate



Comparing candidate_1 against baseline configuration.

Baseline

```
1 {
2   "agentName": "contoso-eval-en",
3   "agentVersion": "1",
4   "system_prompt": "You are a customer support assistant for
Contoso Atlas Cloud. Answer in English using the customer's
inquiry and the supplied policy context.\nFor a policy
question, if the knowledge_base_retrieve knowledge tool is
connected, use it to check the evidence before answering. If
no tool or context is available, state the limitation rather
than inventing policy details.\nThe response must be a single
JSON object containing only the following four fields, with
no other text or Markdown.\n{\n  \"answer\": \"Answer text\",
  \"citations\": [\n    \"Document identifier\"],
  \"route\": \"answer\",
  \"needs_human\": false\n}\nroute must be one of
answer, clarify, escalate, or refuse. citations is an array
of strings. needs_human is true only when route is escalate
and false otherwise.\n",
5   "skills": [],
6   "tools": []
7 }
```

Candidate: candidate_1

```
1 {
2   "agentName": "contoso-eval-en",
3   "agentVersion": "1",
4   "system_prompt": "You are a customer support assistant for
Contoso Atlas Cloud. Answer in English using the customer's
inquiry, any supplied policy context, and the evidence rules
below.\n\nYour final response must be a single valid JSON
object and nothing else. Do not include Markdown, code
fences, tool traces, arrays outside the JSON object, or
citation footnote markup inside the answer text.\n\nRequired
JSON shape:\n{\n  \"answer\": \"Answer text\",
  \"citations\":
  [\n    \"Document identifier\"],
  \"route\": \"answer\",
  \"needs_human\": false\n}\n\nField rules:\n- Include exactly
these four fields: answer, citations, route, needs_human.\n-
route must be one of: answer, clarify, escalate, refuse.\n-
needs_human must be true only when route is \"escalate\"; it
must be false for answer, clarify, and refuse.\n- citations
must be an array of document identifiers only, such as
\"ATLAS-SLA-001\", \"ATLAS-SUB-001\", or \"ATLAS-REF-001\".
Do not put URLs, inline citation markers, or quoted snippets
in citations.\n- The answer field should be concise but
complete, and should not contain unsupported policy
claims.\n\nEvidence and tool-use rules:\n- For any policy
question, if the knowledge_base_retrieve tool is connected,
call it before answering.\n- Use the retrieved policy
content, supplied policy context, and the embedded policy
```

< 1 of 1 >

Close

Read the instruction changes. Use View changes to identify which response behaviors change. Reject invented policy, unsupported certainty, and non-instruction configuration changes; assess content and effect rather than length. View full-size image (file: [web/assets/portal/en/10-optimizer-changes.png](#))

Save a candidate only after reviewing it.

1. Select the candidate and inspect **View changes**. Reject instructions that invent policy or unsupported certainty.
2. Copy the **complete revised instructions**. If the view shows only changed lines, ask the operator to confirm the full text rather than using those fragments as the instruction file.
3. Save **UTF-8 plain text** to `.lab/lab-en/candidate.txt`. Do not use Word/rich text; check that the name is not `candidate.txt.txt`. Exclude diff (change comparison) `+/-` markers, UI explanations, and scores.
4. Record real job/candidate IDs and any manual edits/reasons in `notes.md`. Use this file next, not the repository's `prompts/en/candidate.txt`.

Do not continually create v3, v4 and later releases. Use the next step's CLI to create one v2 for comparison, so **do not use both Promote candidate and the CLI creation command**. Creating a version is separate from publishing it or approving activation.

Discuss: Present the reviewed candidate and explain the changed behaviors, expected improvements, and possible regressions.

Completion criteria: Record real job/candidate IDs, the reviewed instruction file, and change reasons. If no candidate is worth retaining, record "keep v1; no improvement observed" and continue at 10. Do not repeatedly run the same job to manufacture improvement.

Next path: With a reviewed candidate file, continue to [09 reevaluation](#). Without one, go to [10 cleanup](#). Do not create v2 just to have another version.

09. Reevaluate and compare v1/v2

What: Reevaluate changed instructions under the same definition and record retain v1, accept v2, or hold.

Why: A v2 label is not improvement evidence. Consider actual errors, latency/tokens, and statistical uncertainty alongside quality.

How • where: Prepare explicit v2 and its same-definition run in the terminal, then compare every case in Foundry Compare runs. Do not publish to production.

For a prepared environment, confirm who runs commands first. Version creation checks ownership records. If you have no participant-specific configuration and execution records, the owning operator runs these commands. Do not bypass that check by replacing the user name in their configuration. Participants compare the two supplied runs in the portal.

1. Create v2 from the reviewed instructions. Keep models, tools, data, evaluators, and Judge fixed, and preserve v1.

```
python -m lab --config .lab/lab-en/.env native-agent --version 2 --prompt .lab/lab-en/candidate.txt --confirm
```

Confirm `agent_name: lab-en-iq` and `version: "2"`. If a different v2 already exists or the model deployment changed, stop and preserve the original records. A new version alone does not demonstrate improvement.

2. Add a v2 run to the same evaluation configuration. The helper below is a supplied Python program. It uses the official Azure AI Projects/OpenAI SDK to reuse the baseline's dataset and evaluator settings, rather than creating a different evaluation.

Placeholder	Where to find your value
YOUR_PROJECT_ENDPOINT	AZURE_AI_PROJECT_ENDPOINT in .lab/lab-en/.env
YOUR_SUBSCRIPTION_ID	The subscription ID from 01, also AZURE_SUBSCRIPTION_ID in that .env
YOUR_EVALUATION_ID	The selected evaluation_id from the native-evals output in 06
YOUR_BASELINE_RUN_ID	That evaluation's completed version-1 run_id, not an Optimizer job ID

Replace all four values and execute this single-line command. Submitting a new run incurs charges.

```
python scripts/add_foundry_eval_run.py --endpoint "YOUR_PROJECT_ENDPOINT" --subscription "YOUR_SUBSCRIPTION_ID" --evaluation "YOUR_EVALUATION_ID" --baseline "YOUR_BASELINE_RUN_ID" --version 2 --name candidate-v2 --out .lab/lab-en/artifacts/foundry-evaluations/candidate-v2.json
```

Confirm status: completed and all 12 cases. If the default 30-minute wait expires, the command returns exit code 2 with Still running. If the receipt contains a run ID, repeat the **identical command and --out path** to collect that run. For an unknown submission without a run ID or an existing remote name, follow [duplicate-submission recovery](#). Do not delete the receipt or rename the run to resubmit.

The helper verifies thresholds, Judge, mappings, and each output item's Agent version and instructions.

3. Return to Foundry and compare. In the same evaluation's **Evaluation runs**, select the v1 and v2 rows and **Compare runs**. Explicitly choose **v1 as Baseline**; selection order must not reverse the comparison.

Microsoft Foundry / proj-fev11-dea3cec5

Search with AI (K)

New Foundry Home Discover Build Operate Manage Docs

Create

- Agents
- Models
- Services
- Tools
- Knowledge
- Memory
- Guardrails
- Data

Optimize

- Evaluations
- Fine-tune

← Statistically compare results

Total runs selected: 2 Analysis method: PairedTTest

Baseline: contoso-eval-en

Name	contoso-eval-en	candidate-v2
TaskAdherence	1.00 Baseline	1.00 Inconclusive
Relevance	4.42 Baseline	4.33 Inconclusive

1-2 of 2

Confirm the comparison direction and results. Select v1 as Baseline, then compare scores, pass counts, and statistical results. Inconclusive means a difference was not established; it is not evidence of equivalence. View full-size image (file: web/assets/portal/en/20-evaluation-comparison.png)

Compare responses, scores, and reasons for all 12 cases side by side. Require no decline in all-criteria passes or either metric's pass count/mean, and at least one strict measured improvement. Check factual truth, routing, and response format too, and record latency, tokens, and statistical results.

Explain the result: Identify the actual gain, unchanged criteria, any regression, and the remaining uncertainty. Observed improvement is not a guarantee that every future stochastic run will improve.

In `notes.md`, record v1/v2 pass counts and means per metric, errors, latency/tokens, statistical results, and your decision side by side. If improvement criteria are not met, **retain v1 or hold the decision**. Do not adopt a candidate with actual policy errors merely because generic evaluators passed it.

Latency is time spent waiting for a response; **tokens** measure the amount of text processed by the model. A small score improvement can come with increased waiting time and cost.

Decision	When it fits
Consider v2 for adoption	All 12 cases use the same conditions, quality does not regress and at least one measure improves, no policy errors are found, and latency/token tradeoffs were reviewed. Production approval is separate.
Retain v1	No improvement, or previously correct behavior got worse. This is a valid learning outcome.
Hold the decision	Errors, missing results, or different conditions prevent a fair judgment. Record the missing evidence.

Completion criteria: Record the two actual run IDs, complete same-criteria comparison, and the reasons for accepting, retaining, or holding. Do not accumulate release numbers. Production approval and independent generalization are separate and are not granted by this dev12 exercise.

10. Save results and delete resources

What: Preserve results, then remove your lab resources or hand them to an approved retention owner.

Why: Search and logs can remain after the browser closes. Deleting the wrong shared group can also remove someone else's resources.

How · where: Verify exact targets, ownership, and authorization in Azure Portal and the terminal, then verify absence after deletion.

Closing a browser, deleting an Agent, or running `cleanup` does not by itself stop all resource-group charges.

Preserve records before deletion

- 1. Record the actual project, Agent, dataset version/hash, evaluation/run/job IDs, and decision in `notes.md`.
- 2. Use **Download/Export** where offered in the evaluation and Optimizer views. If unavailable, retain existing receipts and detailed results; do not claim an export you could not obtain.
- 3. Retain `.lab/lab-en/` plans, manifest, approvals, and run records in an organization-approved private location. Do not commit raw credentials, cookies, signed URLs, or account details. Public records should contain only the necessary synthetic cases, scores, and reasons.
- 4. Inspect Evaluations and Optimization runs for active jobs. Use **Cancel** where supported and verify the terminal state. Hand unresolved active jobs to the cleanup owner.

Distinguish dedicated and shared environments

Environment	Cleanup scope
Dedicated resource group created for you in 02	Follow the group-deletion procedure below. Reconfirm that no unrelated resources are present and deletion is authorized.
Dedicated environment explicitly approved for retention	Do not run deletion commands. Record remaining resources, the retention reason, cost owner, and review date; verify that the resources remain and hand them over.
Operator-provided shared project	Do not delete the resource group, Foundry resource, shared models, or Search service. Remove only assigned objects and hand remaining resources to the operator.
Setup or execution stopped partway through	Compare <code>config.json</code> , the manifest, and actual Azure resources. Failure does not mean nothing was created.

Follow only the path matching your environment. In a shared environment without your own configuration and local ownership records, hand CLI cleanup to the operator too. Do not copy another person's records to perform deletion.

For shared resources with your own local ownership records, first inspect this **plan-only** command:

```
python -m lab --config .lab/lab-en/.env cleanup
```

Read mode: `LOCAL_PLAN_ONLY`, `actions`, `never_deleted`, and `manual_follow_up`. Verify names, project, and ownership. Run the following only after obtaining authorization for exactly those deletions. If your environment name differs, use its matching prefix.

```
python -m lab --config .lab/lab-en/.env cleanup --confirm-prefix lab-en
```

Confirm `OWNED_OBJECTS_ABSENT`. This removes recorded Agent versions, search objects, and connections, but **not resource groups, model deployments, Search hosting, logs, or RBAC**. It does not automatically remove every portal-created dataset, evaluation, Optimizer job, Playground conversation, or model-smoke response. Delete authorized items individually in their owning view; if no delete action is offered, hand retention or project cleanup to the operator.

Delete a dedicated group and verify its absence


Replace `YOUR_LAB_RESOURCE_GROUP` with `names.resource_group` from `config.json`. Verify subscription, ownership tags, and the entire inventory, not just a familiar prefix.

```
az group show --subscription "YOUR_SUBSCRIPTION_ID" --name "YOUR_LAB_RESOURCE_GROUP" --query "{name:name,location:location,tags:tags}" -o json
az resource list --subscription "YOUR_SUBSCRIPTION_ID" --resource-group "YOUR_LAB_RESOURCE_GROUP" --query "[].{name:name,type:type}" -o table
```

Resource-group deletion is irreversible and removes the group's Foundry resources, model deployments, Search, and monitoring together. After saving evidence and receiving approval for that exact group, choose one method:

The default Application Insights Smart Detection action group can also serve alerts in other resource groups. A dedicated lab group does not establish that every resource is independent. Review shared dependencies with the owner before deletion, or retain the group; do not delete alerts, permissions, or locks simply to bypass a check.

1. **Portal:** Azure Portal → Resource groups → exact group → **Delete resource group**. Read the deletion inventory, type the requested group name, and confirm.
2. **CLI:** Run the command below and review the target again at the confirmation prompt. Do not append `--yes` to bypass confirmation.



Read the proposed inventory and confirmation field first. Verify the group name and every resource. Only after authorization, enter the group name and proceed with Delete. Use Cancel if the target is wrong or deletion is not approved. View full-size image (file: web/assets/portal/en/28-delete-review.png)

```
az group delete --subscription "YOUR_SUBSCRIPTION_ID" --name "YOUR_LAB_RESOURCE_GROUP"
```

An accepted request is not completed deletion. Check portal notifications and group status, then verify that this command successfully prints false:

```
az group exists --subscription "YOUR_SUBSCRIPTION_ID" --name "YOUR_LAB_RESOURCE_GROUP"
```

Authentication or network errors are not false and do not prove deletion. If true, inspect deletion progress. For **Locks**, permissions, or dependencies in other groups, follow [deletion troubleshooting](#). Do not remove organizational locks without authorization.

Afterward, open **Cost Management → Cost analysis** for the subscription, time range, and group. Billing updates can lag, and earlier usage charges do not disappear. Separately check logs/storage in other groups and service-specific soft-deleted resources. Permanent deletion/purge requires organizational policy and separate authorization.

Final completion criteria: Verify an authorized dedicated-group deletion and record the time, or hand over an explicitly retained dedicated/shared-resource inventory with reasons, cost owner, and retention review date. Do not delete .lab first and lose ownership evidence. See the [operator cleanup worksheet](#) and [deletion troubleshooting](#).

Facilitator guide · improve the measured result

[Ten participant steps](#) · [Operator handoff](#) · [Troubleshooting](#)

Teach **evaluate → learn → improve → reevaluate** using business-specific tasks and real managed results. Do not manufacture a weak baseline or promise that every future evaluation will improve.

Before class

Decide whether participants perform 01–03 or receive an operator-prepared environment. Self-setup requires access, cost authorization, and a cleanup owner. Prepared classes need a complete [handoff](#) and verified project, Agent, retrieval, and dataset registration.

Separate environment preparation from learning evaluation in beginner classes. Prefer operator preparation for 02–03, but still have participants verify their own account, project, and completion evidence. Assign SDK-command responsibility before class; do not distribute the operator's sign-in or ownership records.

Start together with the [six basic terms](#) and [question → answer → scoring diagram](#). Ask participants to explain that **the Agent answers and the Judge scores**, and that **instructions, not model weights, will change**. Do not require memorized terminology or high scores to begin.

Steps 04–06 use one evaluation wizard, not three evaluations. Read one scored example together in 07 before independent analysis. If a group shares an Agent, assign one owner for v2 creation and paid submissions.

New lab Agents use `lab-en-iq` or `lab-ko-iq`. If the environment prefix differs, use the actual name printed by the command. Verify [model roles and identity](#), and reserve time for cleanup in 10.

Both v1 and v2 are full immutable Agent versions. Their model, tools, reasoning, and strict output schema match; only instructions differ. Do not weaken v1 to create improvement. The default path creates one reviewed v2 and reevaluates it; only operators who verified draft support use a separate draft workflow.

Keep English and Korean data/Agents separate and compare each language's actual executions independently. Translating another language's result does not constitute a new run.

Verify completion across all ten steps

Step	Required action	Completion evidence
01 Environment	Check account, subscription, permissions, Python, CLI, and language variables	Matching identity and successful local commands
02 Provisioning	Plan, authorize, preflight, apply, and inspect status	APPLIED, generated <code>.env</code> , and runtime preflight PASS
03 Agent setup	Verify model/retrieval and create <code>native-agent v1</code>	Actual JSON response and Agent tool call
04 Dataset	Register or reuse unchanged three-column JSONL	All 12 rows, registration/version, and SHA-256
05 Criteria	Relevance 4, TaskAdherence 1, and the actual Judge	Saved remote definition and query-only input
06 Baseline	Foundry Evaluation against pinned v1	Completed actual run ID and all 12 items
07 Analysis	Read actual answers, scores, reasons, and policies	Concrete improvement hypothesis, including failures
08 Optimization	Instruction-only optimization and candidate-diff review	Job/candidate IDs, reviewed file and provenance, or the reason to retain v1 without a different candidate
09 Reevaluation	With a reviewed candidate, evaluate same-criteria v2 and open Compare runs	Complete paired evidence and retain/accept/hold decision; without a different candidate, continue to 10
10 Cleanup	Preserve results and ownership, then perform authorized deletion or retention	Authorized group deletion or an approved dedicated/shared-resource retention handoff

Use **Individual turns / One time / Existing dataset**. The upload preview can show only five rows; the experiment still contains 12. Version selection can clear the target checkbox: reselect it and confirm one target.

Discuss the actual cases

Ask participants to connect **question → actual answer → evaluator score and reason → policy/reference answer**. `conversation_id` → `User view` shows the conversation; the evaluator reasons belong in **Detailed metrics result**, such as `Relevance.reason` and `TaskAdherence.reason`.

Discuss at least one weak case and one strong case. If no failure is observed, use the lowest-scoring case and say no failure was found. Never weaken v1, omit a difficult row, change a reference answer or substitute local scores.

Keep the selected lab's twelve synthetic questions, actual responses and managed evaluator reasons in separate private run records. Preserve unsuccessful outcomes; remove tokens, cookies, signed URLs and private account metadata if results need to be shared.

Interpret improvements correctly

Relevance 4/5 is not 80% accuracy. TaskAdherence 1 is Pass, not a poor one-out-of-five score. Report each evaluator separately, all-criteria passes, errors and coverage.

`scripts/compare_foundry_eval.py` requires complete matched cases, identical data/model/tools/settings, per-item version/instruction attestation, no lower pass counts or means, and at least one strict measured quality improvement. Malformed responses remain failures; the helper never repairs them into a pass.

Read answers for unsupported certainty, incorrect policy arithmetic, action claims and wrong routing even if the generic evaluator passed them. A request demanding an impossible certainty must not be satisfied with invented facts.

Optimizer ranking is internal selection evidence, not the direct managed reevaluation mean or pass rate. A higher ranking alone does not establish a better v2.

Recover without changing the experiment

Symptom	Response
Wrong identity, 401/403	Confirm the approved account/project; do not widen shared permissions
Catalog or deployment succeeds but Agent fails	Record the actual API/runtime error; verify the intended role
Portal Add run item-schema error	Use the official SDK helper with the existing baseline data source, not a different scorer
No custom evaluators available	Switch Custom only OFF or select View built-in evaluators
Duplicate local receipt or existing remote run	Retrieve that same run; do not rename it to force another submission
Candidate changes model/tools/reasoning/schema	Reject an instruction-only comparison claim

Symptom	Response
Draft becomes a numbered release	Stop; draft support is not enabled as expected
Existing released v2 differs from local source	Do not overwrite it or create v3; preserve the release and inspect the mismatch
Missing output or metrics	Report incomplete execution; do not reduce the denominator
Statistical comparison is Inconclusive	Report that result; do not claim significance or equivalence

Relevance maps to `{{sample.output_text}}`; TaskAdherence maps to `{{sample.output_items}}`. Do not overwrite service-generated bindings or feed `ground_truth` into the Agent. The actual remote Judge overrides any stale local default.

After a timeout, identify **which existing job ID to resume** first. Use the [issue record](#) to prevent deletion/recreation or paid retries “until it succeeds.” Do not mark a blocked step completed.

Correct common misconceptions

Claim	Correction
“V2 is a prompt file number only.”	It is a complete Agent version; this experiment changes only its instructions
“A more powerful model makes v2 better than old-model v1.”	Both versions must use the same model for the instruction comparison
“A better v2 can be guaranteed before execution.”	Only observed improvement can be established; future stochastic results are not guaranteed
“Keep only successful rows.”	Keep every case, including failures and missing/error outcomes
“The source schema certifies correct business decisions.”	It controls output structure, not policy truth or proper routing
“Public evidence must include cookies and raw logs.”	Publish actual synthetic responses and reasons through an allowlist, not authentication data

Finish with a separate lab record

Hand over the actual environment, run IDs, comparison conditions and final decision privately. Do not append execution results or validation history to reusable guides; correct instructions only when they are wrong.

Report the same-model result, native statistical comparison, any latency/token tradeoffs and the remaining limits of a reused synthetic dev set. Candidate activation is a separate operator decision; no production publication is part of this workshop.

Before ending class, perform [the cleanup/retention checks in 10](#) together. Confirm authorized deletion or a cost owner and review date for retained resources. Do not delete the records needed to establish ownership or reopen the environment.

Operator setup, authorization, and cleanup

[Ten steps from the beginning](#) · [Facilitator](#) · [Troubleshooting](#)

Use this reference for participant steps **02 provisioning authorization, 03 Agent setup, and 10 deletion**. Participants preparing their own environments can use it too. Supply actual values and completion evidence rather than leaving prerequisites implicit.

Distinguish new and prepared environments

Situation	Action
No project exists	Start with 01 account setup and use the authorized bootstrap plan → preflight → apply → status path.
The operator already prepared this lab	Supply actual project, Agent, data, and access through the handoff ; participants do not duplicate creation.
Only an existing shared project is available	Bootstrap does not silently adopt existing groups. Its owner must define the preparation/change scope and approve project-specific access.


Verify CLI, SDK, and portal identities separately. Never transfer the operator's token or sign-in session to a participant. The owning operator runs bootstrap-bound creation commands. Other participants use their own assigned portal roles; SDK access requires separate configuration and permissions matching their own identity.

For groups new to Azure and Foundry, operator-prepared environments are recommended.

Participants check their own account/computer in 01, receive completion evidence for 02–03, and continue at 04. Without participant-specific SDK configuration, the operator handles ID lookup in 06, v2 creation/reevaluation commands in 09, and owned-object cleanup in 10. Participants inspect settings/results using their own portal access.

Use separate environments/Agents for individual labs. For a group sharing one Agent, designate **one person to create v2 and submit paid jobs**. Record that responsibility so several participants do not attempt different v2 releases or duplicate the same evaluation.

Default Agent names are lab-en-iq for English and lab-ko-iq for Korean. If the environment name changes, use the actual name returned by the CLI and record it in the handoff.



Check the isolated resource group. Open the group named in the handoff and confirm its subscription, region, and resource inventory. Verify that the Foundry project and supporting resources belong to this lab. View full-size image (file: [web/assets/portal/en/00-resource-group.png](#))

Verify model roles and actual deployments

Role	Configuration key	Default model/version
Agent	MODEL_DEPLOYMENT	gpt-6-sol / 2026-09-22, GlobalStandard 100
Managed Evaluation Judge	JUDGE_DEPLOYMENT	gpt-6-luna / 2026-09-22, GlobalStandard 100
Optimizer/search planner	OPTIMIZER_DEPLOYMENT / IQ_PLANNER_DEPLOYMENT	gpt-5.5 / 2026-04-24, one shared GlobalStandard 100 deployment
Policy embeddings	EMBEDDING_DEPLOYMENT	text-embedding-3-small / 1, GlobalStandard 10

Numbers are requested ARM capacity units, not a universal TPM conversion. These defaults require availability checks; they are not a deployment guarantee for every subscription. Inspect deployment names, model names, versions, and readiness in Foundry **Models + endpoints/Build → Models**. `.env` contains actual deployment names rather than product names.

Distinguish additional monitoring resources. Application Insights can automatically add a [default Failure Anomalies alert and Smart Detection action group](https://learn.microsoft.com/azure/azure-monitor/alerts/proactive-failure-diagnostics#alert-rule-creation) (<https://learn.microsoft.com/azure/azure-monitor/alerts/proactive-failure-diagnostics#alert-rule-creation>). Bootstrap checks read-only that the alert targets only the owned Application Insights component and that the linked action group uses only the default role receivers. It does not adopt resources by name or modify/delete a shared action group in another resource group. If those links are not yet observable during provisioning, inspect the same `bootstrap` status again; do not edit the manifest or delete the alert to bypass the check.

Set TPM before starting model calls

The [participant step-02 recommendations](#) require **100,000 TPM** per Agent, Judge, and shared Optimizer/planner deployment, and **10,000 TPM** for embeddings. This assumes sequential jobs in one environment; quota pools remain model/SKU/region-specific. For simultaneous English/Korean environments, sum both allocations for each model. Do not count two roles sharing one deployment as two deployments.

- 1. New dedicated environment:** Check generative capacity 100 and embedding capacity 10 in the updated `bootstrap` plan, then authorize and provision that exact plan. Existing plans do not change automatically. For an unprovisioned lower-capacity plan, prepare a new environment name and authorization rather than editing the hashed plan.
- 2. Separately managed prepared deployment:** Its owner checks **Tokens per Minute Rate Limit** under Foundry **Build → Models → Deployments → deployment name → Details**. Use **Edit** to meet these values only when the change is authorized. If the UI uses thousands of tokens, also check the final TPM display, save, and allow propagation. **Portal-only changes to a bootstrap-managed environment cause plan drift.** Do not edit hashes/approvals to bypass it; prepare a newly authorized dedicated environment with sufficient capacity when necessary.
- 3. Verify the setting:** With that environment's `.env` and language-specific artifacts path, run the read-only check below. Inspect each `*_tpm` entry's `observed`, `expected`, and `reason`.

```
python -m lab --config .lab/lab-en/.env preflight
```

Continue to step-03 smoke/retrieval and later evaluation/Optimizer only after overall `status`: `PASS` and all five TPM checks pass. The checker reads the actual deployment's `rateLimits` entry with key: `token`. If the CLI omits labels, it reads raw ARM metadata for that same deployment instead of guessing TPM from request counts or `sku.capacity`. Unverifiable limits are `BLOCKED`.

TPM accounting uses estimates rather than average billed tokens; **RPM and short-window burst limits** also apply. Allow more headroom for longer inputs/maximum outputs, shared users, and overlapping evaluation/optimization. These recommended minimums do not guarantee a 429-free run. See [official rate-limit guidance](https://learn.microsoft.com/azure/foundry/openai/how-to/quota#understanding-rate-limits) (<https://learn.microsoft.com/azure/foundry/openai/how-to/quota#understanding-rate-limits>).

Use the [real response check in 03](#) to verify Agent and `knowledge_base_retrieve` calls. Optimizer has a separate [supported-model list](https://learn.microsoft.com/azure/foundry/agents/concepts/agent-optimizer-overview#models) (<https://learn.microsoft.com/azure/foundry/agents/concepts/agent-optimizer-overview#models>); Agent/Judge availability does not establish generator support.

Microsoft Foundry / proj-fev11-dea3cec5

Search (🔍)

Create

Agents

Models

Services

Tools

Knowledge

Guardrails

Data

Optimize

Evaluations

Fine-tune

Models

Deployments Batch jobs

Serverless deployments Managed compute deployments

Search

PTU Calculator

Deploy

Name	Model	Version	Deployment status	Deployment type	Created on ↓	Created by
☒ lab-judge-luna-dea3cec5	gpt-6-luna	2026-09-22	Succeeded	Global Standard	9/30/26, 11:19:02 ...	
☐ lab-agent-luna-dea3cec5	gpt-6-luna	2026-09-22	Succeeded	Global Standard	9/30/26, 11:18:59 ...	
☐ lab-embedding-dea3cec5	text-embedding-3...	1	Succeeded	Global Standard	9/30/26, 9:58:45 PM	
☐ lab-planner-dea3cec5	gpt-5.5	2026-04-24	Succeeded	Global Standard	9/30/26, 9:58:44 PM	
☐ lab-judge-dea3cec5	gpt-5.4-mini	2026-03-17	Succeeded	Global Standard	9/30/26, 9:58:42 PM	
☐ lab-agent-dea3cec5	gpt-4.1-mini	2025-04-14	Succeeded	Standard	9/30/26, 9:58:40 PM	

1-6 of 6

Prev Next

Check model, version, and status. Map each deployment to the Agent, Judge, Optimizer, and embedding roles. Deployment names must match `.env`; verify actual calls to the prepared deployments as well. View full-size image (file: `web/assets/portal/en/02-model-deployments.png`)

Prepare an executable Agent

First complete Python installation, virtual-environment creation, login, and language variables using the [copyable commands in 01](#). Do not activate a nonexistent `.venv` or treat fake IDs in `.env.example` as a live environment.

Use `prompts/en/baseline.txt` for English and `prompts/baseline.txt` for Korean. Never weaken the baseline to manufacture improvement. After `iq prepare` → `iq probe` succeeds, create a policy-connected v1:

```
python -m lab --config .lab/lab-en/.env native-agent --version 1 --confirm
```

`native-agent` uses `ensure_fixed_release` to allow only **1 and 2**, comparing local workspace ownership with remote metadata. It reuses matching configurations, rejects a different v2, and never creates v3. It also retains a model-deployment snapshot and rejects drift. Receipts `agents/native-v1.json`, `agents/native-`

v2.json, and agents/native-model.json belong to that environment's artifacts folder.

native_response_format() projects the common schema into the supported strict-generation subset. Citation uniqueness and needs_human/route consistency still require post-generation checks. Valid JSON does not establish policy correctness.

Creating v2 is not completing evaluation. The default path creates one unpublished v2 from reviewed instructions, runs the same managed evaluation, and then decides whether to accept it. Experienced operators may first test draft-... versions as in the historical rehearsal, but must verify subscription support. Draft validation is not an unstated prerequisite for the default path.

```
python -m lab --config .lab/lab-en/.env native-agent --version 2 --prompt .lab/lab-en/candidate.txt --confirm
```

Do not combine portal Promote and CLI creation. V1/v2 are complete Agent versions; the same Agent, model, tools, and output settings must differ only in instructions.

Retrieve IDs with native-evals --name lab-en-learning-loop. The reevaluation helper copies the baseline data source and checks the remote contract. Do not replace the remote definition with a stale local JUDGE_DEPLOYMENT.

Freeze the comparison conditions

Item	Fixed requirement
Data	Same selected-language dev12 bytes, registration/version, questions and references
Relevance	Managed builtin.relevance, scale 1–5, threshold 4
TaskAdherence	Managed builtin.task_adherence, binary 0/1, pass 1
Judge	Same explicit Luna deployment in both evaluator definitions
Agent	Same Sol model/version, tools, reasoning and strict output schema
Change	Instructions only; releases remain v1/v2
Optimizer	Instruction target only, model/tool-description changes off, bounded candidate count
Lab records	Keep same-condition v1/candidate results for all twelve cases in private run folders, separate from reusable guides

Catalog evaluator versions are not proof that the private service rubric is fully pinned. Preserve the actual definition, settings and this limitation. A passing gate on reused dev12 does not guarantee future scores, independent generalization or production approval.

The standalone `scripts/compare_foundry_eval.py` selects report language and the default dev12 file using `LAB_LANGUAGE` or `--language ko/--language en`. Match explicitly supplied datasets to that language; do not relabel another corpus's results.

Check evaluation response mappings

Participants select the [two evaluators and Judge in 05](#) and preserve automatic mappings. Reading Raw JSON or guessing SDK fields is not a beginner prerequisite.

Setting	Required value
Agent user input	Only <code>{{item.query}}</code> ; leave the instruction override unset
Relevance response	<code>response={{sample.output_text}}</code>
TaskAdherence response	<code>response={{sample.output_items}}</code>
Judge/thresholds	Actual <code>JUDGE_DEPLOYMENT</code> ; Relevance 4, TaskAdherence binary pass 1

If a required field is **Unassigned**, compare the remote definition, target Agent, and data columns. Never map the reference answer as the response or copy old UI bindings. Consult the [official portal evaluation guide](https://learn.microsoft.com/azure/foundry/how-to/evaluate-generative-ai-app) (<https://learn.microsoft.com/azure/foundry/how-to/evaluate-generative-ai-app>) without changing this lab's fixed conditions. The helper in 09 checks the remote contract too and does not suppress errors to submit a run.

Complete the actual authorization file

Open `approval.example.json` and **Save As `approval.json`**. This worksheet explains fields; it is not an approved request. Keep the actual organizational authorization private. Use double quotes for JSON strings, but not for numbers, true, false, or null.

Field	How to complete it
<code>schema_version</code> , <code>scope_sha256</code> , <code>models</code> , <code>retention_days</code>	Preserve generated values. Editing the model array or plan invalidates existing authorization.
<code>approved</code>	Set <code>true</code> only after actual approval.
<code>approved_by</code>	The operator sign-in name matching <code>expected_user</code> in <code>config.json</code> .

Field	How to complete it
approved_at, expires_at	Actual authorization start/end in timezone-aware ISO 8601, such as 2026-10-03T09:00:00+09:00. Do not reuse the example date. The current time must fall within the interval.
currency, budget_amount, budget_policy	Approved three-letter currency and positive numeric budget. Keep policy BOUNDED ; the budget owner determines the amount.
acknowledge_no_monetary_cap	Keep false . Unlimited spending is not the default lab policy.
max_hosting_hours	Approved positive integer hours. If 8, the approval interval must also be no longer than 8 hours.
max_wait_seconds	Positive integer wait limit, for example 3600. Ending the wait does not cancel a remote job.
max_calls, max_candidates	Approved call/candidate limits. The candidate example is 2; consider internal Agent, Judge, and search calls when setting the call limit.
allow_global_inference	true only when GlobalStandard processing is authorized; it is not NCUS-only processing.
allow_resource_creation, allow_rbac_assignments	Each is true only when this plan's resources and resource-scoped role assignments are authorized.
acknowledge_continuous_hosting, acknowledge_unknown_cost	Each is true only after acknowledging continuing hosting charges and an unconfirmed final cost.
allow_training, allow_global_training	false for this path.
max_epochs, max_training_jobs	0 for this path.
accept_deprecated_models	Default false ; do not change it without reviewing a service deprecation notice.

Validate with `preflight --approval ...` and read `approval_reason` on failure. Authorization covers only the exact plan hash, models, and retention scope. It does not authorize another class, an automatic retry, or deletion.

Enforcement boundary: Bootstrap checks authorization validity, provisioning scope, and its wait bound. JSON budget and call limits do not automatically cap all Azure Portal or Optimizer spending. Budget alerts are not cutoffs. The operator monitors actual usage and cancels jobs or removes resources when required.

Verify permissions and provider registration

Follow the [current Foundry RBAC guidance \(https://learn.microsoft.com/azure/foundry/concepts/rbac-foundry\)](https://learn.microsoft.com/azure/foundry/concepts/rbac-foundry). **Foundry User/Owner/Account Owner/Project Manager** may still appear under their former **Azure AI User/Owner/Account Owner/Project Manager** names.

Use **Access control (IAM) → Check access** to inspect active assignments and Scope. Older UI versions may use View my access. The [actual subscription/access illustration in 01](#) demonstrates inspection of existing access, not adding roles.

Identity or task	Required scope to verify
New-environment operator	This bootstrap checks effective subscription permissions for group/deployment/resource writes and <code>Microsoft.Authorization/roleAssignments/write</code> . An already-authorized provisioning operator performs it.
Portal participant	Agent/evaluation operations in the assigned project; an administrator reviews the required project-scoped roles, such as Foundry User.
Policy-setup operator	Search schema/upload permissions, such as Search Service Contributor and Search Index Data Contributor, on this Search service.
Project managed identity	Search Index Data Reader on this Search service, model invocation, and telemetry submission to the assigned monitoring resource.
Search managed identity	Model invocation on the assigned Foundry resource.
Log reader	Read access to the assigned Application Insights/Log Analytics resources, not subscription-wide log access by default.

In Azure Portal → **Subscriptions → selected subscription → Resource providers**, verify **Registered** for `Microsoft.CognitiveServices`, `Microsoft.Search`, `Microsoft.OperationalInsights`, and `Microsoft.Insights`. An authorized subscription administrator registers missing providers through organizational procedures. Bootstrap never auto-registers providers or expands permissions.

Allow registration/role propagation, then rerun the same preflight. For continuing 403, compare identity, role scope, and network access. Do not grant new subscription-wide Owner, disable network or organizational policies, or transfer tokens as a shortcut.

Costs and operational boundaries

Evaluation invokes the Agent and Judge; optimization makes additional internal calls. Search and monitoring can continue to incur charges after the lab page closes. Distinguish measured tokens, cost estimates, and actual billing; unknown billing is not zero.

In Azure Portal → **Cost Management** → **Cost analysis**, select the subscription, resource group, and date range. Configure organizational budget alerts where needed, but do not present them as enforced caps. Record the cost/cleanup owner and end time in the handoff.

Hand off actual values without gaps

Copy this table into the private class record and fill it with actual values. “The environment is ready” is not a sufficient handoff.

Required item	Actual values or evidence
Identity	User, tenant ID, and subscription ID; never passwords or tokens
Responsibilities	Owners for provisioning, 06 ID lookup, 09 v2 creation/reevaluation, portal exercises, cost, and cleanup; one paid-submission owner per shared Agent
Local paths	Environment name, bootstrap config, runtime <code>.env</code> , <code>LAB_LANGUAGE</code> , and <code>LAB_ARTIFACTS_DIR</code> ; distinguish participant-specific configuration and which computer runs commands
Azure environment	Resource group, Foundry account/project, project endpoint, and Search service
Models	Actual names, product/version, SKU, and capacity for all four deployments
Policy search	Source/language, eight uploaded documents, knowledge-base/connection names, and <code>retrieval_verified</code>
Agent	Actual name, pinned v1, strict output, and a successful real Agent tool call
Dataset	JSONL path, 12 rows, SHA-256, registration name, and version
Evaluation	Relevance 4, TaskAdherence 1, Judge deployment, and query-only input; add actual evaluation/v1-run IDs after the first run
Authorization	Scope hash, expiry, budget, candidate limit, and retention

Required item	Actual values or evidence
Finish	Exact deletable group/objects, deletion approver, planned time, and shared-resource retention

Keep English and Korean source data and private lab records separate; do not relabel another language's results as a new execution.

Raw service files can contain account metadata, tokens or signed URLs. Keep those local; **the synthetic evaluation results themselves are not secret**. Publish only allowlisted fields, including failures, and exclude credentials—not inconvenient outcomes. Historical attempts stay in the private audit, not as additional current-version reports.

Verify deletion and retention

Use [step 10](#) as the exit checklist. Provisioning approval is not deletion approval, and workshop documentation does not authorize deleting another person's resources.

Target	Deletion or handoff check
Active evaluation/Optimizer jobs	Cancellation availability, actual terminal state, and unresolved job IDs
Locally owned objects	Exact <code>cleanup</code> plan and verified absence afterward
Portal-created objects	Individually inspect datasets, evaluations, conversations, Optimizer records, and empty Agent entries
Dedicated resource group	Review all inventory, obtain authorization, delete, and verify a successful <code>az group exists</code> returns <code>false</code>
Shared/external resources	Retained item, reason, owner, and retention end date
Billing/soft delete	Delayed charges, prior usage, and separate service retention/purge policy
Evidence	Keep outcomes and failures privately for the required period; do not delete <code>lab</code> first

Lab troubleshooting

[Ten participant steps](#) · [Operator authorization and handoff](#)

Start with the blocked step and check the symptom, action and completion criteria. Do not silently change data, evaluators or models, or submit the same operation under a new name to bypass an error.

Symptom	Where to look
Terms, installation, sign-in or file paths are unclear	Basic terms · Environment and sign-in
Authorization, permissions or model capacity blocks provisioning	Provisioning
Policy retrieval or Agent responses fail	Policies and Agent
Evaluation is unfinished or scores/reasons are missing	Evaluation and results
Optimize is unavailable or candidate selection is unclear	Optimizer
Deletion fails or resources must be retained	Cleanup and retention

How to use this guide

Keep timestamps, steps, commands, target IDs, error codes and actions in your private lab record. Preserve original responses and receipts rather than editing failures into successes. Do not append run-specific validation logs or results to reusable guides.

Environment, sign-in and local files

Symptom	Action	Completion criteria
<code>.venv/bin/activate</code> or <code>Activate.ps1</code> is missing	Create the virtual environment in the repository and follow the OS-specific setup order .	<code>python -m pip --version</code> and <code>python -m lab --help</code> work in the same environment.
<code>ModuleNotFoundError</code>	Confirm the intended environment and run <code>python -m pip install -r requirements.lock</code> with its Python.	Required imports and dataset checks work in that environment.

Symptom	Action	Completion criteria
PowerShell activation is blocked by policy	Use <code>.\.venv\Scripts\python.exe</code> without changing organizational policy.	Python runs without a policy change.
Portal and CLI accounts differ	These are separate sign-ins. Compare user, tenant and subscription from <code>az account show</code> with the portal; complete personal authentication yourself.	All three approved values match. Do not substitute another identity or copy tokens to bypass authentication.
<code>AzureCliCredential</code> fails with both tenant and subscription	Use the repository authentication path: verify the tenant first, then specify the subscription on the credential.	Calls use the intended subscription without dropping tenant checks.
Artifacts go to the other language's folder	Set <code>LAB_LANGUAGE</code> and <code>LAB_ARTIFACTS_DIR</code> before starting Python.	Dataset, configuration, workspace and results share the same language; existing records are not overwritten.

Provisioning, authorization and quota

Symptom	Action	Completion criteria
Plan reports <code>BLOCKED_AWAITING_APPROVAL</code>	Inspect <code>plan_status</code> and <code>mutations_performed</code> , then prepare the actual authorization .	Distinguish local planning from Azure provisioning.
Approval fails despite <code>approved: true</code>	Check hashes, models, approver, validity, budget, acknowledgments and <code>approval_reason</code> .	Readiness is <code>READY</code> and authorization is <code>READY_FOR_APPROVED_APPLY</code> .
Provider is not Registered	The subscription administrator registers it through organizational procedures.	Rerun the same preflight after all four providers are registered.
Contributor cannot assign roles	Resource creation and role assignment are separate permissions. Use the authorized provisioning owner.	Verify effective permissions; do not add subscription Owner to bypass the issue.
Model, SKU, version, quota or capacity error	Inspect the planned role requirements and preflight reason. A replacement needs a separate plan and authorization.	Confirm support and capacity, then verify actual calls.
Runtime <code>*_tpm</code> is <code>BLOCKED</code>	Check the real allocation using the TPM setup instructions .	All five TPM checks and overall preflight report <code>PASS</code> .
<code>unknown</code> resource after creation	Check whether this is verified default Smart Detection linked to your Application Insights, another workload, or changed receivers. Allow link propagation during creation.	Inspect the same <code>bootstrap_status</code> ; do not delete alerts or edit the manifest to bypass it.

Symptom	Action	Completion criteria
Apply times out or its outcome is unknown	Preserve the original config, manifest and deployment ID; read <code>bootstrap status</code> and the group's Deployments.	Wait for active work; do not resubmit an unknown outcome.
The local environment already exists	Resume the same lab with its original configuration. A different lab needs a new environment name and authorization.	Do not adopt or overwrite existing plans, settings or resources.

Use `--retry` only for a verified terminal owned failure and a separate retry allowance. `repair-dependencies` and `repair-trace-routing` also require matching failure evidence and scoped authorization. `APPLIED` describes infrastructure, not answer quality, production approval or confirmed trace ingestion.

Policy retrieval and Agent setup

Symptom	Action and completion criteria
Deployment exists but model calls fail	Catalog/deployment readiness and runtime support differ. Inspect the API error and intended role, then verify the actual response in 03 .
Setup remains <code>created_not_retrieval_tested</code>	Only creation is complete. Confirm <code>retrieval_verified</code> and actual references from <code>iq probe</code> .
Direct retrieval succeeds but the Agent tool returns 403	Check Search-read/model-call permissions and connection audience for the project's managed identity. User retrieval alone is insufficient.
<code>native-agent</code> reports a configuration mismatch	Inspect the receipt and remote version first. Equivalent MCP allowlist representations are normalized, but different tools, permissions, models or output settings remain blocked. Do not delete the receipt or create another version to bypass it.
Existing v2 has different instructions or model	Fixed-comparison protection is working. Compare the original owned environment, instructions and model snapshot; do not overwrite v2 or create v3.
Retrieval differs between languages	Check the selected language's source documents, index, analyzer and hashes. Renaming another language's index is not reuse.
Private-network or <code>PublicNetworkAccess=Disabled</code> failure	Use an approved VNet/VPN/execution environment and check DNS. Do not disable the firewall or private endpoint.

Foundry evaluation and safe resumption

Add run item-schema error

If the following message appears, inspect the existing evaluation's Agent-target data source:

```
Unable to create data source configuration from item schema
```

The `scripts/add_foundry_eval_run.py` [helper in 09](#) copies a completed baseline's remote data source and changes only the explicit candidate version. Confirm the same evaluation ID, registered dataset, Judge, thresholds and mappings, plus the actual version/instructions for all twelve items. Do not substitute new data or a local Judge.

IDs and duplicate submissions

Symptom	Action
Evaluation ID or run ID is unknown	Read <code>native-evals --name</code> . Evaluation uses <code>eval_...</code> ; the baseline is a completed version-one <code>evalrun_...</code> .
Several evaluations have the same name	Compare creation time, Agent, dataset and runs; do not select by name alone.
Helper exits with Still running	If the receipt has a run ID, resume collection with the identical command and <code>--out</code> .
Receipt has no run ID, or a remote run with the same name exists	Preserve the original and check submission acceptance and the actual run. Do not delete receipts, rename or automatically resubmit.
Changing the version clears its checkbox	Reselect explicit v1 and confirm one checked target.
Dataset preview has five rows	Compare the twelve-row source, registered version and complete result count; the preview alone is insufficient.
Scores or reasons are not visible	Read the response under <code>conversation_id → User view</code> ; reasons are in Detailed metrics result under <code>Relevance.reason</code> and <code>TaskAdherence.reason</code> .
Output items or scores are missing	Preserve failures/missing items and keep the denominator at twelve. Do not invent zero scores or passes.

Agent Optimizer and interpretation

Symptom	Action and completion criteria
No custom evaluators available	Use Custom only OFF or View built-in evaluators rather than creating a custom evaluator to bypass a filter.
Optimize or the selected model is unavailable	Check New Foundry, prompt Agent, project permissions, Preview availability and role-specific model support (https://learn.microsoft.com/azure/foundry/agents/concepts/agent-optimizer-overview#models).
A candidate changes model or tool descriptions	Select Instruction only and disable model comparison. Non-instruction changes are not a same-condition comparison.
Candidate embeds policy dates, numbers or evaluation answers	Preserve the original service output and remove embedded answers or reject the candidate. Record manual changes, reasons and provenance, then reevaluate separately.
Only the baseline is selected, with no different candidate	Record why v1 is retained and continue to 10. Do not invent v2 or repeat the job to manufacture improvement.
Internal ranking differs from reevaluation	Internal 0–1 ranking and separate evaluation scores/pass rates differ. Inspect actual answers, policy errors and regression too.
TaskAdherence is 1	That is a binary pass, not a one-out-of-five Relevance score.
Statistics are Inconclusive	A difference was not established; this is not evidence of significance or equivalence.
Quality improves but latency/tokens increase	Record the tradeoff too. Unknown billing is not zero cost.

Cleanup, retention and continuing costs

Symptom	Action and completion criteria
Retention is explicitly authorized	Do not run deletion commands. Record the inventory, cost owner and review date; keep ownership/configuration needed to reopen the environment.
Search charges remain after cleanup	<code>cleanup</code> handles recorded objects, not all hosting/models/groups/logs. Follow authorized cleanup or retention in 10 .
<code>.env</code> or workspace is missing	Compare the configuration, manifest and Azure inventory first. Do not fabricate ownership records to delete resources.
Monitoring in the group serves another group	Review shared action-group dependencies with the owner. A dedicated group does not authorize deleting shared resources.

Symptom	Action and completion criteria
Locks, permissions or dependencies block deletion	Read Locks, Activity log and the failed operation; request action from the owner. Do not remove organizational locks without authorization.
The group remains after a deletion request	Inspect asynchronous progress. Only a successful <code>az group exists</code> result of <code>false</code> verifies absence.
Costs appear after deletion	Check the usage period, billing delay and remaining resources in other groups. Earlier usage charges do not disappear.

Service-specific soft-delete retention and permanent deletion/purge require separate policy and authorization. See the [resource-group deletion documentation \(https://learn.microsoft.com/azure/azure-resource-manager/management/delete-resource-group\)](https://learn.microsoft.com/azure/azure-resource-manager/management/delete-resource-group).

Synthetic English evaluation data

[Dataset step](#) · [Operator handoff](#) · Korean data guide (file: docs/ko/data-guide.html)

One unchanged dataset for the comparison

Use **data/en/optimizer/dev.jsonl** (file: data/en/optimizer/dev.jsonl), exactly **12 JSONL records**, for Foundry Evaluation, Agent Optimizer, and reevaluation. In a new project, follow [04](#) to register lab-en-dev12 version 1. In a prepared project, reuse its operator-confirmed registration.

Keep its exact bytes, SHA-256, questions, references, and language fixed throughout the comparison. Keep Korean and English data separate rather than mixing them in one comparison.

Purpose and limits

Contoso Atlas Cloud is fictional. Its policies, inquiries and reference answers are synthetic—not private customer data or the terms of Microsoft, Azure or a real provider.

The purpose is **evaluation against your own business tasks and criteria**, not a public benchmark alone. Representative tasks expose policy-date boundaries, ambiguity, unsupported certainty and action claims. A 12-case development set demonstrates the learning loop; it does not certify production or independent generalization.

Preserve the original source corpus

data/en/cases.jsonl (file: data/en/cases.jsonl) contains 100 cases: train 56, validation 12, dev 12 and test 20. Only dev12 is used in this workshop. Other source subsets remain available but are not additional participant exercises.

Do not select only easy rows, duplicate questions, merge subsets, or replace references after seeing results. The eight policy documents (file: data/en/knowledge/documents.json) retain stable ATLAS-* IDs and effective dates. The participant or authorized operator prepares policy access with the [commands in 03](#).

Upload and response schema

Each JSONL line has exactly three columns:

Column	Type	Purpose
query	String	The only Agent input
context	String	Source-policy reference for supported evaluators and case review
ground_truth	JSON string	Structured reference answer, not a generation prompt

There is no prefilled response. Foundry invokes the pinned Agent to obtain one. Do not turn the file into a JSON array, rename columns, append references to the user message or precompute answers.

The reference answer and Agent response use exactly four keys:

Key	Contract
answer	Nonempty English string
citations	Unique stable IDs of supporting policy documents
route	answer, clarify, escalate or refuse
needs_human	Boolean; true exactly when route is escalate

The source response schema (file: `schemas/response.schema.json`) also checks routing consistency. Native structured output controls the JSON shape, not the factual truth of an answer. Numeric retrieval IDs are not policy citations.

The Agent cannot actually submit tickets, change subscriptions, approve credits or delete records. Classification as **escalate** does not mean a person was contacted. A request for unsupported certainty does not authorize fabricated facts.

Register once; pin the same version

Use **Foundry → Build → Evaluations → Create → Agent**, an explicit baseline version, **Individual turns / One time / Existing dataset**. Reselect the target checkbox if version selection clears it.

Choose the existing dev12 registration. In a fresh authorized project only, upload the exact language-specific file and record its new dataset version. The preview may show five rows; the file and evaluation scope still contain 12.

Leave the custom prompt override unset. Use `query → query` if a field-mapping screen appears. Service mappings are **Relevance** `response={{sample.output_text}}` and **TaskAdherence** `response={{sample.output_items}}`. These are not extra JSONL columns.

Keep **Relevance 1–5 / threshold 4**, **TaskAdherence binary 0/1 / pass 1**, and the same Luna Judge. Selecting two evaluators does not mean both consume every reference field or certify every business rule.

Compare complete evidence

Compare the questions, responses, evaluator scores, and reasons for all 12 cases across same-model v1 and v2. Record your actual run IDs and dataset hash to identify the comparison.

The helper checks matching questions/references and actual per-item Agent version/instructions. Missing/error rows cannot become successful rows or disappear from the denominator. JSON/route checks are supplementary validation, **not a replacement local Judge**.

Agent Optimizer's internal ranking is distinct from the separate managed run's means and pass counts. Review candidate instructions and preserve model/tools/reasoning/schema. Compare one reviewed v2 without continually adding releases. Operators may use separate drafts after verifying support.

Before sharing results, remove credentials, signed URLs, and private account details. Keep every case, including errors and missing results.

`cleanup` does not remove every registered dataset. Follow [10](#) to retain or delete registrations/evaluation records according to dedicated or shared-project ownership.