

Foundry Learning Loop Lab · Complete Print Edition

Follow the six participant steps for Foundry Evaluation, Agent Optimizer, and same-criteria comparison. Operator, facilitator, data, and verification references are included but are not extra labs. Cross-document links stay within this book; file references use package paths.

As of 2026-10-01

Included documents

1. [Participant guide](#)
2. [Appendix 1 · Facilitator guide](#)
3. [Appendix 2 · Operator setup](#)
4. [Appendix 3 · Verification and sources](#)
5. [Appendix 4 · Data guide](#)

Participant guide

Great agents start with evaluation · v1

CONTOSO ATLAS CLOUD · ONE LEARNING PATH

Evaluate your business cases, improve instructions with Agent Optimizer, and verify a better v2 using the same criteria.

01 [Prepare the dataset](#)

02 [Select evaluation criteria](#)

03 [Run Foundry Evaluation](#)

04 [Read scores and reasons](#)

05 [Improve instructions with Agent Optimizer](#)

06 [Reevaluate and compare v1/v2](#)

Use **your own representative tasks and business criteria**, not a public benchmark alone. Synthetic Contoso policies demonstrate **evaluate → learn → improve → reevaluate** without publishing private customer data.

The [operator](#) prepares the isolated project, Agent, policy tool, deployments and cost authorization before class. Infrastructure, Judge calibration and additional governance are not extra participant exercises.

Version meaning: v1/v2 are complete Foundry Agent versions. Here both use `gpt-6-sol`; only instructions change. Do not weaken v1 or promise an improvement before measuring it. Develop candidates as drafts and keep released versions at **v1 and v2**.

How to read the pictures: These real English portal captures illustrate control locations from an earlier UI session. Their model names, versions and scores are not the current Sol verification. The [latest v2 report](#) is the authoritative measured comparison and includes every actual question,

response, score and reason. Korean guide text does not mean a Korean run was executed.

01. Prepare the dataset

Use **data/en/optimizer/dev.jsonl** (file: `data/en/optimizer/dev.jsonl`) unchanged: **12 JSONL rows**, not a JSON array. The existing registration is `contoso-eval-en-dev12` **version 1**. Do not reupload it in the prepared project.

Column	Type	Use
<code>query</code>	String	The only Agent input
<code>context</code>	String	Policy reference for supported evaluators and case review
<code>ground_truth</code>	JSON string	Structured reference answer, never a generation prompt

Record all 12 rows, the registration/version and file SHA-256. Keep the bytes identical through baseline, optimization and reevaluation. The source has other reserved subsets; they are not part of this workshop.

The Agent returns exactly `answer`, `citations`, `route`, `needs_human`. `answer` is English, citations are supporting policy IDs, `route` is `answer/clarify/escalate/refuse`, and the Boolean `needs_human` is true only for `escalate`. The Agent cannot actually submit, refund, delete or grant access.

Open **Foundry New experience → Build → Evaluations → Create → Create new evaluation**. For the prepared English target select **Agent** → `contoso-eval-en-sol`, pin the baseline explicitly to **v1**, and confirm one checked target. **Pin currently latest** is safe only while latest really is v1; selecting it can clear the checkbox.

Choose **Individual turns**, **One time**, and **Existing dataset**. Select the registered dev12. In a fresh authorized project only, use Upload new dataset for the exact file. The preview may show five rows; the dataset still has 12.

Microsoft Foundry / proj-fev11-dea3cec5

Search with AI (96K)

New Foundry

Home

Discover

Build

Operate

Manage

Docs

Create

Agents

Models

Services

Tools

Knowledge

Memory

Guardrails

Data

Optimize

Evaluations

Fine-tune

Create new evaluation

Target: Agent

Scope

Frequency

Data

Select an existing dataset or generate a synthetic one to evaluate your model or agent's behavior.

Dataset source

Synthetic generation

Choose a deployment and questions will be automatically generated for each target when you submit the evaluation

Existing dataset

Fully customizable advanced option that requires data mapping and prompt setting in evaluators on your own

Benchmarks

Industry-standard benchmarks with built-in evaluators

Existing traces

Evaluate based on traces already generated by this agent

Back

Next

Configure agents

Criteria

Review

Select your file

Download data sample

Upload new dataset

Search by file name

Name	Version	Created by	Created on
eval-data-2026-09-30_144335_a656c.UTC	1		9/30/26, 11:43:35 PM
contoso-eval-en-dev12	1		9/30/26, 11:20:29 PM

1-2 of 2

Dataset preview (Top 5 rows)

context	ground_truth	query
[ATLAS-SUB-001] Subscription-3. Canceling a subscription stops the next automatic renewal; access continues until the end timestamp of the current paid period.	("answer":"By 'reduce', do you mean stopping automatic renewal, or con on a lower plan from the next renew Those are different choices, so I ne clarify your intention first. This conv has not changed the	Next month's cost is a concern. Please reduce my subscription. I haven't said yet whether I'll stop using it entirely or keep using a lower tier.
[ATLAS-REF-001] Refund-2. A first purchase of a prepaid monthly subscription made before 2026-09-01 00:00:00 KST retains the previous 7 x 24-hour application window, including	("answer":"The instruction in the att notice does not change the policy. / 28 purchase retains the previous se window, whose final timestamp wa: September 4 at 08:00 KST. Septembt past the deadline; the month of the	My first monthly payment was on August 28, 2026 at 08:00 KST. I haven't used it, and I'm contacting you on September 6. The notice I attached says 'Everyone contacting us in
[ATLAS-REF-001] Refund-1. A first purchase of a prepaid monthly subscription made at or after 2026-09-01 00:00:00 KST is	("answer":"The standard conditions both no paid-credit use and no paid production jobs. If even one paid pr job was run, the situation you descr	I used the service for just one day after my first monthly payment on September 8, 2026. Paid-credit usage is zero, but I

Find the dataset controls.

Select the current operator-provided registration.

The preview is not proof of the total row count.

View full-size image (file: web/assets/portal/en/15-evaluation-dataset.png)

Completion signal: The draft targets explicit v1 and the unchanged 12-row dataset; count, version and hash are recorded. [Data contract](#).

02. Select evaluation criteria

In **Configure agents**, leave the custom prompt override unset. Input is `{{item.query}}` **only**; never append context or `ground_truth`. If field mapping appears, use `query → query`.

Retain exactly these two managed evaluators:

Evaluator	Meaning	Setting
Relevance	Addresses the question, 1–5	Threshold 4
TaskAdherence	Follows the task, binary 0/1 Pass/Fail	Pass 1 , not threshold 4

Choose the explicit Judge deployment and preserve service-generated mappings: Relevance `response={{sample.output_text}}`; TaskAdherence `response={{sample.output_items}}`. Read the definition's Raw JSON rather than substituting old UI bindings.

Role	Model/version	Deployment
Agent	gpt-6-sol / 2026-09-22	lab-agent-sol-dea3cec5
Evaluation Judge	gpt-6-luna / 2026-09-22	lab-judge-luna-dea3cec5
Optimizer generator	gpt-5.5 / 2026-04-24	lab-planner-dea3cec5

Sol's actual Agent/tool invocation was verified. Model catalog visibility alone is insufficient, and the [Optimizer support list](#) (<https://learn.microsoft.com/azure/foundry/agents/concepts/agent-optimizer-overview#models>) is role-specific. Catalog evaluator version links do not prove that a private service rubric version is pinned; preserve the actual definition and disclose that limit.

Microsoft Foundry / proj-fev11-dea3cec5

Search with AI (🔍)

New Foundry

Home

Discover

Build

Operate

Manage

Docs

Create

Agents

Models

Services

Tools

Knowledge

Memory

Guardrails

Data

Optimize

Evaluations

Fine-tune

← Create new evaluation

Target: Agent

Scope

Frequency

Data

Configure agents

Criteria

Review

Mapped fields from your dataset:

query: {{item.query}} response: {{sample.output_text}} context: {{item.context}} ground_truth: {{item.ground_truth}}

tool_calls: {{sample.tool_calls}} tool_definitions: {{sample.tool_definitions}}

Instructions

What evaluators are not auto-selected here?

Evaluators from Azure OpenAI come with text-based field mappings to give you more flexibility beyond one to one mapping. You will need to specify field relationship when you construct your system, user and assistant prompts. Don't worry, we have templates so you won't have to start from scratch.

Choose what criteria you want to evaluate — quality, risk and safety, and more.

Judge model *

Model: lab-judge-luna-dea3cec5 Global Standard deployment

We've auto-suggested the following 2 evaluators

Agents (1) Remove all

TaskAdherence X

Quality (1) Remove all

Relevance X

Add new evaluator

Back

Next

Read each metric's own scale.

Keep Relevance threshold 4, binary TaskAdherence pass 1, and the current Luna Judge.

View full-size image (file: web/assets/portal/en/16-evaluation-criteria.png)

Completion signal: The two evaluator scales, thresholds, mappings and actual Judge deployment are recorded. The saved remote definition, not a stale local environment default, determines the Judge.

03. Run Foundry Evaluation

Review **v1 + original dev12 + query-only input + two evaluators + Luna Judge**. The current prepared comparison is `contoso-en-sol-learning-loop`. Use the existing completed baseline when following the recorded rehearsal, not a second submission.

For a new authorized class, **Review → Submit** once. Keep the actual evaluation ID and run ID. Wait for completion and account for all 12 cases, including errors or missing results. A configured draft, HTTP creation receipt or successful model smoke is not a completed evaluation.

Microsoft Foundry / proj-fev11-dea3cec5

Search with AI (🔍)

New Foundry Home Discover Build Operate Manage Docs

Create

- Agents
- Models
- Services
- Tools
- Knowledge
- Memory
- Guardrails
- Data

Optimize

- Evaluations
- Fine-tune

Create new evaluation

Target: Agent

Scope

Frequency

Data

Configure agents

Criteria

Review

Add a name and description to this evaluation run.

Evaluation name *

contoso-en-learning-loop

Back Submit

Summary

Agent evaluation

Targets

contoso-eval-en: v1

Agent

Edit

Scope

Individual turns

Scope

Edit

Frequency

One time

Frequency

Edit

Dataset

contoso-eval-en-dev12

Version 1

Existing dataset

Edit

Evaluators

TaskAdherence

Agents

Edit

Relevance

Quality

Edit

Review before Submit.

This picture illustrates the controls, not the current Sol run. Confirm the actual pinned version and saved evaluation contract. View full-size image (file: `web/assets/portal/en/17-evaluation-review.png`)

Completion signal: The real Foundry run is Completed and exposes 12 output items. A failed or incomplete run stays failed/incomplete. No local custom Judge substitutes for this managed Evaluation.

04. Read scores and reasons

Read summary counts and **Detailed metrics result**, especially `Relevance.reason` and `TaskAdherence.reason`. `conversation_id → User view` opens the actual question/answer, not an inline Judge-reason panel.

Select a failing or lowest-score case and a good case. Connect **question → actual response → score/reason → policy and parsed reference answer**. If no failure exists, say so; never weaken the baseline to produce one.

Microsoft Foundry / proj-fev11-dea3cec5

Search with AI (英K)

New Foundry

Home

Discover

Build

Operate

Manage

Docs

Create

Agents

Models

Services

Tools

Knowledge

Memory

Guardrails

Data

Optimize

Evaluations

Fine-tune

← contoso-eval-en

Make recurring

Delete

Run details

Status

Completed

Create time

10/1/2026, 12:21:37 AM

Duration

3m 59.84s

Created by

Agent ID

contoso-eval-en:1

See all properties

Raw JSON

User logs

Download user logs

Results

Download results

Overall metric results

P50 latency	P95 latency	TaskAdherence	Relevance
5.9s	8.8s	100% 12 / 12	83% 10 / 12

Detailed metrics result

Analyze Results

Index	query	conversation_id	trace_id	Relevance
0	Next month's cost is a concern. Please reduce my subscription. I...	conv_07e180556bd61ca800UPWqFCcJqavhmEPmHcTJBXGDHbqHYFz	-	Fail: 3
1	My first monthly payment was on August 28, 2026 at 08:00 KST. I...	conv_0bb9168bd3bf11ec00yPvdLhJy3IE1IHfKA6uCYJS9YCcWpLx	-	Pass: 5
2	I used the service for just one day after my first monthly payment on...	conv_05f494a8836c0111004DizSunhTEdBq7cEwzgpPeP3f4M5UKYcT	-	Pass: 4
3	I typed only a title on the Support > Contact support screen and left...	conv_09e82a7f1e3db55ea00h7GlpAlJkG8N9Ngm0JrFoOUAXyEm0k5r	-	Pass: 4
4	There is an unfamiliar access event on the administrator's audit screen...	conv_05740a8abb5d88ad00G5SaTwellR5t4uMTWWfIdYBV0kq4lgQVM	-	Pass: 5
5	Assume I paid the same monthly base fee in two months, with...	conv_098f970e5fa2d1da00zhmGbfVSMhdOfzI7lLUcjB4PGVLH8UFJ	-	Pass: 5
6	I'm on paid Pro, but only beta webhook delivery failed; the...	conv_099677dc019f239900fUJqcz6Qpnwkk97SDxfp0SUv8pXxquQ	-	Pass: 5
7	Since there is a 99.9% monthly commitment, can I write in a...	conv_0a857d1c7aec9b6800pluwLFZNNSVYwbrpuCehFCbQhFyecUuo	-	Pass: 5
8	I learned that the refund window includes the first five business days...	conv_0eb63b9bf392360b00tONh20KhDm3f7uwl1mCf4HefDQVGR7NA	-	Pass: 5

Find scores and reasons.

This earlier UI image is not current verification. Read the current run IDs and full case evidence in the latest report. View full-size image (file: web/assets/portal/en/18-evaluation-results.png)

Relevance 4/5 is not 80% accuracy. TaskAdherence 1 means Pass, not a low five-point score. Missing values are not zeros or successful rows. Generic evaluators do not certify every business rule.

conv_07e180556bd61ca...

Conversation ID: conv_07e180556bd61c...

Trajectories

User view

Graph view

3 spans

1 chat call

1.8s

1,1K

Next month's cost is a concern. Please reduce my subscription. I haven't said yet whether I'll stop using it entirely or keep using a lower tier.

{
 "answer": "I understand your concern about the upcoming cost. You can reduce your subscription to a lower tier or make other changes to your subscription plan to better fit your needs. Please let me know if you'd like help with specific options or if you want assistance with pausing or canceling your subscription entirely. I'm here to support you in managing your subscription in the best way possible.",
 "citations": [],
 "route": "answer",
 "needs_human": false
}

Read the response separately.

Return to Detailed metrics result for the reasons and compare the source policy. Do not copy this older answer as current evidence. View full-size image (file: web/assets/portal/en/19-evaluation-case.png)

Watch for wrong date boundaries, unnecessary assumptions, numeric retrieval IDs instead of document citations, incorrect human routing, invented actions and unsupported certainty. A truthful statement of uncertainty must not be replaced with a fabricated fact to satisfy a Judge.

Discuss: State the run ID, each metric's scale/pass count, a concrete problematic answer and the instruction behavior you want to improve.

Completion signal: You can explain a real weakness and a supported improvement hypothesis, not just an average.

05. Improve instructions with Agent Optimizer

Open **Build** → **Agents** → `contoso-eval-en-sol` → **Optimize Preview** → **Agent**, not Cost. For a fresh job use **Create an optimization run**.

Setting	Choice
Agent version	Explicit baseline 1
Choose targets	Instruction only ; Model and Tool description off
Max candidates	Operator-approved bound; this run uses at most 2, not more releases
Optimization model	lab-planner-dea3cec5 / gpt-5.5
Evaluation model	lab-judge-luna-dea3cec5 / gpt-6-luna
Dataset	Same registered English dev12 version 1
Criteria	Relevance 4; TaskAdherence binary pass 1

Separate the model roles. Use the current operator-provided Sol target, Luna Judge and supported generator; this capture only locates the controls. View full-size image (file: [web/assets/portal/en/07-optimizer-target.png](#))

If Criteria shows **No custom evaluators available**, switch **Custom only OFF** or select **View built-in evaluators**. Open each built-in row, set its correct threshold and Apply. Do not create a custom scorer to bypass a filter.

Create an optimization run

Target

Dataset

Criteria

Review

Choose an existing dataset to test each candidate against. Optimization runs the same dataset across every candidate so results are comparable.

Generate data

Select dataset and criteria

Back

Next

Select dataset

Search by file name

Create dataset from traces

Upload dataset

Name	Version	Created by	Created on
eval-data-2026-09-30_144335_a656c_UTC	1		9/30/26, 11:43:35 PM
contoso-eval-en-dev12	1		9/30/26, 11:20:29 PM

1-2 of 2

PrevNext

Preview

Dataset

contoso-eval-en-dev12

Fields

context

ground_truth

query

Dataset preview (Top 5 rows)


context	ground_truth	query
[ATLAS-SUB-001] Subscription-3. Canceling a subscription stops the next automatic renewal; access continues until the end timestamp of the current	{"answer": "By 'reduce', do you m stopping automatic renewal, or on a lower plan from the next rer Those are different choices, so I clarify your intention first. This c has not changed the	Next month's cost is a concern. Please reduce my subscription. I haven't said yet whether I'll stop using it entirely or keep using a lower tier.
[ATLAS-REF-001] Refund-2. A first purchase of a prepaid monthly subscription made before 2026-09-01 00:00:00 KST retains the previous 7 x 24-hour application window,	{"answer": "The instruction in the notice does not change the polici 28 purchase retains the previou window, whose final timestamp September 4 at 08:00 KST. Septe past the deadline; the month of	My first monthly payment was on August 28, 2026 at 08:00 KST. I haven't used it, and I'm contacting you on September 6. The notice I attached says 'Everyone
[ATLAS-REF-001] Refund-1. A first purchase of a prepaid monthly subscription made at or after 2026-09-01	{"answer": "The standard conditi both no paid-credit use and no p production jobs. If even one paic job was run, the situation you de	I used the service for just one day after my first monthly payment on September 8, 2026. Paid-credit usage is

Reuse the data. Same registered version and all 12 cases; synthetic regeneration or a changed upload would change the experiment. View full-size image (file: web/assets/portal/en/08-optimizer-dataset.png)

Submit once within the approved cost scope; wait at most 60 minutes and then retain the actual state. Reuse the existing job ID on a resumed class. Optimization includes multiple internal calls.

Inspect original/candidate scores and **View changes**. The internal **0–1 ranking** is not the separate managed Evaluation mean or pass percentage. Keep model, tools, reasoning and output schema unchanged; an empty function-tool export does not authorize removal of the MCP policy connection.

Current run: the managed Optimizer retained the strong v1 baseline. Its generated candidates were not automatically promoted. The operator used the observed failures to refine the instructions and verified that reviewed draft separately. The current v2 source is therefore **operator-reviewed after Agent Optimizer**, not the service's automatically recommended candidate.



Read the candidate, not just the ranking. This image is an older UI example. The current report identifies the actual job and selected candidate. View full-size image (file: web/assets/portal/en/09-optimizer-results.png)

Changes: candidate_1 ★ Best candidate

Comparing candidate_1 against baseline configuration.

Baseline

```
1 {
2   "agentName": "contoso-eval-en",
3   "agentVersion": "1",
4   "system_prompt": "You are a customer support assistant for
Contoso Atlas Cloud. Answer in English using the customer's
inquiry and the supplied policy context.\nFor a policy
question, if the knowledge_base_retrieve knowledge tool is
connected, use it to check the evidence before answering. If
no tool or context is available, state the limitation rather
than inventing policy details.\nThe response must be a single
JSON object containing only the following four fields, with
no other text or Markdown.\n{\n  \"answer\": \"Answer text\",
  \"citations\": [\n    \"Document identifier\",
  ],
  \"route\": \"answer\",
  \"needs_human\": false\n}\nroute must be one of
answer, clarify, escalate, or refuse. citations is an array
of strings. needs_human is true only when route is escalate
and false otherwise.\n",
5   "skills": [],
6   "tools": []
7 }
```

Candidate: candidate_1

```
1 {
2   "agentName": "contoso-eval-en",
3   "agentVersion": "1",
4   "system_prompt": "You are a customer support assistant for
Contoso Atlas Cloud. Answer in English using the customer's
inquiry, any supplied policy context, and the evidence rules
below.\n\nYour final response must be a single valid JSON
object and nothing else. Do not include Markdown, code
fences, tool traces, arrays outside the JSON object, or
citation footnote markup inside the answer text.\n\nRequired
JSON shape:\n{\n  \"answer\": \"Answer text\",
  \"citations\": [\n    \"Document identifier\",
  ],
  \"route\": \"answer\",
  \"needs_human\": false\n}\n\nField rules:\n- Include exactly
these four fields: answer, citations, route, needs_human.\n-
route must be one of: answer, clarify, escalate, refuse.\n-
needs_human must be true only when route is \"escalate\"; it
must be false for answer, clarify, and refuse.\n- citations
must be an array of document identifiers only, such as
\"ATLAS-SLA-001\", \"ATLAS-SUB-001\", or \"ATLAS-REF-001\".
Do not put URLs, inline citation markers, or quoted snippets
in citations.\n- The answer field should be concise but
complete, and should not contain unsupported policy
claims.\n\nEvidence and tool-use rules:\n- For any policy
question, if the knowledge_base_retrieve tool is connected,
call it before answering.\n- Use the retrieved policy
content, supplied policy context, and the embedded policy
```

1 of 1 Close

Review the actual current diff. Reject invented policy, unsupported certainty or configuration changes. Longer instructions are not automatically better. View full-size image (file: web/assets/portal/en/10-optimizer-changes.png)

Do not continually create v3, v4 and later releases. The operator tests the candidate as an explicitly pinned draft. **Promote/create v2 only after the real same-criteria check demonstrates improvement without quality regressions.** Promotion is not permission to publish to production or change production traffic.

Discuss: Which instruction behavior changed, what should improve, and what could regress? Identify the actual candidate rather than copying an old screenshot.

Completion signal: A real managed Optimizer result and reviewed candidate exist. Independent reevaluation is still required.

06. Reevaluate and compare v1/v2

Use the **same Foundry evaluation definition**, same dev12 and Luna Judge. In this portal, Add run previously failed with `Unable to create data source configuration from item schema`. The official Azure AI Projects/OpenAI Evals helper submits a real managed run; it is not local scoring.

With the operator's prepared venv and verified `az login`, replace these four placeholders with **your own** project and baseline IDs:

```
export AZURE_AI_PROJECT_ENDPOINT="OPERATOR_PROJECT_ENDPOINT"
export AZURE_SUBSCRIPTION_ID="OPERATOR_SUBSCRIPTION_ID"
export FOUNDRY_EVALUATION_ID="YOUR_EVALUATION_ID"
export FOUNDRY_BASELINE_RUN_ID="YOUR_BASELINE_RUN_ID"
```

The operator may validate an explicit draft before releasing v2. Once v2 is released, use:

```
python scripts/add_foundry_eval_run.py --endpoint "$AZURE_AI_PROJECT_ENDPOINT" --subscription
"$AZURE_SUBSCRIPTION_ID" --evaluation "$FOUNDRY_EVALUATION_ID" --baseline "$FOUNDRY_BASELINE_RUN_ID" -
-version 2 --name candidate-v2 --out .lab/foundry-evaluations/candidate-v2.json
```

The receipt protects against duplicates. Repeat the identical command/path only to collect the same run. The helper also blocks an existing remote name, verifies remote thresholds/Judge/mappings, and confirms the actual version and system instructions in every result row.

In **Evaluation runs**, select both rows and **Compare runs**. Explicitly choose **v1 as Baseline**; selection order must not reverse the comparison.

Microsoft Foundry / proj-fev11-dea3cec5

Search with AI (K)

New Foundry Home Discover Build Operate Manage Docs

Create

- Agents
- Models
- Services
- Tools
- Knowledge
- Memory
- Guardrails
- Data

Optimize

- Evaluations
- Fine-tune

← Statistically compare results

Total runs selected 2Analysis method PairedTTest

Baseline * contoso-eval-en

Name	contoso-eval-en	candidate-v2
TaskAdherence	1.00 Baseline	1.00 Inconclusive
Relevance	4.42 Baseline	4.33 Inconclusive

1-2 of 2

PrevNext

Confirm the comparison direction. This capture locates the Baseline control; it is not the current Sol measurement. Use the latest report's exact two run IDs. View full-size image (file: [web/assets/portal/en/20-evaluation-comparison.png](#))

The current [v2 verification](#) publishes all 12 case responses, scores and reasons alongside the v1 control. Require no decline in all-criteria passes or either metric's pass count/mean, and at least one strict measured improvement. Check factual truth, routing and response format too. Report latency, tokens and the native statistical conclusion separately.

Explain the result: Identify the actual gain, unchanged criteria, any regression, and the remaining uncertainty. Observed improvement is not a guarantee that every future stochastic run will improve.

Completion signal: The latest v2 has a complete, honest same-model comparison. There is one current report, not accumulating release numbers. Production approval and independent generalization are separate and are not granted by this dev12 exercise.

Finish: Give the operator the actual run/job IDs, data hash and measured decision. Preserve raw receipts locally; publish synthetic answers and evaluator reasons without credentials, cookies, signed URLs or private account details. [Facilitator recovery](#) · [Data contract](#).

Facilitator guide · improve the measured result

[Six participant steps](#) · [Operator handoff](#) · [Latest v2 verification](#)

Teach **evaluate** → **learn** → **improve** → **reevaluate** using business-specific tasks and real managed results. Do not manufacture a weak baseline or promise that every future evaluation will improve.

Before class

The operator provides `contoso-eval-en-sol`, Sol model deployment, the existing read-only policy tool, the registered English dev12, a Luna Judge and a supported `gpt-5.5` Optimizer deployment. Verify model roles and identity using the [operator guide](#).

Both v1 and v2 are full immutable Agent versions. Their model, tools, reasoning and strict output schema match; only instructions differ. V1 already has useful grounding, routing, date and uncertainty safeguards. Candidate development uses drafts, not an ever-growing sequence of released versions.

Keep English and Korean data/Agents separate. The Korean guide explains the procedure; the current published measurements are **English-data results only**.

Keep exactly six participant steps

Step	Required action	Completion evidence
01 Dataset	Inspect the unchanged three-column JSONL and existing registration	All 12 rows, version and SHA-256
02 Criteria	Relevance 1–5 / threshold 4; TaskAdherence binary 0/1 / pass 1; Luna Judge	Saved remote definition and query-only Agent input
03 Baseline	Run actual Foundry Evaluation against pinned Sol v1	Completed run ID and all output items
04 Analysis	Read actual answers, metric scores, reasons and source policy	A concrete improvement hypothesis, including failures
05 Optimization	Actual instruction-only Agent Optimizer; inspect the candidate diff	Real job/candidate IDs, unchanged non-instruction configuration

Step	Required action	Completion evidence
06 Reevaluation	Same definition/data/Judge, explicitly pinned candidate; Compare runs	Matched item-level evidence, improvement checks and honest decision

Use **Individual turns / One time / Existing dataset**. The upload preview can show only five rows; the experiment still contains 12. Version selection can clear the target checkbox: reselect it and confirm one target.

Discuss the actual cases

Ask participants to connect **question → actual answer → evaluator score and reason → policy/reference answer**. `conversation_id` → `User` view shows the conversation; the evaluator reasons belong in **Detailed metrics result**, such as `Relevance.reason` and `TaskAdherence.reason`.

Discuss at least one weak case and one strong case. If no failure is observed, use the lowest-scoring case and say no failure was found. Never weaken v1, omit a difficult row, change a reference answer or substitute local scores.

The latest public record includes all 12 synthetic questions, both actual responses and the managed evaluator reasons. Those results are not credentials. Exclude tokens, cookies, signed URLs and private account metadata, not unsuccessful outcomes.

Interpret improvements correctly

Relevance 4/5 is not 80% accuracy. TaskAdherence 1 is Pass, not a poor one-out-of-five score. Report each evaluator separately, all-criteria passes, errors and coverage.

`scripts/compare_foundry_eval.py` requires complete matched cases, identical data/model/tools/settings, per-item version/instruction attestation, no lower pass counts or means, and at least one strict measured quality improvement. Malformed responses remain failures; the helper never repairs them into a pass.

Read answers for unsupported certainty, incorrect policy arithmetic, action claims and wrong routing even if the generic evaluator passed them. A request demanding an impossible certainty must not be satisfied with invented facts.

Optimizer ranking is internal selection evidence, not the direct managed reevaluation mean or pass rate. A higher ranking alone does not establish a better v2.

Recover without changing the experiment

Symptom	Response
Wrong identity, 401/403	Confirm the approved account/project; do not widen shared permissions
Catalog or deployment succeeds but Agent fails	Record the actual API/runtime error; verify the intended role
Portal Add run item-schema error	Use the official SDK helper with the existing baseline data source, not a different scorer
No custom evaluators available	Switch Custom only OFF or select View built-in evaluators
Duplicate local receipt or existing remote run	Retrieve that same run; do not rename it to force another submission
Candidate changes model/tools/reasoning/schema	Reject an instruction-only comparison claim
Draft becomes a numbered release	Stop; draft support is not enabled as expected
Existing released v2 differs from local source	Do not overwrite it or create v3; preserve the release and inspect the mismatch
Missing output or metrics	Report incomplete execution; do not reduce the denominator
Statistical comparison is Inconclusive	Report that result; do not claim significance or equivalence

Relevance maps to `{{sample.output_text}}`; TaskAdherence maps to `{{sample.output_items}}`. Do not overwrite service-generated bindings or feed `ground_truth` into the Agent. The actual remote Judge overrides any stale local default.

Correct common misconceptions

Claim	Correction
"V2 is a prompt file number only."	It is a complete Agent version; this experiment changes only its instructions
"A more powerful model makes v2 better than old-model v1."	Both versions must use the same model for the instruction comparison
"A better v2 can be guaranteed before execution."	Only observed improvement can be established; future stochastic results are not guaranteed
"Keep only successful rows."	Keep every case, including failures and missing/error outcomes

Claim	Correction
"The source schema certifies correct business decisions."	It controls output structure, not policy truth or proper routing
"Public evidence must include cookies and raw logs."	Publish actual synthetic responses and reasons through an allowlist, not authentication data

Finish with one current report

Publish **the latest v2 verification and its frozen v1 control**, not a sequence of obsolete current-version reports. Keep earlier raw receipts for audit without relabeling them as this pair.

Report the same-model result, native statistical comparison, any latency/token tradeoffs and the remaining limits of a reused synthetic dev set. Candidate activation is a separate operator decision; no production publication is part of this workshop.

Operator prerequisites · the fixed Sol v1/v2 pair

[Participant guide](#) · [Facilitator](#) · [Latest v2 verification](#)

Prepare the environment before class. The participant exercise is **dataset → evaluators → Foundry Evaluation → case analysis → Agent Optimizer → reevaluation**, not infrastructure construction.

Use the existing isolated project

Use the operator-owned North Central US project, its existing policy connection, and the approved synthetic Contoso data. Confirm the actual account, tenant and subscription for the CLI, SDK and portal separately. Do not recreate resources, reuse production traffic or broaden shared permissions because a client is signed in as the wrong identity.

The English target is `contoso-eval-en-sol`. The original English and Korean source corpora remain separate. A Korean guide is not proof of a new Korean service run.

Microsoft Azure

Search resources, services, and docs (G+I)

Copilot

Home

rg-foundry-eval-v11-20260930-dea3cec5

Resource group

Generate Bicep code to duplicate this resource group.

Export resource groups using Bicep or Terraform +1

Search

Create

Manage view

Delete resource group

Refresh

Export to CSV

Open query

Assign tags

Move

Delete

Export template

Group by none

Overview

Activity log

Access control (IAM)

Tags

Resource visualizer

Events

Settings

Cost Management

Monitoring

Automation

Help

Essentials

Subscription (move)

Subscription ID

Tags (edit)

lab : foundry-eval-v11

lab_environment : lab-en-20260930

Show more (2)

Deployments

: 1 Failed, 2 Succeeded

Location

: North Central US

JSON View

Resources

Recommendations

Filter for any field...

Type equals all

Location equals all

Add filter

Name ↑	Type	Location
ai-fev11-20260930-dea3cec5	Foundry	North Central US
proj-fev11-dea3cec5 (ai-fev11-20260930-dea3cec5/proj-fev11-dea3cec5)	Foundry project	North Central US
appi-fev11-20260930-dea3cec5	Application Insights	North Central US
Failure Anomalies - appi-fev11-20260930-dea3cec5	Smart detector alert rule	Global
log-fev11-20260930-dea3cec5	Log Analytics workspace	North Central US
srch-fev11-20260930-dea3cec5	Search service (Foundry IQ)	North Central US

Showing 1 - 6 of 6. Display count: 20

Give feedback

Resource-group orientation only.

This earlier portal picture is not the current v2 result. Verify the actual project and current state from the operator handoff. View full-size image (file: web/assets/portal/en/00-resource-group.png)

Verify the three model roles

Role	Actual deployment	Model/version
Agent	lab-agent-sol-dea3cec5	gpt-6-sol / 2026-09-22
Managed Evaluation Judge	lab-judge-luna-dea3cec5	gpt-6-luna / 2026-09-22
Agent Optimizer instruction generator	lab-planner-dea3cec5	gpt-5.5 / 2026-04-24

Sol was verified through an actual pinned Foundry prompt-agent invocation and successful knowledge_base_retrieve call, not just a model catalog entry. The generated response used the four-field JSON contract. Both versions keep the same model, read-only tools, reasoning setting and structured-output schema.

The Optimizer model is a separate role with its own supported-model list (https://learn.microsoft.com/azure/foundry/agents/concepts/agent-optimizer-overview#models). Do not assume that a model usable as an Agent or Judge is automatically an allowed optimization model.

Microsoft Foundry / proj-fev11-dea3cec5

Search (36K)

Create

Agents

Models

Services

Tools

Knowledge

Guardrails

Data

Optimize

Evaluations

Fine-tune

Models

Deployments

Batch jobs

Serverless deployments

Managed compute deployments

Search

PTU Calculator

Deploy

Name	Model	Version	Deployment status	Deployment type	Created on	Created by
lab-judge-luna-dea3cec5	gpt-6-luna	2026-09-22	Succeeded	Global Standard	9/30/26, 11:19:02 ...	
lab-agent-luna-dea3cec5	gpt-6-luna	2026-09-22	Succeeded	Global Standard	9/30/26, 11:18:59 ...	
lab-embedding-dea3cec5	text-embedding-3...	1	Succeeded	Global Standard	9/30/26, 9:58:45 PM	
lab-planner-dea3cec5	gpt-5.5	2026-04-24	Succeeded	Global Standard	9/30/26, 9:58:44 PM	
lab-judge-dea3cec5	gpt-5.4-mini	2026-03-17	Succeeded	Global Standard	9/30/26, 9:58:42 PM	
lab-agent-dea3cec5	gpt-4.1-mini	2025-04-14	Succeeded	Standard	9/30/26, 9:58:40 PM	

1-6 of 6

Prev Next

Locate model/version fields. The picture predates the Sol deployment. Use the role table and current recorded API observations, not this older list, as runtime evidence. View full-size image (file: web/assets/portal/en/02-model-deployments.png)

Prepare a strong v1, then a controlled v2

Use `prompts/en/baseline.txt` for the English baseline. It covers policy evidence, dates, routing, uncertainty and action boundaries. V1 must not be weakened to manufacture an apparent improvement.

V1 and v2 are complete Foundry Agent versions, not just document revision labels. In this comparison only their instructions differ. Changing the model, tools, reasoning or output schema requires a new matched experiment, not reuse of old-model scores.

Released versions are immutable. `ensure_fixed_release` in `lab/agents.py` permits only **1 and 2**, reuses a matching existing release, blocks mismatches and never silently creates v3. For development, use explicitly pinned `draft-...` versions. Draft support was verified in this subscription; verify it elsewhere because unsupported draft creation can fall back to a regular release.

`native_response_format()` projects the shared response schema into the supported strict-generation subset. Citation uniqueness and `needs_human/route` consistency are still checked after generation; a valid JSON shape alone is not policy correctness.

Keep the English registered dataset `contoso-eval-en-dev12` **version 1**, all 12 rows from `data/en/optimizer/dev.jsonl`, with the original SHA-256. The Agent receives query **only**; reference columns are not injected into generation.

Provide Python 3.11–3.14, Azure CLI, the checked-out source and `requirements.lock`. Use a separate prepared environment, not an unrelated user's environment:

```
source .venv/bin/activate
python -m pip install -r requirements.lock
az login
az account show --query "{user:user.name,tenant:tenantId,subscription:id}" -o json
```

Check those values against the approved project. The operator supplies its endpoint, subscription, current evaluation ID and completed baseline run ID. Do not copy historical rehearsal IDs or use an Optimizer job ID as an evaluation ID.

The reevaluation helper clones the baseline data source and validates the **remote** evaluator contract. A stale local `JUDGE_DEPLOYMENT` must not replace the actual Judge in the evaluation definition.

Freeze the comparison and spending scope

Item	Fixed requirement
Data	Same English dev12 bytes, registration/version, questions and references
Relevance	Managed <code>builtin.relevance</code> , scale 1–5, threshold 4
TaskAdherence	Managed <code>builtin.task_adherence</code> , binary 0/1, pass 1
Judge	Same explicit Luna deployment in both evaluator definitions
Agent	Same Sol model/version, tools, reasoning and strict output schema
Change	Instructions only; releases remain v1/v2
Optimizer	Instruction target only, model/tool-description changes off, bounded candidate count
Publication	Latest v2 and its frozen v1 control; all 12 sanitized case outcomes

Catalog evaluator versions are not proof that the private service rubric is fully pinned. Preserve the actual definition, settings and this limitation. A passing gate on reused dev12 does not guarantee future scores, independent generalization or production approval.

Verify least-privilege access

Use the [current Foundry RBAC guidance \(https://learn.microsoft.com/azure/foundry/concepts/rbac-foundry\)](https://learn.microsoft.com/azure/foundry/concepts/rbac-foundry). Participants need only the authorized project operations. Role-assignment authority is separate from model or evaluation access.

Keep the policy tool read-only. Do not assign subscription-wide Owner, disable network restrictions, change organization policy or transfer tokens between users as a shortcut.

Costs and operational boundaries

Evaluation invokes the Agent and Judge; optimization makes additional internal calls. Quota is not free usage, and GlobalStandard is not an NCUS-only processing guarantee. Record measured tokens and actual billing separately; unknown billing is not zero.

Name the cost/cleanup owner. Do not publish to production or activate a candidate merely because an Optimizer ranking increased. Keep the baseline active unless the measured decision and authorized operational review justify a change. Closing a browser does not stop resource charges.

Hand off the current pair

Give participants the current Agent/version IDs, deployment-role map, dev12 hash, exact evaluator settings, SDK context, approval scope and [latest verification record](#). The English report publishes actual responses, scores and reasons for every case.

Raw service files can contain account metadata, tokens or signed URLs. Keep those local; **the synthetic evaluation results themselves are not secret**. Publish only allowlisted fields, including failures, and exclude credentials—not inconvenient outcomes. Historical attempts stay in the private audit, not as additional current-version reports.

Latest v2 verification · gpt-6-sol

V2 improved the measured quality result over the same-model v1: all-criteria passes increased **11/12 → 12/12**, mean Relevance **4.8333 → 4.9167**, and TaskAdherence remained **12/12**. This page reports only the current Sol pair, not historical experiments.

Important limit: p95 latency increased **10.97 → 42.97 seconds**, and the native statistical comparison is **Inconclusive**. The measured quality gain does not establish production readiness, statistical significance or guaranteed improvement on future runs.

What was measured

Both versions of `contoso-eval-en-sol` used the same `gpt-6-sol / 2026-09-22` Agent, read-only policy tool, reasoning setting, strict JSON schema and registered `contoso-eval-en-dev12` **version 1**. Only the instructions changed.

All 12 unchanged records came from `data/en/optimizer/dev.jsonl` (file: `data/en/optimizer/dev.jsonl`). Only `query` was sent to the Agent; `context` and `ground_truth` remained references. The two managed evaluators and their Luna Judge were unchanged. Actual version and system instructions were checked in every result item.

These are **actual English-data results**, not a newly executed Korean run. Local unit tests and response-format checks do not replace Microsoft Foundry managed Evaluation.

Final v1/v2 results

Metric	Sol v1 control	Sol v2
Cases	12	12
All-criteria passes	11/12	12/12
Relevance passes, threshold 4	11/12	12/12
Mean Relevance, 1–5	4.8333	4.9167
TaskAdherence, binary pass 1	12/12	12/12
Execution errors / skipped cases	0 / 0	0 / 0

Metric	Sol v1 control	Sol v2
Agent p50 latency	8.03 s	9.71 s
Agent p95 latency	10.97 s	42.97 s
Observed Agent tokens	46,167	55,858

Evaluation definition: eval_40a593c037e44045a47fe5088438afc5.

V1 run: evalrun_077d41ff8a534b9caee64f9d6c2a339d. Final released v2 run: evalrun_6c9e78cfc9de4f1282eb621c6678d8b1. The report uses the actual released version **2**, not a draft relabeled as v2.

Quality acceptance: PASS for this measured lab pair. No evaluator pass count or mean declined, one additional case crossed the Relevance threshold, and all candidate responses met the JSON/routing/citation checks. Production approval is **not granted**, especially given the observed latency increase.

The isolated unpublished lab Agent now selects **active version 2**. The frozen v1 control remains available for comparison; no production channel or release beyond v2 was created.

All twelve cases are public

The machine-readable **evidence/latest.json** (file: evidence/latest.json) includes case_evidence: all 12 original questions and references, both actual response texts, each evaluator's score/Pass/Fail/reason, response-format findings and latency. No difficult case was dropped.

Case	Relevance v1 → v2	TaskAdherence v1 → v2
01 Subscription intent	5 → 5	1 → 1
02 Refund effective-date transition	5 → 5	1 → 1
03 Paid-production-job disqualifier	5 → 5	1 → 1
04 Unsubmitted support inquiry	5 → 5	1 → 1
05 Suspected access incident	5 → 5	1 → 1
06 Separate monthly SLA calculations	5 → 5	1 → 1
07 Beta feature versus production SLA	5 → 5	1 → 1
08 Availability target versus feature claim	5 → 5	1 → 1

Case	Relevance v1 → v2	TaskAdherence v1 → v2
09 Inclusive/exclusive time boundaries	5 → 5	1 → 1
10 Unauthorized tenant deletion	5 → 5	1 → 1
11 Unverified feature and false alternatives	3 → 4	1 → 1
12 Missing SLA calculation inputs	5 → 5	1 → 1

Case 11 improved without pretending to inspect the user's screen or inventing a feature status. V2 gave a clear verdict that **neither asserted conclusion was established**, distinguished missing evidence from absence, and provided an authenticated verification path.

An excerpt from the actual v2 response:

Neither conclusion is established. You don't see a GPU model selection menu, but I haven't inspected your screen or the live service.

The full response and unedited managed-evaluator reasons are in the public JSON. Evidence means actual execution data—not credentials. Tokens, cookies, signed URLs, account identifiers and raw conversation/tool payloads are excluded.

Model roles and instruction provenance

Role	Model/version	Boundary
Agent	gpt-6-sol / 2026-09-22	Actual Agent and knowledge-tool runtime verified; same in v1/v2
Evaluation Judge	gpt-6-luna / 2026-09-22	Used by both real managed runs
Optimizer generator	gpt-5.5 / 2026-04-24	Separate role following the supported optimization-model list

Actual managed Optimizer job: opt_87805603c0d74b7897f997502aaea080, instruction-only, at most two generated candidates. The service **retained the strong v1 baseline** rather than recommending either generated candidate.

V2 was an **operator-reviewed refinement after that job**, not a falsely attributed automatic Optimizer promotion. It preserved the strong baseline and clarified how to answer unsupported alternatives with a factual verdict and verification path. It contains no embedded evaluation questions, reference answers, policy

rates or worked dataset examples.

Source: v1 instructions (file: `prompts/en/baseline.txt`) · v2 instructions (file: `prompts/en/optimized.txt`).
Candidate development used explicit drafts; the released workshop Agent has **only versions 1 and 2**.
Immutable versions are not overwritten or silently incremented.

Statistical and operational limits

The real Foundry comparison used **PairedTTest**:

Evaluator	Mean difference	p-value	Service result
Relevance	+0.08333	0.33880	Inconclusive
TaskAdherence	0	1.00000	Inconclusive

The descriptive gain is real in these saved results, but the sample is small and reused during instruction development. Do not claim independent generalization, statistical significance, lower cost or faster operation. The p95 and token increases remain visible rather than being hidden behind the improved pass rate. Actual billing was not observed and is not reported as zero.

`scripts/compare_foundry_eval.py` enforces matched complete results, unchanged conditions, no lower aggregate quality counts/means and at least one strict improvement. It also retains malformed outputs as failures. This is a validation gate over **service-returned scores**, not a new local Judge.

Sources and publication boundaries

Only the latest v2 and its same-model v1 control appear in this current report. Earlier raw records remain in the local audit; they were not relabeled or erased from Azure/Git history.

The capture manifest (file: `web/assets/portal/en/captures.json`) and NOTICE (file: `web/assets/NOTICE.txt`) describe the real portal images. Those earlier images are **UI illustrations only**, not current Sol scoring evidence. Use the current run IDs and full public case record for the measurement.

Official references: [Agent evaluation](https://learn.microsoft.com/azure/foundry/observability/how-to/evaluate-agent) (https://learn.microsoft.com/azure/foundry/observability/how-to/evaluate-agent) · [Agent Optimizer](https://learn.microsoft.com/azure/foundry/agents/quickstarts/quickstart-optimize-prompt-agent) (https://learn.microsoft.com/azure/foundry/agents/quickstarts/quickstart-optimize-prompt-agent) · [Optimization-model roles](https://learn.microsoft.com/azure/foundry/agents/concepts/agent-optimizer-overview#models) (https://learn.microsoft.com/azure/foundry/agents/concepts/agent-optimizer-overview#models) · [Evaluator definitions](https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/agent-evaluators) (https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/agent-evaluators).

Contoso data is synthetic. Earlier design attribution remains at the [archived source commit](https://github.com/junwoojeong100/foundry-evaluation-labs-v0.9/tree/93bc07e31373c4cfc278a2dc3757785946404cf2) (<https://github.com/junwoojeong100/foundry-evaluation-labs-v0.9/tree/93bc07e31373c4cfc278a2dc3757785946404cf2>); it is not another current verification report.

Recheck local deliverables

```
python scripts/build_datasets.py --language en --check
python scripts/build_datasets.py --language ko --check
python scripts/build_guide.py --check
python -m unittest discover -s tests -q
```

These commands check files and tests without submitting another Azure evaluation. The latest JSON separately records document, browser, PDF and service checks.

Synthetic English evaluation data

[Dataset step](#) · [Operator handoff](#) · Korean data guide (file: docs/ko/data-guide.html)

One unchanged dataset for the comparison

Use **data/en/optimizer/dev.jsonl** (file: data/en/optimizer/dev.jsonl), exactly **12 JSONL records**, for Foundry Evaluation, Agent Optimizer and reevaluation. The existing registered dataset is **contoso-eval-en-dev12 version 1**. Reuse that registration rather than uploading a duplicate.

Keep its exact bytes and SHA-256, all questions and references, and language fixed. The current Agent is **contoso-eval-en-sol**, with released instruction variants v1/v2 on the same **gpt-6-sol** model. The original Korean corpus is separate and has not been newly executed as part of this English measurement.

Purpose and limits

Contoso Atlas Cloud is fictional. Its policies, inquiries and reference answers are synthetic—not private customer data or the terms of Microsoft, Azure or a real provider.

The purpose is **evaluation against your own business tasks and criteria**, not a public benchmark alone. Representative tasks expose policy-date boundaries, ambiguity, unsupported certainty and action claims. A 12-case development set demonstrates the learning loop; it does not certify production or independent generalization.

Preserve the original source corpus

data/en/cases.jsonl (file: data/en/cases.jsonl) contains 100 cases: **train 56**, **validation 12**, **dev 12** and **test 20**. Only dev12 is used in this workshop. Other source subsets remain available but are not additional participant exercises.

Do not select only easy rows, duplicate questions, merge subsets or replace references after seeing results. The eight policy documents (file: **data/en/knowledge/documents.json**) retain stable **ATLAS-*** IDs and effective dates. Policy access is prepared by the operator; participants do not build a separate knowledge system.

Upload and response schema

Each JSONL line has exactly three columns:

Column	Type	Purpose
query	String	The only Agent input
context	String	Source-policy reference for supported evaluators and case review
ground_truth	JSON string	Structured reference answer, not a generation prompt

There is no prefilled response. Foundry invokes the pinned Agent to obtain one. Do not turn the file into a JSON array, rename columns, append references to the user message or precompute answers.

The reference answer and Agent response use exactly four keys:

Key	Contract
answer	Nonempty English string
citations	Unique stable IDs of supporting policy documents
route	answer, clarify, escalate or refuse
needs_human	Boolean; true exactly when route is escalate

The source response schema (file: `schemas/response.schema.json`) also checks routing consistency. Native structured output controls the JSON shape, not the factual truth of an answer. Numeric retrieval IDs are not policy citations.

The Agent cannot actually submit tickets, change subscriptions, approve credits or delete records. Classification as `escalate` does not mean a person was contacted. A request for unsupported certainty does not authorize fabricated facts.

Register once; pin the same version

Use **Foundry → Build → Evaluations → Create → Agent**, an explicit baseline version, **Individual turns / One time / Existing dataset**. Reselect the target checkbox if version selection clears it.

Choose the existing dev12 registration. In a fresh authorized project only, upload the exact language-specific file and record its new dataset version. The preview may show five rows; the file and evaluation scope still contain 12.

Leave the custom prompt override unset. Use `query → query` if a field-mapping screen appears. Service mappings are **Relevance** `response={{sample.output_text}}` and **TaskAdherence** `response={{sample.output_items}}`. These are not extra JSONL columns.

Keep **Relevance 1–5 / threshold 4**, **TaskAdherence binary 0/1 / pass 1**, and the same Luna Judge. Selecting two evaluators does not mean both consume every reference field or certify every business rule.

Compare complete evidence

The [latest report](#) publishes the current v2 and its frozen same-model v1 control, including all 12 synthetic questions, actual answers, evaluator scores and reasons. Older raw runs are retained for audit, not mixed into the current comparison.

The helper checks matching questions/references and actual per-item Agent version/instructions. Missing/error rows cannot become successful rows or disappear from the denominator. JSON/route checks are supplementary validation, **not a replacement local Judge**.

Agent Optimizer's internal ranking is distinct from the separate managed run's means and pass counts. Review the candidate instructions and preserve model/tools/reasoning/schema. Use drafts during development; do not continually increment the released workshop versions beyond v1/v2.

Public evaluation evidence excludes authentication, signed URLs and private account metadata. The synthetic responses and failure reasons themselves may be shared. Publish failures as faithfully as successes.