

Foundry Learning Loop Lab · Complete Print Edition

Follow the six participant steps for Foundry Evaluation, Agent Optimizer, and same-criteria comparison. Operator, facilitator, data, and verification references are included but are not extra labs. Cross-document links stay within this book; file references use package paths.

As of 2026-10-01

Included documents

1. [Participant guide](#)
2. [Appendix 1 · Facilitator guide](#)
3. [Appendix 2 · Operator setup](#)
4. [Appendix 3 · Verification and sources](#)
5. [Appendix 4 · Data guide](#)

Participant guide

Great agents start with evaluation · v1

CONTOSO ATLAS CLOUD · ONE LEARNING PATH

Evaluate your business cases in Foundry, improve instructions with Agent Optimizer, and measure what changed. Follow these six steps.

01 [Prepare the dataset](#)

02 [Select evaluation criteria](#)

03 [Run Foundry Evaluation](#)

04 [Read scores and reasons](#)

05 [Improve instructions with Agent Optimizer](#)

06 [Reevaluate, compare, and finish](#)

Why private/domain evaluations? Your representative tasks, policy boundaries, and known failures matter more to your application than a public benchmark alone. The synthetic Contoso scenario demonstrates **evaluate → learn → improve → reevaluate**, not production certification.

Before starting: The [operator](#) supplies the isolated project, prepared agent and baseline version, deployment names, and cost/data authorization. Infrastructure is a prerequisite, not another participant lab. Allow processing/discussion time; wait at most **60 minutes per optimization job**.

Keep one language per experiment. Use the English file and English agent here, or the separate Korean pair. Switching this website's language does not translate the agent or dataset.

Real English UI and data: Screens illustrate the completed English rehearsal, not your own results or a Korean run. Use your own run IDs and read the [measured results and limitations](#). Previewing a wizard is not executing an evaluation.

01. Prepare the dataset

What: A small, repeatable set of domain tasks. **Why:** The same questions expose what changed after instruction optimization. **How:** Inspect the supplied JSONL, then upload all 12 cases to an evaluation draft.

Use `data/en/optimizer/dev.jsonl` (file: `data/en/optimizer/dev.jsonl`), unchanged. It has **12 JSONL rows**, not a JSON array. The source corpus has 100 cases, but only this dev12 file is used in the workshop. Other subsets are unused here.

Column	Type	Use
query	String	The only input sent to the agent
context	String	Original policy excerpts kept as evaluator/reviewer reference
ground_truth	JSON string	Structured reference answer, not a generation prompt or precomputed response

Confirm all 12 records locally and record the file's **SHA-256** from the operator's handoff. Keep the same bytes, registered dataset version, and hash for baseline, optimization, and reevaluation. Do not trim, duplicate, translate, or add a fake `response` column. See the [data contract](#).

The prepared agent returns exactly four response keys. These are contract examples, not measured answers:

Key	Contract
answer	Nonempty English text; for example, "Please confirm the purchase date."
citations	Supporting policy IDs, for example ["ATLAS-REF-001"]
route	answer, clarify, escalate, or refuse
needs_human	Boolean; true only for escalate

The assistant advises; it does not execute refunds, create tickets, or grant access. A reference to escalation is not evidence that a human was contacted.

Portal action: Open [Foundry \(https://ai.azure.com/\)](https://ai.azure.com/) in the **New experience**, select the operator's project, then **Build → Evaluations → Create → Create new evaluation**.

1. For a fresh baseline agent still at v1, choose **Agent** → `contoso-eval-en` → **Pin currently latest**, not Latest available, before changes. **Reselect the cleared checkbox and confirm 1 target before Next.** The recorded lab retains candidate v2 but has restored active v1: use explicit versions and the saved baseline, not assumptions about "latest." The failed gpt-6-luna probe was another agent.
2. Set **Scope: Individual turns**, not Full conversations, and **Frequency: One time**, not recurring.
3. Under **Data**, choose **Existing dataset**, not the default Synthetic generation, Benchmarks, or Existing traces.
4. Select `contoso-eval-en-dev12` **version 1**, registered through this wizard from the unchanged 12-row `data/en/optimizer/dev.jsonl`. Do not upload it again. In a fresh authorized project where it is absent, use **Upload new dataset** → **name** → **Choose file** → **Upload** with that exact file and record the resulting version.
5. Confirm the dataset is selected and previewed, then record its actual registered version. In the observed English wizard, the matching schema **auto-resolved Field mapping**, so the next screen was **Configure agents**. If a schema-dependent mapping screen appears, use **query** → **query**.

Dataset selection. Use Existing dataset and the selected upload. The preview shows only the first five rows; that is not proof that the dataset contains five cases. View full-size image (file: `web/assets/portal/en/15-evaluation-dataset.png`)

Completion signal: The draft points to the explicit baseline version and the selected 12-row dataset; its local count/hash are recorded. Preview is preparation, not execution. If the file or version differs, correct the draft before proceeding.

02. Select evaluation criteria

What: Two managed evaluators with different result types. **Why:** A comparison must respect each metric's scale.

How: Use Relevance threshold 4 and TaskAdherence binary pass 1, with the verified Judge deployment.

In **Configure agents**, leave the **custom prompt override unset**. Agent input must be **query only**. Keep **context** and **ground_truth** separate as evaluator references; do not append either to the question to make an answer look better.

In **Criteria**, the observed portal suggested **23 evaluators**. Remove the extras and retain **Relevance** and **TaskAdherence** only. The Optimizer may display the latter as **Task Adherence**.

Evaluator	What to inspect	Required setting
Relevance	Whether the response addresses the question; 1–5	Edit Relevance evaluator → Threshold 4 → Update
TaskAdherence	Whether the response follows task instructions; Binary Pass/Fail, raw 0/1	Pass = 1 . If a generic threshold control is shown: Edit TaskAdherence evaluator → Threshold 1 → Update

The [official agent-evaluator definitions](https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/agent-evaluators) (<https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/agent-evaluators>) identify TaskAdherence as **binary**. Read its service-returned Pass/Fail; do not interpret a passing 1 as a one-out-of-five score.

In **Model: Judge model**, choose `lab-judge-luna-dea3cec5`, not the agent model. Record the two evaluator names, result types, settings, Judge deployment/model version, and dataset version/hash. Keep exactly two evaluators; this correction adds no extra participant exercise.

Catalog references are not service-version pins. Catalog links showed `relevance v14` and `task_adherence v17`, but submitted service criteria had `evaluator_version empty/default`. Do not claim that private service rubric versions are fully pinned. Retain the saved evaluation definition and disclose this limit.

Observe mappings; do not override generated values. The corrected submission confirmed **Relevance** `response={{sample.output_text}}` and **TaskAdherence** `response={{sample.output_items}}`. UI defaults also showed `query={{item.query}}` and TaskAdherence's `tool_definitions={{sample.tool_definitions}}`. Inspect **Raw JSON**; do not force TaskAdherence back to the UI's earlier output-text binding. These are service bindings, not extra JSONL columns or generation inputs.

Three model roles are separate. The request is to use `gpt-6-luna` wherever that role is actually supported, not to force one model into every role.

Role	Workshop selection boundary
Agent answering the questions	lab-agent-dea3cec5 → gpt-4.1-mini / 2025-04-14. Use verified contoso-eval-en v1 with the existing read-only knowledge connection; the gpt-6-luna Responses/Agent attempt failed here
Evaluation Judge	lab-judge-luna-dea3cec5 → gpt-6-luna / 2026-09-22, verified in the actual managed baseline and candidate evaluations
Optimizer instruction generator	lab-planner-dea3cec5 → gpt-5.5 / 2026-04-24. The official supported optimization-model list does not include gpt-6-luna

Capability verification is not agent quality. The earlier gpt-6-luna native Relevance check used **one authored compatibility fixture: passed 1 / total 1 / errors 0**; Chat Completions returned READY. That fixture and the completed configuration pilot are separate from the corrected comparison. The [operator model note](#) records their distinct roles.

Direct Responses and a **pinned** gpt-6-luna **prompt-agent v2** returned **HTTP 500 in this environment**. Keep the verified gpt-4.1-mini agent until the provider issue is resolved. This is a failed runtime verification here, not a claim that the model is unsupported everywhere. Catalog visibility is not runtime proof.

The existing gpt-4.1-mini model version has an Azure retirement date of **2027-04-14**. Public documentation's **Deprecated** label can mean restrictions for fresh subscriptions; verify the actual account rather than assuming universal availability. The separate Optimizer generator follows the [official allowed model list](https://learn.microsoft.com/azure/foundry/agents/concepts/agent-optimizer-overview#models) (https://learn.microsoft.com/azure/foundry/agents/concepts/agent-optimizer-overview#models).

Microsoft Foundry / proj-fev11-dea3cec5

Search with AI (96K)

New Foundry

Home

Discover

Build

Operate

Manage

Docs

Create

Agents

Models

Services

Tools

Knowledge

Memory

Guardrails

Data

Optimize

Evaluations

Fine-tune

← Create new evaluation

Target: Agent

Scope

Frequency

Data

Configure agents

Criteria

Review

Mapped fields from your dataset:

query: {{item.query}}

response: {{sample.output_text}}

context: {{item.context}}

ground_truth: {{item.ground_truth}}

tool_calls: {{sample.tool_calls}}

tool_definitions: {{sample.tool_definitions}}

Instructions

What evaluators are not auto-selected here?

Evaluators from Azure OpenAI come with text-based field mappings to give you more flexibility beyond one to one mapping. You will need to specify field relationship when you construct your system, user and assistant prompts. Don't worry, we have templates so you won't have to start from scratch.

Choose what criteria you want to evaluate — quality, risk and safety, and more.

Judge model *

Model: lab-judge-luna-dea3cec5

Global Standard deployment

We've auto-suggested the following 2 evaluators

Agents (1)

Remove all

TaskAdherence

Quality (1)

Remove all

Relevance

Add new evaluator

Back

Next

Read each metric's own scale. Relevance uses threshold 4; TaskAdherence uses binary pass 1. Keep only these two evaluators and verify Model: Judge model separately. View full-size image (file: [web/assets/portal/en/16-evaluation-criteria.png](#))

Completion signal: Relevance **1–5 / threshold 4**, TaskAdherence **binary 0/1 / pass 1**, query-only input, and the verified Judge are recorded, together with generated mappings and the service-version limitation. If setup differs, correct it before submission rather than silently changing metrics or models.

03. Run Foundry Evaluation

What: An actual Microsoft Foundry managed agent evaluation. **Why:** Configuration and model connectivity are not measured task performance. **How:** Review the exact contract, submit once, and follow the real evaluation/run IDs to completion.

For a fresh authorized workshop, continue to **Review** using the [direct agent-evaluation workflow](#) (<https://learn.microsoft.com/azure/foundry/observability/how-to/evaluate-agent>). The corrected English contoso-en-learning-loop baseline is already portal-submitted: open its recorded run instead of submitting again. Keep the earlier pilot unchanged and separate from this comparison.

Review item	Required value
Target	contoso-eval-en, Pin currently latest resolved to version 1 before configuration changes
Scope / frequency	Individual turns / One time
Data	contoso-eval-en-dev12 version 1 , all 12 cases, original file hash
Agent configuration	query only; no custom prompt override
Criteria / Judge	Relevance threshold 4; TaskAdherence binary pass 1 (generic Threshold 1 if shown), generated mappings, lab-judge-luna-dea3cec5
Evaluation name	contoso-en-learning-loop, the corrected comparison definition; not the retained pilot name

Microsoft Foundry / proj-fev11-dea3cec5

Search with AI (Beta)

New Foundry

Home

Discover

Build

Operate

Manage

Docs

Create

Agents

Models

Services

Tools

Knowledge

Memory

Guardrails

Data

Optimize

Evaluations

Fine-tune

← Create new evaluation

✓ Target: Agent

✓ Scope

✓ Frequency

✓ Data

✓ Configure agents

✓ Criteria

⚙ Review

Add a name and description to this evaluation run.

Evaluation name *

contoso-en-learning-loop

Back

Submit

Summary

Agent evaluation

Targets

contoso-eval-en: v1

Agent

Edit

Scope

Individual turns

Scope

Edit

Frequency

One time

Frequency

Edit

Dataset

contoso-eval-en-dev12

Version 1

Existing dataset

Edit

Evaluators

TaskAdherence

Agents

Edit

Relevance

Quality

Edit

Review before submission. Confirm the corrected definition name, pinned agent, dataset, Relevance 4, TaskAdherence binary pass 1, and Judge. A pilot's old settings are not the corrected contract, and a review screen is not execution evidence. View full-size image (file: web/assets/portal/en/17-evaluation-review.png)

The final action is **Submit**, used once for the authorized baseline. The English submission below already exists; do not submit it again or rewrite the retained pilot.

Completed canonical portal baseline: `contoso-en-learning-loop`, evaluation `eval_94feef6f6f644fabb22a5680f5f24fb1`, run `evalrun_cde9948ac9d946929661bc3d9e60432a`. The HTTP 201 creation receipt was followed by actual completion on `contoso-eval-en` pinned v1 and `contoso-eval-en-dev12` v1, with Relevance threshold 4, TaskAdherence binary pass 1, and the same `gpt-6-luna` Judge. These are baseline results, not candidate results.

The **EVALUATION details page** provides **Raw JSON** and **Evaluation runs**. Keep the definition/run IDs, URL, and submission time. The older pilot remains in the [operator record](#), not this comparison. The visible Add run candidate wizard currently fails; use step 06 rather than repeating it.

Completion signal: The actual run page reports **Completed**, and you can inspect its results. Still account for **all 12 cases**, including failed, errored, or missing rows. A failed or partial run stays failed/partial in your record; a completed run is not automatically a quality pass.

If access, model support, or scoring fails, keep the IDs and error details, use the [facilitator's recovery reference](#), and mark later work not run. Do not replace a failed managed evaluation with authored scores.

04. Read scores and reasons

What: Per-metric scores and row-level explanations. **Why:** An average can hide a policy mistake or a missing answer. **How:** Read both evaluator summaries, then compare a failure or lowest-score case with a good case.

On the **Evaluations run page**, read the summaries and **Detailed metrics result**, including `Relevance.reason` and `TaskAdherence.reason`; scroll horizontally if necessary. A row's **conversation_id → User view** opens its question and actual JSON response, **not an inline Judge-reason panel**. Return to Detailed metrics result for evaluator reasons.

Actual canonical baseline only: Do not substitute pilot aggregates or label these as candidate results.

Baseline result	Observed value
Execution	Completed; 12 rows
Overall	10 passed / 2 failed / 0 errored
Relevance	10/12 passed, threshold 4
TaskAdherence	12/12 passed, binary 1

Record for each evaluator	Interpretation
Total scope	n = 12, including rows without usable scores

Record for each evaluator	Interpretation
Scored / total, missing, errors	Coverage, not extra successful cases; do not silently drop rows
Relevance result	1–5 , threshold 4 . A score of 4/5 is not 80% accuracy
TaskAdherence result	Binary 0/1: 1 = Pass, 0 = Fail . A missing/error result is not a scored 0
Pass count and score summary	Report separately for Relevance and TaskAdherence; do not merge them into a new accuracy claim
Row-level reason	The evaluator's explanation to check against the actual answer and reference

Microsoft Foundry / proj-fev11-dea3cec5

Q Search with AI (26K)

New Foundry

Home

Discover

Build

Operate

Manage

Docs

Create

Agents

Models

Services

Tools

Knowledge

Memory

Guardrails

Data

Optimize

Evaluations

Fine-tune

← contoso-eval-en

Make recurring

Delete

Run details

Status

Completed

Create time

10/1/2026, 12:21:37 AM

Duration

3m 59.84s

Created by

Agent ID

contoso-eval-en:1

See all properties

Raw JSON

User logs

Download user logs

Results

Download results

Overall metric results

P50 latency	P95 latency	TaskAdherence	Relevance
5.9s	8.8s	100% 12 / 12	83% 10 / 12

Detailed metrics result

Analyze Results

Index	query	conversation_id	trace_id	Relevance
0	Next month's cost is a concern. Please reduce my subscription. I...	conv_07e180556bd61ca800UPWgFCcJqavhmEPmHcTjBXGDhbqHYFFz	-	Fail: 3
1	My first monthly payment was on August 28, 2026 at 08:00 KST. I...	conv_0bb9168bd3bf11ec00jyPvdLhly3iE1iHfKa6uCYJS9YCcWplx	-	Pass: 5
2	I used the service for just one day after my first monthly payment on...	conv_05f494a8836c0111004DizSunhTEdBq7cEwzgpPeP3f4M5UKYcT	-	Pass: 4
3	I typed only a title on the Support > Contact support screen and left...	conv_09e82e7f1e3db55a00h7GlpAUkG8N9Ngm0jrFoQUAXyEm0k5r	-	Pass: 4
4	There is an unfamiliar access event on the administrator's audit screen...	conv_05740a8abb5d88ad00G5SaTweilR5t4uMTWWWtdYBV0kq4lgQVM	-	Pass: 5
5	Assume I paid the same monthly base fee in two months, with...	conv_098f970e5fa2d1da00zhmGbrfVSMhdOFzI7ILUcjB4PGVLH8UFJ	-	Pass: 5
6	I'm on paid Pro, but only beta webhook delivery failed; the...	conv_099677dc019f239900fUjqcz6Qpnwkk97SDxfp0SUV8pXxqrUq	-	Pass: 5
7	Since there is a 99.9% monthly commitment, can I write in a...	conv_0a857d1c7aec9b6800pluwLFZNNSVYwbrpuCehFCbQhFYecEUo	-	Pass: 5
8	I learned that the refund window includes the first business day...	conv_0eb63b9bf392360b00tONh2DhBm9f7ZuclmF4HrEDrGf67NA	-	Pass: 5

Read the canonical baseline.

Match the recorded run ID: 12 rows, 10 passed, 2 failed, 0 errored overall; Relevance 10/12 and binary TaskAdherence 12/12. Read reasons in Detailed metrics result. These are not candidate results. View full-size image (file: web/assets/portal/en/18-evaluation-results.png)

Select **at least one failing or lowest-score case and one good case**, if a good case exists. For each, connect the **actual question → actual answer → evaluator score/reason → policy context and parsed ground_truth reference**. If no good case exists, say so rather than inventing one.

Check concrete domain behavior: Did the answer apply the right policy date? Ask for genuinely missing information? Cite supporting policy? Claim an approval or action that never happened? A reference answer need not match word for word.

The canonical v1 baseline's two Relevance failures: TaskAdherence passed for both; these are not candidate regression findings.

Case	Observed answer and reference distinction
atlas-dev-001	The subscription-reduction answer was vague; the reference required explicit clarification
atlas-dev-011	The answer honestly described an undocumented feature as unverified, but Relevance judged it incomplete

Inspect the actual reasons and reference policy. **Do not change policy facts or replace justified uncertainty with unsupported confidence merely to chase Judge scores.**

conv_07e180556bd61ca...

Conversation ID: conv_07e180556bd61c...

Trajectories

User view

Graph view

3 spans1 chat call1.8s1.1Ks

Next month's cost is a concern. Please reduce my subscription. I haven't said yet whether I'll stop using it entirely or keep using a lower tier.

{
 "answer": "I understand your concern about the upcoming cost. You can reduce your subscription to a lower tier or make other changes to your subscription plan to better fit your needs. Please let me know if you'd like help with specific options or if you want assistance with pausing or canceling your subscription entirely. I'm here to support you in managing your subscription in the best way possible.",
 "citations": [],
 "route": "answer",
 "needs_human": false
}

conversation_id → User view: inspect the actual response.

This shows the atlas-dev-001 question and JSON response, not an inline Judge-reason panel. Return to Detailed metrics result for Relevance.reason and TaskAdherence.reason, then compare the reference policy. View full-size image (file: web/assets/portal/en/19-evaluation-case.png)

Limits: Relevance and TaskAdherence are useful, compact criteria, not a complete policy or safety certification. Not every evaluator consumes context or ground_truth; those columns remain available for supported reference use and your review. Do not claim that selecting two criteria checks every response-contract rule.

Discuss: State the actual run ID, n = 12, each metric's coverage/pass count, one problematic claim, and one good response. Explain which instruction behavior you want to improve and why.

Completion signal: You can support an improvement hypothesis with an actual answer and reason, while keeping errors visible. If no failure is observed, use the lowest-score case and report that result honestly; do not weaken the baseline to manufacture a lesson.

05. Improve instructions with Agent Optimizer

What: One instruction-only Agent Optimizer candidate. **Why:** A controlled change makes the later comparison interpretable. **How:** Keep the agent model, tools, dataset, and Judge fixed; inspect the proposed instruction diff before creating a separate lab version.

Actual English job succeeded: `opt_e44bcf5701a348deb62a1cd4f9cb3910`, **one candidate**, instruction-only, same dataset, `gpt-5.5` generator, `gpt-6-luna` Judge, no model comparison. Optimizer UI scores were **0.635 → 0.646**, with **UI-reported +0.010** and **264,260 reported tokens**. These are service optimization results, not direct candidate-evaluation scores; do not recalculate a different delta from the displayed values or resubmit the job.

Open **Build → Agents → `contoso-eval-en` → Optimize Preview**. On first use, choose **Optimize my agent → Agent**, not **Cost**. After jobs exist, the menu may read **Optimize → Agent**. The first-run tax-agent benchmark is a **product example**, not Contoso results.

In **Create an optimization run**, follow the [official prompt-agent Optimizer workflow](https://learn.microsoft.com/azure/foundry/agents/quickstarts/quickstart-optimize-prompt-agent) (<https://learn.microsoft.com/azure/foundry/agents/quickstarts/quickstart-optimize-prompt-agent>):

Setting	Exact workshop choice
Agent version	<code>contoso-eval-en</code> version 1 , the same pinned baseline used in direct Evaluation
Choose targets	Uncheck Model ; select Instruction only ; Tool description off
Max candidates	1 . Clear Model first: model-only selection can disable this field
Optimization model	<code>lab-planner-dea3cec5</code> , <code>gpt-5.5 / 2026-04-24</code>
Evaluation model	<code>lab-judge-luna-dea3cec5</code> , the same verified <code>gpt-6-luna</code> Judge used in direct Evaluation
Model comparison	Off; do not change the model being optimized

Create an optimization run

1 Target

Optimization tunes your agent's instructions and tool descriptions. Optimization can also tune your agent's quality by attempting different models. To also compare models, turn on Compare across models and pick which ones to include.

Goal
Improve quality
Raise evaluation quality above your current version, with token cost left free.

Agent: Version:

Optimization approach

☐ Auto
Foundry chooses which eligible attributes to optimize based on the available budget.

☒ Choose targets
Choose the attributes Foundry should attempt to optimize within the available budget. Targets that do not apply to your agent are skipped.

☒ Instruction ☐ Tool description ☐ Model

Optimization model Max candidates:

Evaluation model

Next

2 Data

3 Review

How optimization works

Scored and ranked
Each candidate is scored against your criteria and ranked by performance.

Verify before applying
Review the top candidate before applying changes to your agent.

Choose the target before the candidate limit. Instruction only keeps model changes out of the experiment. The optimization model generates instructions; the evaluation model scores responses. View full-size image (file: `web/assets/portal/en/07-optimizer-target.png`)

Choose **Next** → **Select dataset and criteria** and `contoso-eval-en-dev12` **version 1**, not Generate data or an edited upload. Use the same Judge and keep exactly Relevance and TaskAdherence.

In Optimizer **Criteria**, if **No custom evaluators available** appears, switch **Custom only OFF** or click **View built-in evaluators**. Selecting either evaluator row opens its **Configure...** dialog: set **Relevance Threshold 4** or **TaskAdherence Threshold 1**, then **Apply**. TaskAdherence remains binary pass 1. Do not create custom evaluators to bypass this filter; record settings without claiming private rubric-version pins.

Create an optimization run

Target

Dataset

Criteria

Review

Target

Dataset

Criteria

Review

Choose an existing dataset to test each candidate against. Optimization runs the same dataset across every candidate so results are comparable.

Generate data

Select dataset and criteria

Generate a dataset and evaluator from production traces or a prompt.

Choose an existing dataset, then configure evaluation criteria.

Back

Next

Select dataset

Search by file name

Create dataset from traces

Upload dataset

Name	Version	Created by	Created on
eval-data-2026-09-30_144335_a656c UTC	1		9/30/26, 11:43:35 PM
contoso-eval-en-dev12	1		9/30/26, 11:20:29 PM

1-2 of 2

Preview

Dataset

contoso-eval-en-dev12

Fields

contextground_truthquery

Dataset preview (Top 5 rows)

context	ground_truth	query
[ATLAS-SUB-001] Subscription-3. Canceling a subscription stops the next automatic renewal; access continues until the end timestamp of the current	{"answer": "By 'reduce', do you m stopping automatic renewal, or on a lower plan from the next rer Those are different choices, so I clarify your intention first. This c has not changed the	Next month's cost is a concern. Please reduce my subscription. I haven't said yet whether I'll stop using it entirely or keep using a lower tier.
[ATLAS-REF-001] Refund-2. A first purchase of a prepaid monthly subscription made before 2026-09-01 00:00:00 KST retains the previous 7 x 24-hour application window,	{"answer": "The instruction in the notice does not change the polici 28 purchase retains the previou window, whose final timestamp September 4 at 08:00 KST. Septe past the deadline; the month of	My first monthly payment was on August 28, 2026 at 08:00 KST. I haven't used it, and I'm contacting you on September 6. The notice I attached says 'Everyone
[ATLAS-REF-001] Refund-1. A first purchase of a prepaid monthly subscription made at or after 2026-09-01	{"answer": "The standard condi both no paid-credit use and no p production jobs. If even one paic job was run, the situation you de	I used the service for just one day after my first monthly payment on September 8, 2026. Paid-credit usage is

Reuse the registered dataset. Confirm its version and 12-case scope. A five-row preview does not change the full dataset count. View full-size image (file: [web/assets/portal/en/08-optimizer-dataset.png](#))

For a fresh authorized job, review and submit once. For the recorded job, follow its existing ID and wait **at most 60 minutes** within the approved scope. One job can make multiple internal calls; it is not one free model request. If blocked or timed out, record the status and stop rather than submitting a duplicate.

For the succeeded job, read **original versus candidate**, **per-evaluator results**, and **View changes**. The Optimizer's **0–1 ranking** is not the same as Relevance's 1–5 score or a pass percentage. A higher ranking does not establish policy improvement or replace direct reevaluation.

Microsoft Foundry / proj-fev11-dea3cec5

Search with AI (96K)

New Foundry

Home

Discover

Build

Operate

Manage

Docs

Create

Agents

Models

Services

Tools

Knowledge

Memory

Guardrails

Data

Optimize

Evaluations

Fine-tune

opt_e44bcf5701a348deb62a1cd4f9cb3910

Succeeded

Download JSON

Promote candidate

Run details

Improvement

Goal: Improve quality

candidate_1

Avg. score: 0.646

Baseline: 0.635

Task-weighted average (0-1)

View changes

Download config

Run overview

Candidates: 1

Best score: 0.646

Eval model: lab-judge-luna-dea3cec5

Dataset: contoso-eval-en...

Elapsed time: 9m 36s

Total tokens: 264.3K

View details

Additional settings

Score comparison

Task-weighted average (0-1)

candidate_1: 0.646

baseline: 0.635

View overall score details

Candidate results

Iteration	Candidate	Optimizations	Score	Score delta	Avg tokens	Avg latency	Changes	Config	Score details
0	baseline	--	0.635	--	2.8K	6.1s	Version 1	↓	evalrun_f9836666...
1	candidate_1	System prompt	0.646	+0.010	3.3K	8.1s	View changes	↓	evalrun_0424454...

1-2 of 2

Prev Next

Read the candidate's evidence. Check actual job/candidate IDs and per-evaluator results. A recommendation is not a promised improvement or production approval. View full-size image (file: web/assets/portal/en/09-optimizer-results.png)

Changes: candidate_1

★ Best candidate

×

Comparing candidate_1 against baseline configuration.

Baseline

```

1 {
2   "agentName": "contoso-eval-en",
3   "agentVersion": "1",
4   "system_prompt": "You are a customer support assistant for
Contoso Atlas Cloud. Answer in English using the customer's
inquiry and the supplied policy context.\nFor a policy
question, if the knowledge_base_retrieve knowledge tool is
connected, use it to check the evidence before answering. If
no tool or context is available, state the limitation rather
than inventing policy details.\nThe response must be a single
JSON object containing only the following four fields, with
no other text or Markdown.\n{\n  \"answer\": \"Answer text\",
  \"citations\": [\"Document identifier\"], \"route\":
  \"answer\", \"needs_human\": false\n}\nroute must be one of
answer, clarify, escalate, or refuse. citations is an array
of strings. needs_human is true only when route is escalate
and false otherwise.\n",
5   "skills": [],
6   "tools": []
7 }
```

Candidate: candidate_1

```

1 {
2   "agentName": "contoso-eval-en",
3   "agentVersion": "1",
4   "system_prompt": "You are a customer support assistant for
Contoso Atlas Cloud. Answer in English using the customer's
inquiry, any supplied policy context, and the evidence rules
below.\n\nYour final response must be a single valid JSON
object and nothing else. Do not include Markdown, code
fences, tool traces, arrays outside the JSON object, or
citation footnote markup inside the answer text.\n\nRequired
JSON shape:\n{\n  \"answer\": \"Answer text\", \"citations\":
  [\"Document identifier\"], \"route\": \"answer\",
  \"needs_human\": false\n}\n\nField rules:\n- Include exactly
these four fields: answer, citations, route, needs_human.\n-
route must be one of: answer, clarify, escalate, refuse.\n-
needs_human must be true only when route is \"escalate\"; it
must be false for answer, clarify, and refuse.\n- citations
must be an array of document identifiers only, such as
\"ATLAS-SLA-001\", \"ATLAS-SUB-001\", or \"ATLAS-REF-001\".
Do not put URLs, inline citation markers, or quoted snippets
in citations.\n- The answer field should be concise but
complete, and should not contain unsupported policy
claims.\n\nEvidence and tool-use rules:\n- For any policy
question, if the knowledge_base_retrieve tool is connected,
call it before answering.\n- Use the retrieved policy
content, supplied policy context, and the embedded policy
```

< 1 of 1 >
Close

Inspect the instruction diff. Check policy reasoning, clarification, authority boundaries, and the English response contract. Longer instructions are not inherently better. View full-size image (file: [web/assets/portal/en/10-optimizer-changes.png](#))

Keep the same model, tools, connections, and response contract. Do not remove tools just because an exported candidate configuration omits them. If the proposed change is not instruction-only, report the mismatch rather than claiming a controlled improvement.

Actual Promote created contoso-eval-en v2, with exact candidate instructions and unchanged model/knowledge tool verified. Active-version changes affect **all channels**, so this was restricted to the **isolated, unpublished lab, never production**. After HOLD, active v1 was restored; v2 remains available with its evaluation evidence. Do not promote again or delete it.

Discuss: Which instruction changed, which failure might it address, and what could regress? The Optimizer's internal result selects a candidate for verification; it does not replace the next direct Evaluation.

Completion signal: The real job succeeded and verified lab v2 exists. Its separate native SDK evaluation completed as recorded in step 06; neither promotion nor the Optimizer UI score is production approval.

06. Reevaluate, compare, and finish

What: A real Foundry managed evaluation of v2 under the existing definition. **Why:** The portal candidate form currently fails before submission. **How:** Run one Azure AI Projects/OpenAI Evals SDK helper command, then read and compare the remote runs in Foundry.

Observed portal blocker: Add run → Pin v2 → Individual turns led to **Configure agents: Config required → Add custom prompt / User prompt**. Even `{{item.query}}` failed Submit client-side with Unable to create data source configuration from item schema. **No candidate remote run was submitted by that attempt.** Do not create another evaluation definition to bypass it.

From the repository root in bash/zsh/WSL2, replace **all four placeholders** below. The operator supplies the project endpoint and subscription. Copy **your own baseline's** evaluation/run IDs from its portal URL or **Raw JSON**. Do not copy the IDs from this guide's recorded rehearsal.

```
export AZURE_AI_PROJECT_ENDPOINT="OPERATOR_PROJECT_ENDPOINT"
export AZURE_SUBSCRIPTION_ID="OPERATOR_SUBSCRIPTION_ID"
export FOUNDRY_EVALUATION_ID="YOUR_EVALUATION_ID"
export FOUNDRY_BASELINE_RUN_ID="YOUR_BASELINE_RUN_ID"
```

Prerequisites: Activate the prepared venv with `source .venv/bin/activate`, install locked dependencies with `python -m pip install -r requirements.lock`, then complete `az login` and the [account/subscription verification](#). Do not proceed with a different account or missing prerequisites.

Run this one SDK operation. It clones the baseline `data_source`, retains existing evaluation criteria, changes only the target version in the run input, verifies the same model/tools, and submits a **real Foundry evaluation run under the same evalID**. It is **not a local Judge or bespoke cloud scoring**.

```
python scripts/add_foundry_eval_run.py --endpoint "$AZURE_AI_PROJECT_ENDPOINT" --subscription
"$AZURE_SUBSCRIPTION_ID" --evaluation "$FOUNDRY_EVALUATION_ID" --baseline "$FOUNDRY_BASELINE_RUN_ID" -
-version 2 --name candidate-v2 --out .lab/foundry-evaluations/candidate-v2.json
```

The first submission persists a private **receipt**. Repeat the identical command and `--out` path only to collect that same run. If the receipt is missing but the run name already exists remotely, the helper stops instead of creating a duplicate. Open the existing run or recover its original receipt; do not change names, paths, or evaluation definitions to chase a favorable result.

Verified SDK result: under `eval_94feef6f6f644fabb22a5680f5f24fb1`, candidate run `evalrun_f3b710fc835d444fb8aa0d2bb7797bdf`, `contoso-eval-en v2`, is **Completed: 12 rows, 11 passed / 1 failed / 0 errored**, versus v1's **10 passed / 2 failed / 0 errored**. The helper confirmed the same dataset, evaluators, Judge, model, and tools; only instructions/version changed. This is a real Foundry run, not the Optimizer's internal score.

Open **Build → Evaluations →** `contoso-en-learning-loop` **→ Evaluation runs**, select both run checkboxes, then **Compare runs**. **Explicitly select the original** `contoso-eval-en` **in the Baseline dropdown**. The page defaults to the first selected row—`candidate-v2` in this recorded case—so verify direction before reading differences/tests. The values below use original v1 as Baseline; private rubric-version pinning remains unproven.

Actual comparison, n = 12	Original <code>contoso-eval-en</code> v1	<code>candidate-v2</code>
All-criteria pass / fail / error	10 / 2 / 0	11 / 1 / 0
Relevance passes; mean (1–5)	10/12; 4.4167	11/12; 4.3333
TaskAdherence passes; binary mean	12/12; 1.0	12/12; 1.0
Relevance rows 1 / 2 / 6 / 11	3 / 5 / 5 / 3	4 / 4 / 4 / 3
Latency p50 (ms)	5,891.09	7,287.52
Latency p95 (ms)	8,817.33	16,038.35
Agent tokens	35,187	43,751

Microsoft Foundry / `proj-fev11-dea3cec5`

Search with AI (🔍)

New Foundry ☒ Home Discover **Build** Operate Manage Docs

Create

- Agents
- Models
- Services
- Tools
- Knowledge
- Memory
- Guardrails
- Data

Optimize

- Evaluations**
- Fine-tune

← Statistically compare results

Total runs selected: 2Analysis method: PairedTTest

Baseline * ⓘ
`contoso-eval-en`

Name	<code>contoso-eval-en</code>	<code>candidate-v2</code>
TaskAdherence	1.00 Baseline	1.00 Inconclusive
Relevance	4.42 Baseline	4.33 Inconclusive

1-2 of 2

PrevNext

Actual native comparison. Set **Baseline → contoso-eval-en**, not the first-selected `candidate-v2`. PairedTTest is Inconclusive for both metrics. The higher pass count does not erase the lower Relevance mean or higher latency/tokens: adoption remains HOLD. View full-size image (file: `web/assets/portal/en/20-evaluation-comparison.png`)

Decision executed: HOLD adoption; active version restored to 1. After explicitly selecting the original Baseline in Compare, keeping the original in this lab uses **Agent Details → Agent configuration → Active version → Edit → Version 1**; the operator already completed this. **Keep candidate v2, both evaluation runs, and receipts—do not delete evidence.** Lower mean, row regressions, higher latency/tokens, and Inconclusive tests do not prove overall improvement; do not change policy facts to chase scores.

Closing explanation: Explain HOLD using the extra passing row, rows 2/6 regressions, lower mean, higher latency/tokens, and Inconclusive tests—not just the Optimizer's +0.010 or native pass count.

Completion signal: Actual run IDs, measured tradeoffs, and the HOLD decision are recorded. Further review and new representative cases belong to the next authorized improvement cycle, not an extra required lab or retries until favorable.

Reused dev12 does not establish statistical significance or production approval. **No production Publish was performed and no production channels/traffic are configured.** Foundry still automatically provides **RBAC-only Responses/preview endpoints without Publish**; unpublished does not mean no endpoints.

Finish: Give the operator the actual run/job/version references, dataset version/hash, observed failures, and remaining blockers. Agree who owns ongoing resource costs and authorized cleanup; closing the browser does not stop charges. Share only approved synthetic examples and redacted summaries, never credentials or private environment files.

[Data contract](#) · [Facilitator and recovery](#) · [Operator handoff](#) · [Measured verification and sources](#)

Facilitator guide · the evaluation learning loop

[Six participant steps](#) · [Operator prerequisites](#) · [Measured verification](#)

Teach **evaluate** → **learn from failure** → **improve** → **reevaluate**. Require an actual question, answer, evaluator reason, and policy reference—not just “the score improved.” Representative company-owned tasks, policy edge cases, and known failures make domain evaluation useful; this workshop demonstrates the method with synthetic Contoso data only.

Before the session

The operator prepares the environment, sample agent, and SDK prerequisites. Participants start with the dataset. Use **Microsoft Foundry managed Evaluation, instruction-only Agent Optimizer**, then **one SDK command for the real candidate run and portal Compare runs**. It is not a local Judge. Keep exactly two evaluators and six steps.

- Confirm the project and authorization from the [operator handoff](#). English preparation is verified: `contoso-eval-en version 1`, `lab-agent-dea3cec5`, existing **read-only knowledge connection preserved**. The Korean counterpart is operator-prepared `contoso-eval-ko`; no new Korean execution or measured results are claimed.
- English `contoso-eval-en-dev12 version 1` is registered from unchanged `data/en/optimizer/dev.jsonl`, 12 cases. Retain its SHA-256. Korean uses `data/optimizer/dev.jsonl`, also 12 cases, with separately confirmed `contoso-eval-ko-dev12` registration.
- Distinguish the one-fixture `gpt-6-luna / 2026-09-22` **compatibility check**, the completed configuration pilot, and the clean corrected comparison. Direct Responses/pinned prompt-agent v2 failed with HTTP 500 here, not everywhere. Keep `gpt-4.1-mini / 2025-04-14` as the verified agent runtime.
- Use `lab-judge-luna-dea3cec5` for the `gpt-6-luna` Judge and `lab-planner-dea3cec5` for the `gpt-5.5 / 2026-04-24` instruction generator. The generator is not the agent being optimized or the Judge. `gpt-6-luna` is not on the [supported optimization-model list](https://learn.microsoft.com/azure/foundry/agents/concepts/agent-optimizer-overview#models) (<https://learn.microsoft.com/azure/foundry/agents/concepts/agent-optimizer-overview#models>).
- Confirm Agent Optimizer access and one candidate. If a required role is unverified, do not lead participants through a submission as though support were established.
- Complete [SDK prerequisites](#): prepared `venv`, `requirements.lock`, `az login` and verified account/subscription, operator project endpoint/subscription, and baseline URL/Raw JSON IDs. Do not treat a missing helper/dependency as permission to change the scoring method.

Screenshots illustrate the **English UI and English-data rehearsal**, not a participant's current result or a Korean run. Use them to locate controls. Read current measured status from verification; do not turn historical images or a reading checkbox into execution evidence.

Canonical baseline complete: contoso-en-learning-loop has **12 rows, 10 passed / 2 failed / 0 errored; Relevance 10/12, binary TaskAdherence 12/12** on pinned v1 and the same Judge. The earlier contoso-en-baseline-luna-judge remains a **CONFIGURATION PILOT**. Actual IDs are in [step 03](#); do not mix pilot aggregates into the canonical comparison.

Optimizer succeeded with one candidate: opt_e44bcf5701a348deb62a1cd4f9cb3910, same dataset, gpt-5.5 generator / gpt-6-luna Judge, no model comparison. UI **0.635 → 0.646, UI-reported +0.010, 264,260 reported tokens** are optimization results, not the final direct comparison. Promote created verified lab-only v2 with exact candidate instructions and unchanged model/knowledge tool.

Promotion warning: active-version changes affect **all channels** and belong only to the **isolated, unpublished lab, never production**. After HOLD, active v1 was restored; candidate v2 remains available. Use explicit versions and the saved baseline rather than assuming “latest” is active.

Native SDK execution verified; adoption HOLD: candidate evalrun_f3b710fc835d444fb8aa0d2bb7797bdf, v2, completed under eval_94feef6f6f644fabb22a5680f5f24fb1 with **12 rows, 11 passed / 1 failed / 0 errored**, versus v1 **10 passed / 2 failed / 0 errored**. Matched dataset/evaluators/Judge/model/tools leave only instructions/version changed. Relevance passes rose 10/12→11/12 but mean fell **4.4167→4.3333**; TaskAdherence stayed **12/12, binary mean 1.0**. The portal's **PairedTTest is Inconclusive for both metrics**, not proof of overall improvement.

Observed latency worsened: **p50 5,891.09→7,287.52 ms; p95 8,817.33→16,038.35 ms**. Agent tokens rose **35,187→43,751**. Keep pinned v1 pending further review/new representative cases; do not present Optimizer +0.010 or the native pass-count increase alone as success.

Keep the six-step path explicit

Participant step	Portal action and exact choice	Completion signal
01 Dataset	For a fresh agent still latest v1: New experience → Build → Evaluations → Create, Pin currently latest, reselect 1 target ; Individual turns / One time / Existing dataset. For the recorded promoted agent, inspect its existing v1 baseline	Correct pinned version/data, n = 12 and hash; never confuse latest v2 with baseline v1
02 Criteria	Query only, no custom override; retain Relevance threshold 4 + TaskAdherence binary pass 1 , generic	Two criteria with distinct scales and settings recorded

Participant step	Portal action and exact choice	Completion signal
	TaskAdherence Threshold 1 if shown; explicit Judge	
03 Baseline	Open the completed contoso-en-learning-loop baseline; do not resubmit	12 rows and canonical metric results; separate from the pilot
04 Analysis	Read Detailed metrics result reasons; conversation_id → User view opens the question/JSON response, not a Judge panel	Actual response, Relevance.reason / TaskAdherence.reason, and policy reference connected
05 Optimization	Recorded job succeeded; inspect original/candidate and diff. For fresh setup: Custom only OFF → built-in rows → Configure... → Apply, Relevance 4 / TaskAdherence 1	One real candidate, verified lab v2; latest-active warning respected
06 Comparison	Reuse the completed scripts/add_foundry_eval_run.py run's receipt; Evaluation runs → both checkboxes → Compare runs → Baseline dropdown: original contoso-eval-en, not first-selected candidate-v2	Verified execution, measured tradeoffs, PairedTTest Inconclusive both; HOLD adoption, no production approval

The dataset preview shows the **first five rows**, not the total. In the observed English wizard, the matching schema **auto-resolved Field mapping**, then proceeded to **Configure agents** → **leave custom prompt override unset** → **Criteria**. If a schema-dependent mapping screen appears, use **query** → **query**. Never append context or ground_truth to generation. Use Upload new dataset only when the authorized project lacks the intended registration.

Remove extras from the **23 suggested evaluators**. **Relevance is 1–5, threshold 4. TaskAdherence is Binary Pass/Fail, raw 0/1, pass 1**; set its generic UI **Threshold to 1** if shown. The [official definitions](https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/agent-evaluators) (<https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/agent-evaluators>) confirm this distinction. The pilot's TaskAdherence Threshold 4 yielded score=1, passed=true; the generic setting did not define a meaningful 1–5 rubric.

Use lab-judge-luna-dea3cec5. Catalog links showed relevance **v14** / task_adherence **v17**, but service criteria had evaluator_version **empty/default**: do not teach those links as fully pinned private rubric versions.

The corrected submission confirmed **Relevance** response={{sample.output_text}} and **TaskAdherence** response={{sample.output_items}}. UI defaults also showed query={{item.query}} and TaskAdherence's tool_definitions={{sample.tool_definitions}}. Inspect **Raw JSON** and **do not override generated mappings** with the earlier UI output-text value for TaskAdherence. They are service bindings, not extra dataset fields or generation inputs.

For a fresh baseline, **Pin currently latest** must resolve to v1 before changes, then reselect the cleared checkbox. The recorded candidate is already promoted v2: v1 remains in the saved baseline run, and step 06 explicitly targets version 2 with the SDK. `gpt-4.1-mini` retires in Azure on **2027-04-14**; public **Deprecated** status may restrict fresh subscriptions.

Optimizer filter, not missing evaluator support: If Criteria says **No custom evaluators available**, switch **Custom only OFF** or click **View built-in evaluators**. Selecting each row opens **Configure...**; set Relevance **4** or TaskAdherence **1**, then **Apply**. This differs from the direct Evaluation wizard's Edit/Update controls and requires no custom evaluator.

Discuss evidence, not a desired score

Spend a few minutes after the baseline, after inspecting the Optimizer diff, and after direct reevaluation. These are discussions of existing results, not additional submissions.

Discuss	Ask for
Baseline	Run ID; n = 12 ; each metric's scored/total, pass count, missing/errors; a failure or lowest-score case and a good case if one exists
Candidate	Actual job/candidate ID; original versus candidate per-evaluator results; changed instruction and plausible regression
Reevaluation	New direct run ID under the same definition; same dataset/version/hash and recorded evaluator/Judge configuration; actual before/after answers/reasons and rubric-version limits

Use this unfilled discussion prompt:

"In corrected run , *metric* __ *uses scale* __ *and pass rule* __ (*Relevance* ≥ 4 ; *TaskAdherence* =1). It has __ scored rows out of 12, __ passes, and __ missing/errors. Answer __ conflicts/agrees with policy __ because . Our next decision is __."

If no failure or no good case exists, report that observation. Do not manufacture one, weaken the baseline, copy a screenshot's values, or demand an improved candidate.

Explain what the metrics do not prove

Completed is not a quality pass. Relevance's **1–5 score is not percent accuracy**. TaskAdherence is **binary 0/1**, not a five-point score: 1 means Pass, 0 means Fail. Missing/errors are not scored zeros. Report counts, coverage, reasons, and each metric's proper result type.

Relevance asks whether the response addresses the question. TaskAdherence assesses the agent's task instructions and constraints. Neither certifies every policy, citation, authority boundary, or safety property. Some evaluators do not use the reference columns; participants still compare the actual answer/reason with the original policy and reference answer.

The application values cases such as policy-date boundaries, necessary clarification, unsupported approval claims, and safe handling of restricted requests. Use an actual row to show why a generic public benchmark cannot substitute for those domain decisions.

Canonical v1 baseline Relevance failures: `atlas-dev-001` gave vague subscription-reduction guidance where the reference required explicit clarification; `atlas-dev-011` honestly called an undocumented feature unverified but was judged incomplete. TaskAdherence passed both. These are not candidate regression findings; do not change policy facts or remove justified uncertainty to chase Judge scores.

Screenshot 19 is **conversation_id → User view** for `atlas-dev-001`: question and actual JSON response, **not an inline Judge-reason panel**. Return to **Detailed metrics result** for `Relevance.reason / TaskAdherence.reason`, then read the reference policy. Do not invent quotations or reasons not present in the run.

Keep model/tools, language, 12 cases, dataset/version/hash, **Relevance 4 / TaskAdherence binary pass 1**, and Judge fixed. The SDK helper clones baseline `data_source` and reuses existing criteria, changes only target version, verifies model/tools, and submits a real Foundry run in the same evalID. Repeating the identical command and receipt path collects safely rather than resubmitting. No bespoke cloud scoring or new definition is involved.

The portal **Add run** attempt for Pin v2 / Individual turns required **Configure agents → Config required → Add custom prompt / User prompt**. Even `{{item.query}}` failed client-side with `Unable to create data source configuration from item schema`; **no remote candidate run was submitted by that attempt**. Use the SDK route, not a new evaluation definition or handcrafted score. Private service rubric versions remain unproven pinned.

Recorded row changes: Relevance row 1 improved **3→4**, rows 2 and 6 declined **5→4**, and row 11 stayed **3** while maintaining honest uncertainty. A threshold crossing can raise pass count even as the mean falls. Native **PairedTTest: Inconclusive** for Relevance and TaskAdherence does not prove improvement or equivalence. Reused dev12 is not statistical significance, independent generalization, or production approval.

Recover without changing the experiment

Inspect the existing run/job and its exact inputs before doing anything else. If submission status is uncertain, locate the request with the operator rather than creating another. Keep failures and missing rows visible.

Symptom	Safe response
Wrong account/project, 401/403, network restriction	Operator checks identity and authorized access; do not broaden shared permissions or bypass restrictions
Catalog model is visible but runtime/evaluator fails	Record role, deployment, error, and request/run ID. Catalog visibility is not runtime proof
gpt-6-Luna Responses/pinned agent v2 HTTP 500	Keep that runtime unverified here and the agent on verified gpt-4.1-mini ; do not undo the distinct native Relevance Judge capability result or claim model-wide lack of support
Upload preview has five rows	Confirm the unchanged local file has 12 and the correct registered dataset is selected
Wrong schema or input	Use the exact three-column file; ground_truth stays a JSON string and only query reaches the agent
Pinning leaves no selected target	The version change clears the row checkbox; reselect it and confirm 1 target before Next
UI mapping differs from submitted TaskAdherence mapping	Keep the generated sample.output_items binding; inspect Raw JSON, do not override it to sample.output_text
Compare runs is disabled	In the same definition's Evaluation runs table, select both baseline and candidate rows; wait for completion before interpreting results
Comparison direction looks reversed	Baseline defaults to the first selected row (candidate-v2 here). Explicitly choose original contoso-eval-en in the Baseline dropdown before interpreting deltas/tests
Partial evaluation or missing score	Report n = 12 with scored/missing/error counts; do not remove rows or fabricate reasons
Optimizer Max candidates disabled	Clear Model and select Instruction only; Tool description off, Max candidates 1
Optimizer Criteria: No custom evaluators available	Switch Custom only OFF or View built-in evaluators; select each row → Configure... → Relevance 4 / TaskAdherence 1 → Apply
Portal candidate Submit shows item-schema error	No remote run was created by that failed form. Use the single SDK command; never create another definition as a workaround

Symptom	Safe response
SDK receipt exists or run is still processing	Keep the same <code>.lab/foundry-evaluations/candidate-v2.json</code> and repeat the identical command to collect. Do not delete/change the receipt to force another submission
SDK identity/model/tool validation fails	Stop and check operator context; do not bypass validation, switch models, or substitute local Judge scores
First-run Optimizer shows a tax-agent benchmark	Label it product example, not Contoso results; use Optimize my agent → Agent
Optimization exceeds 60 minutes	Record the actual job status and stop waiting/submitting; operator owns follow-up and remaining costs
Candidate changes models/tools, or evaluator contract differs	Do not claim instruction-only improvement; record the mismatch and keep baseline
Optimizer or reevaluation is blocked	Record unrun work. Do not substitute a different feature, handwritten candidate, or repeated favorable run

Correct common misconceptions

Claim	Correction
"Upload or Review succeeded, so evaluation completed."	Preparation is not submission. Require real evaluation/run IDs and final results
"Relevance 4/5 means 80% accuracy."	It is an ordinal Relevance score, not an accuracy estimate
"TaskAdherence 1 means a poor one-out-of-five answer."	No: it is binary Pass = 1 . Do not apply a five-point scale or threshold 4
"The three roles must all use the same model."	Runtime, Judge, and instruction generator have separate support constraints
"The candidate ranks higher, so promote to production."	Never. Latest-active promotion affects all channels and is restricted here to the isolated unpublished lab; direct evaluation and production approval are separate
"No improvement means we should keep trying."	Report no improvement and keep baseline. Do not rerun until favorable
"English screenshots prove the Korean path ran."	They illustrate UI only; language-specific execution needs its own evidence
"These private-domain lessons used private customer data."	Contoso is synthetic. Real company data requires separate authorization, privacy controls, and representative review

Finish with an honest result

Ask for the shared evaluation definition ID, distinct baseline/candidate run IDs, optimization job/candidate ID, pinned agent versions and model roles, dataset version/hash, recorded criteria/configuration and service-version limits, per-metric comparison, a case-level explanation, regressions, and the decision. In-progress, failed, or unrun work stays explicitly unfinished.

Final workshop decision: HOLD adoption; keep pinned baseline v1. Pass count rose, but mean Relevance fell, rows regressed, and latency/token usage increased; neither metric's PairedTTest is conclusive. Further review/new representative cases belong to a later authorized cycle, not an added required lab or reruns until favorable. No production approval or statistical significance is claimed.

HOLD was executed: after setting the original Baseline explicitly in Compare, the operator used **Details → Agent configuration → Active version → Edit → Version 1** to restore active v1. Keep candidate v2, both evaluation runs, and receipts; do not delete comparison evidence. No production Publish or production channels/traffic were configured. **RBAC-only Responses/preview endpoints still exist automatically without Publish;** do not teach “unpublished” as “no endpoints.”

Name the cost/cleanup owner and remaining blockers. Share only approved synthetic examples and redacted summaries. For a future company-owned evaluation, separately authorize data use and select representative tasks and known failures; do not claim this 12-case exercise certifies production.

Operator prerequisites · an isolated NCUS workshop


[Participant path](#) · [Facilitator](#) · [Measured verification](#)

Prepare the project and sample agent **before** participants begin. The participant path is dataset preparation, managed Evaluation, failure analysis, instruction-only Agent Optimizer, and matching reevaluation—not infrastructure deployment.

Start with the prepared lab environment

The operator has provisioned a **new, isolated North Central US (NCUS) resource group** for this rehearsal. Confirm its actual identity, ownership, and readiness from the private preparation evidence. Do not recreate it merely because the documentation changed or use a shared production environment.

Keep three claims separate: **resources exist**, **the intended runtime works**, and **task quality is acceptable**. A successful deployment proves neither agent execution nor managed-evaluator support. Record measured preparation and run outcomes in [verification](#), not by copying a screenshot's status.



Prepared resource group. The visible Failed count includes a separate organization-policy deployment failure involving its logging target. Do not hide it, repair shared policy as part of this workshop, or confuse it with the lab deployment's status. View full-size image (file: web/assets/portal/en/00-resource-group.png)

Verify three separate model roles

Use the **actual deployment names below**, not model family names in deployment fields, after confirming the authorized project. The separate English baseline is now `contoso-eval-en v1`, verified working with its existing **read-only knowledge connection preserved**. Keep this role mapping unchanged through the comparison.

Role	Required decision
Agent runtime	<code>lab-agent-dea3cec5</code> → <code>gpt-4.1-mini / 2025-04-14</code> , verified on <code>contoso-eval-en v1</code> . Preserve its read-only knowledge connection; <code>gpt-6-luna Responses/Agent verification</code> failed here
Evaluation Judge	<code>lab-judge-luna-dea3cec5</code> → <code>gpt-6-luna / 2026-09-22</code> , verified in the actual managed baseline and candidate evaluations
Optimizer instruction generator	<code>lab-planner-dea3cec5</code> → <code>gpt-5.5 / 2026-04-24</code> . The official optimization-model list

Role	Required decision
	https://learn.microsoft.com/azure/foundry/agents/concepts/agent-optimizer-overview#models does not list gpt-6-luna

Observed capability evidence: gpt-6-luna version 2026-09-22 is GA in NCUS with **GlobalStandard**. The new deployments are lab-agent-luna-dea3cec5 and lab-judge-luna-dea3cec5; Chat Completions returned **READY**. An actual native Foundry **Relevance** evaluation completed with **passed 1 / total 1 / errors 0** on **one authored compatibility fixture**: evaluation eval_589a069bfd8343cf980b4b88f57771e8, run evalrun_e662077a904543d6bd1420b202a1b078.

That earlier record verifies Judge capability, **not agent quality**. It is separate from the configuration pilot below and from the clean corrected comparison; do not merge their counts or IDs.

Runtime limit: Direct Responses and a **pinned** gpt-6-luna **prompt-agent v2** both returned **HTTP 500 in this environment**. Do not claim model-wide lack of support; runtime verification failed here. Keep the verified gpt-4.1-mini agent until the provider issue is resolved, without silently selecting the experimental agent.

The gpt-4.1-mini / 2025-04-14 Azure retirement date is **2027-04-14**. A **Deprecated** label in public documentation can restrict fresh subscriptions; it does not erase the existing deployment's observed usability. Verify the actual account and existing deployment before handoff rather than promising access to every new subscription.

The user requested gpt-6-luna **where supported**, acknowledging that the three roles need not support it equally. Catalog visibility, deployment success, and a chat response do not verify every API/runtime. Do not silently switch a role after baseline or report a model change as an instruction-only gain.

The optimization model **writes candidate instructions**; it is separate from the agent model being optimized and the Judge that scores responses. Use the same verified Judge for direct Evaluation, Optimizer evaluation, and direct reevaluation.

Microsoft Foundry / proj-fev11-dea3cec5

Search (3K)

Create

Agents

Models

Services

Tools

Knowledge

Guardrails

Data

Optimize

Evaluations

Fine-tune

Models

DeploymentsBatch jobs

Serverless deploymentsManaged compute deployments

SearchPTU CalculatorDeploy

Name	Model	Version	Deployment status	Deployment type	Created on ↓	Created by
☒ lab-judge-luna-dea3cec5	gpt-6-luna	2026-09-22	Succeeded	Global Standard	9/30/26, 11:19:02 ...	
☐ lab-agent-luna-dea3cec5	gpt-6-luna	2026-09-22	Succeeded	Global Standard	9/30/26, 11:18:59 ...	
☐ lab-embedding-dea3cec5	text-embedding-3...	1	Succeeded	Global Standard	9/30/26, 9:58:45 PM	
☐ lab-planner-dea3cec5	gpt-5.5	2026-04-24	Succeeded	Global Standard	9/30/26, 9:58:44 PM	
☐ lab-judge-dea3cec5	gpt-5.4-mini	2026-03-17	Succeeded	Global Standard	9/30/26, 9:58:42 PM	
☐ lab-agent-dea3cec5	gpt-4.1-mini	2025-04-14	Succeeded	Standard	9/30/26, 9:58:40 PM	

1-6 of 6< PrevNext >

Deployments are role inputs, not runtime proof. Check model version and deployment name against the private handoff. A listed deployment alone does not establish Agent, Responses, evaluator, or Optimizer compatibility. View full-size image (file: web/assets/portal/en/02-model-deployments.png)

Prepare the project and sample agent

Use the official [direct agent-evaluation prerequisites](https://learn.microsoft.com/azure/foundry/observability/how-to/evaluate-agent#prerequisites) (https://learn.microsoft.com/azure/foundry/observability/how-to/evaluate-agent#prerequisites) and [prompt-agent Optimizer prerequisites](https://learn.microsoft.com/azure/foundry/agents/quickstarts/quickstart-optimize-prompt-agent#prerequisites) (https://learn.microsoft.com/azure/foundry/agents/quickstarts/quickstart-optimize-prompt-agent#prerequisites). They define the supported project, agent, model, and access requirements. Do not require participants to run a whole legacy bootstrap sequence.

Preparation	Evidence required before handoff
Intended account and project	Actual signed-in tenant/account, subscription, resource group, and project match the authorized lab. Check portal and any operator CLI/SDK identity separately
Prepared Contoso agent	Existing contoso-eval-en version 1 on lab-agent-dea3cec5; read-only knowledge connection preserved. Check the English four-key response contract without changing the baseline configuration

Preparation	Evidence required before handoff
Runtime	Actual response from that pinned agent through the intended runtime; a model chat smoke is not a substitute
Managed Evaluation	contoso-en-learning-loop is Completed : 12 rows, 10 passed / 2 failed / 0 errored; Relevance 10/12 and binary TaskAdherence 12/12. Keep the pilot separate
Agent Optimizer	Optimize Preview access, Instruction-only target, Max candidates 1, lab-planner-dea3cec5, and the same Judge lab-judge-luna-dea3cec5
Dataset	contoso-eval-en-dev12 version 1 , registered through the native Evaluation wizard from unchanged data/en/optimizer/dev.jsonl, exactly 12 rows; retain the original file SHA-256

Configure policy access before class, then hold it constant. Participants neither build a separate knowledge system nor repair infrastructure to obtain a better score. If the prepared agent lacks needed evidence, report the setup blocker rather than appending reference answers to its input.

For a fresh baseline still at v1, select **Pin currently latest**, reselect the cleared checkbox, and confirm **1 target before Next**. The recorded agent retains candidate v2, but **active version is restored to 1**. Use explicit versions and saved evaluation IDs; do not equate “latest” with active. The failed gpt-6-luna v2 was another agent.

In the English pilot, **Existing dataset** with matching schema auto-resolved **Field mapping**, then proceeded to **Configure agents → custom prompt override unset → Criteria**. For the corrected run, remove extras from the 23 suggestions and retain exactly two evaluators: **Relevance Threshold 4, TaskAdherence binary pass 1 (Threshold 1 if a generic control is shown)**.

The [official agent-evaluator definitions](https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/agent-evaluators) (https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/agent-evaluators) identify TaskAdherence as **Binary Pass/Fail, raw 0/1**, not 1–5. The pilot UI accepted TaskAdherence Threshold 4, but actual score=1 returned passed=true; that generic setting was not meaningful. Preserve the pilot and start a clean corrected comparison instead of changing its historical definition.

Use the English dataset with contoso-eval-en. The Korean counterpart is operator-prepared contoso-eval-ko **v1**, using data/optimizer/dev.jsonl, contoso-eval-ko-dev12, and a Korean response contract. **Korean was not re-executed in this rehearsal; no new Korean measured results or registration are claimed**. Verify its preparation separately rather than relabeling English evidence.

Retained CONFIGURATION PILOT — portal execution, not the corrected baseline:

Recorded item	Value
Pilot evaluation name	contoso-en-baseline-luna-judge

Recorded item	Value
Pilot evaluation definition ID	eval_06d9c6cb93df4c7bad2bc3a62da9bc90
Pilot run ID	evalrun_98ac1b10a98d4ebe9d72bed5c66ed0f2
Inputs	contoso-eval-en pinned v1; contoso-eval-en-dev12 version 1; all 12 cases
Pilot configuration issue	Relevance threshold 4; generic TaskAdherence threshold 4 was not meaningful for its binary output . Judge: lab-judge-luna-dea3cec5 / gpt-6-luna
Reported status	Completed; 12 output items
Pilot Agent usage	12 invocations , contoso-eval-en pinned v1 on gpt-4.1-mini
Pilot service-reported Judge usage	24 gpt-6-luna-2026-09-22 invocations; 83,492 tokens ; not corrected-comparison usage

Completed canonical baseline — portal; SDK candidate execution verified below:

Recorded item	Value
Corrected Evaluation name	contoso-en-learning-loop
Corrected evaluation ID	eval_94feef6f6f644fabb22a5680f5f24fb1
Corrected baseline run ID	evalrun_cde9948ac9d946929661bc3d9e60432a
Inputs	contoso-eval-en pinned v1; contoso-eval-en-dev12 version 1; the same 12 cases
Criteria / Judge	Relevance threshold 4; TaskAdherence binary threshold/pass 1; lab-judge-luna-dea3cec5 / gpt-6-luna
Confirmed response bindings	Relevance response={{sample.output_text}}; TaskAdherence response={{sample.output_items}}
Execution / overall	Completed; 12 rows; 10 passed / 2 failed / 0 errored
Per-evaluator results	Relevance 10/12 , threshold 4; TaskAdherence 12/12 , binary 1
Relevance failures	atlas-dev-001: vague subscription-reduction answer versus explicit clarification in the reference. atlas-dev-011: honest unverified-feature answer judged incomplete
Evidence boundary	HTTP 201 was the creation receipt; completion is now separately confirmed. These are baseline-only results; do not change policy facts to chase scores

Baseline Review used **Evaluation name → Submit**. Candidate **Add run** under the same definition did not work end-to-end: Pin v2 / Individual turns led to **Configure agents: Config required → Add custom prompt / User prompt**. Even `{{item.query}}` failed client-side with `Unable to create data source configuration from item schema`. **No candidate remote run was submitted by that portal attempt**. Do not create another definition as a workaround.

Optimizer succeeded with one candidate: `opt_e44bcf5701a348deb62a1cd4f9cb3910`, instruction-only, same dataset, gpt-5.5 generator, gpt-6-luna Judge, no model comparison. UI scores: **0.635 → 0.646, UI-reported +0.010, 264,260 reported tokens**. Preserve the displayed delta rather than deriving a different one from rounded display fields. This is not the direct candidate evaluation.

Portal Promote created lab-only `contoso-eval-en v2`. The operator verified exact candidate instructions and unchanged model/knowledge tool. **Latest active version affects all channels:** only an **isolated, unpublished lab agent** may use this promotion; **never production**. V1 remains the explicit recorded baseline. Do not repeat promotion or treat v2 activation as production approval.

In Optimizer **Criteria, No custom evaluators available** is a filter state: switch **Custom only OFF** or choose **View built-in evaluators**. Select a Relevance or TaskAdherence row to open **Configure...**, set **Relevance 4 / TaskAdherence 1**, then **Apply**. Do not add custom evaluators as a workaround.

Prerequisites for the single SDK step: Supply Python 3.11+ (3.12 recommended), Azure CLI, the operator-owned `scripts/add_foundry_eval_run.py`, and the current `requirements.lock`. From the repository root, use the approved **project endpoint** and subscription, not an endpoint guessed from a model deployment:

```
export AZURE_AI_PROJECT_ENDPOINT="OPERATOR_PROJECT_ENDPOINT"
export AZURE_SUBSCRIPTION_ID="OPERATOR_SUBSCRIPTION_ID"
```

Only if `.venv` is absent, prepare it once with `python3 -m venv .venv`. Activate/install the locked environment and verify the actual account:

```
source .venv/bin/activate
python -m pip install -r requirements.lock
az login
az account set --subscription "$AZURE_SUBSCRIPTION_ID"
az account show --query "{user:user.name,tenant:tenantId,subscription:id}" -o json
```

Check user, tenant, and subscription against the operator's approval and portal identity; do not proceed with another cached identity. Role assignment is not part of this check. Copy `FOUNDRY_EVALUATION_ID` and `FOUNDRY_BASELINE_RUN_ID` from the canonical baseline portal URL or **Raw JSON**, using the actual IDs above, never the pilot or Optimizer job ID.

The [single command in step 06](#) uses Azure AI Projects/OpenAI Evals SDKs to clone baseline `data_source`, retain existing criteria, change only target version to 2, and verify identical model/tools. It submits **one real Foundry evaluation run under the same evalID**, not a local Judge or bespoke cloud scoring. Its

`.lab/foundry-evaluations/candidate-v2.json` receipt is private; repeating the identical command collects the same run. Do not delete/change the receipt or create another definition to force a retry.

Verified end-to-end SDK execution: the helper actually submitted candidate run `evalrun_f3b710fc835d444fb8aa0d2bb7797bdf` under the **same** `eval_94feef6f6f644fabb22a5680f5f24fb1`, targeting `contoso-eval-en v2`. It is **Completed: 12 rows, 11 passed / 1 failed / 0 errored**, versus baseline **10 passed / 2 failed / 0 errored**. The helper confirmed identical dataset/evaluators/Judge/model/tools; only instructions/version changed.

Reuse the receipt and open `contoso-en-learning-loop` → **Evaluation runs** → **baseline/candidate checkboxes** → **Compare runs**. **Set the Baseline dropdown to the original** `contoso-eval-en`. It defaults to the first selected row, which was `candidate-v2` here; confirm direction before interpreting results. The portal item-schema failure remains real; the completed candidate came from the SDK.

Final observed comparison, n = 12	Baseline v1	Candidate v2
All-criteria passes; errors	10/12; 0	11/12; 0
Relevance passes; mean (1–5)	10/12; 4.4167	11/12; 4.3333
TaskAdherence passes; binary mean	12/12; 1.0	12/12; 1.0
Relevance rows 1 / 2 / 6 / 11	3 / 5 / 5 / 3	4 / 4 / 4 / 3
Latency p50 (ms)	5,891.09	7,287.52
Latency p95 (ms)	8,817.33	16,038.35
Agent tokens	35,187	43,751

Native **PairedTTest: Inconclusive** for **both Relevance and TaskAdherence**. Row 1 crossed the pass threshold, but rows 2/6 declined and row 11 stayed at 3 for honest uncertainty rather than an unsupported definitive answer. Do not change policy facts to chase scores. **HOLD adoption; keep pinned v1 pending further review/new representative cases**. Higher pass count does not offset the lower mean and observed latency/token tradeoffs or prove overall improvement. No production approval is granted.

Bound the actual workshop

Agree current authorization for the **actual account, project, data, processing locations, and costs**. Another rehearsal's approval is not transferable. `GlobalStandard` describes a processing arrangement, not a promise that every request is processed only in NCUS.

Work	Bound
Baseline	One correctly configured direct managed evaluation of all 12 cases; the retained pilot is not this baseline
Optimization	One instruction-only job, Max candidates 1 , wait at most 60 minutes
Candidate check	One SDK-submitted real Foundry run in the same evalID , same data/version/hash, Relevance 4 / TaskAdherence binary pass 1 / Judge; receipt-based collection, portal comparison
Agent changes	Lab v2 only; latest-active changes affect all channels. Promotion is restricted to the isolated unpublished agent, never production
Extra features	Unverified features are not run; no alternate exercise is used to claim completion

Setup checks have their own authorized usage; one Optimizer job can include many internal calls. An error is not permission to repeat jobs, widen scope, change regions, or rerun until a preferred result appears.

Verify identities and least-privilege access

Check the current [Foundry RBAC guidance \(https://learn.microsoft.com/azure/foundry/concepts/rbac-foundry\)](https://learn.microsoft.com/azure/foundry/concepts/rbac-foundry) against the operations actually needed. Role-assignment authority is separate from permission to run an agent.

Identity	Access to verify at the narrowest appropriate lab scope
Participant/operator	Open the project, read the pinned agent, upload/select a dataset, create/read evaluations, and use Agent Optimizer
Project managed identity	Required model and existing tool/connection access
Supporting service identities	Only the access required by the prepared agent's actual connections
Operator maintaining the lab	Authorized version creation, resource/cost inspection, and any separately approved lifecycle actions

Do not give participants subscription-wide Owner access as a shortcut, expose credentials, change organization policy, or bypass network restrictions. An account shown in a portal tab is not proof that a different client uses the same identity.

Cost and end-of-session ownership

Costs include agent responses, Judge evaluation, instruction generation, internal Optimizer calls, policy tools, and any continuously running supporting resources or logs. **Unknown cost is not zero**, quota is not a free allowance, and closing a browser does not stop hosting charges.

Name the operator responsible for remaining jobs, cost inspection, and authorized cleanup of workshop resources. Do not turn resource deletion or broad account changes into participant exercises. Keep credentials, private configuration, and raw account identifiers out of shared guides and packages.

Hand off a usable participant contract

Give participants	Required specificity
Portal target	The correct project and <code>contoso-eval-en</code> ; separately prepared Korean target <code>contoso-eval-ko</code>
Agent versions	Active v1 restored , recorded baseline v1 and retained candidate v2; <code>lab-agent-dea3cec5</code> , <code>gpt-4.1-mini / 2025-04-14</code> . Keep explicit versions and all evidence
Model roles	Agent <code>lab-agent-dea3cec5</code> , Judge <code>lab-judge-luna-dea3cec5</code> , Optimizer <code>lab-planner-dea3cec5</code> , with recorded versions and support limits
Dataset	Registered English <code>contoso-eval-en-dev12</code> version 1 , original file, n = 12 and SHA-256; Korean registration requires its own confirmation
Evaluation contract	Same corrected definition; Relevance 1–5 / threshold 4 ; TaskAdherence binary 0/1 / pass 1 , generic TaskAdherence Threshold 1 if shown; query-only input, no custom override
Bounds and ownership	Authorized scope, one candidate, 60-minute optimization wait, cost/cleanup owner, and escalation contact
SDK context	Verified venv/dependencies and <code>az login</code> account; operator project endpoint/subscription; baseline URL/Raw JSON IDs; helper and persistent receipt path
Evidence	Completed baseline/SDK candidate, matched inputs, final metrics/row regressions and PairedTTest Inconclusive; HOLD adoption pending further review/new representative cases. Pilot excluded

Record generated mappings, do not override them. The corrected submission confirmed **Relevance** response={{sample.output_text}} and **TaskAdherence** response={{sample.output_items}}. UI defaults also showed query={{item.query}} and TaskAdherence's tool_definitions={{sample.tool_definitions}}. Inspect **Raw JSON**; do not replace the submitted TaskAdherence binding with the earlier UI output-text value. These are sample bindings, not extra dataset columns.

Version limit: Catalog links showed relevance **v14** and task_adherence **v17**; actual service criteria had evaluator_version **empty/default**. Record names/configuration, Judge and dataset versions, and this limitation. Reusing one definition does not prove private service rubric versions are fully pinned.

Execution worked; adoption remains HOLD. Preserve both runs, receipt, and the measured tradeoffs. Keep pinned baseline v1 for selection; lab v2 activation remains a separate all-channel operator concern. The Optimizer's +0.010 and native pass-count increase are not proven overall improvement, statistical significance, production approval, a repaired portal path, or new Korean results.

Executed final lab state: contoso-eval-en active version was restored to **1** through **Details → Agent configuration → Active version → Edit → Version 1**. Candidate v2 and all evaluation/receipt evidence remain; do not delete them. **No production Publish was performed; no production channels/traffic are configured.** Foundry nevertheless automatically provides **RBAC-only Responses/preview endpoints without Publish**. Endpoint existence is not production publication.

Verification record and limitations · Foundry Evaluation

The **latest result JSON** (file: `evidence/latest.json`) separates documentation checks from the actual English rehearsal. The core workshop is **Foundry Evaluation → Agent Optimizer → same-criteria reevaluation**. It measures changes on the same Contoso business questions, rather than relying on a public benchmark score.

What is checked

Target	Evidence
Data	Separate English and Korean corpora; only each language's 12 dev cases are used in this loop
Actual evaluation	Foundry managed evaluation/run IDs, all 12 output items, metric scores, reasons, and errors
Optimizer	One actual instruction-only job, one candidate, and the original/candidate instruction diff
Reevaluation	Same evaluation definition, dataset version, evaluator settings, Judge, agent model, and tools; only instructions/version change
Portal	User authentication followed by Playwright Headless in the matching lab tenant, with English UI/data
Distribution	Both HTML/print/PDF editions, language switching, mobile layout, links, copying, and reading progress

A documentation **PASS** is not an Azure quality pass or operational approval. The Korean guide describes the same procedure, but the numbers below are **actual English-data results**, not a Korean-language execution.

Actual managed-evaluation results

Versions 1 and 2 of `contoso-eval-en` were evaluated in the newly provisioned dedicated North Central US environment.

Metric	Baseline v1	Candidate v2
Cases	12	12
Cases passing every selected evaluator	10/12	11/12
Relevance passes, threshold 4/5	10/12	11/12
Mean Relevance, 1–5 scale	4.42	4.33
TaskAdherence passes, binary Pass/Fail	12/12	12/12
Execution errors	0	0
Agent p50 latency	5.89 s	7.29 s
Agent p95 latency	8.82 s	16.04 s
Observed agent tokens	35,187	43,751

Adoption is on hold. One more case passed, but mean Relevance declined and agent latency/token usage increased. The portal's **Compare runs → PairedTTest** also reported **Inconclusive** for both metrics. One small synthetic dev comparison does not establish a meaningful overall improvement or generalization. The dedicated lab agent's active version was **restored to 1**; candidate 2 and its evidence remain available.

The actual evaluation definition is `eval_94feef6f6f644fab22a5680f5f24fb1`. Baseline run: `evalrun_cde9948ac9d946929661bc3d9e60432a`. Candidate run: `evalrun_f3b710fc835d444fb8aa0d2bb7797bdf`. Local custom-Judge scores were not substituted for these service results.

Optimizer job `opt_e44bcf5701a348deb62a1cd4f9cb3910` succeeded. Its displayed ranking moved from approximately **0.635 to 0.646**, with a displayed delta of **+0.010**. That is an internal 0–1 aggregate, not the same measure as the separate managed-evaluation pass rates or means above. The job reported **264,260 total tokens**; actual billing requires a separate check.

Model roles and issues found during execution

Role	Model used in this rehearsal	Evidence and limitation
Agent	<code>gpt-4.1-mini · 2025-04-14</code>	Retained the verified working configuration. <code>gpt-6-luna</code> Responses and agent probes returned HTTP 500 in this environment and were not labeled successful
Foundry Evaluation Judge	<code>gpt-6-luna · 2026-09-22</code>	Verified through Chat Completions and actual managed Foundry evaluations

Role	Model used in this rehearsal	Evidence and limitation
Optimizer generation model	gpt-5.5 · 2026-04-24	Retained because gpt-6-luna is not in the current supported optimization-model list (https://learn.microsoft.com/azure/foundry/agents/concepts/agent-optimizer-overview#models)

Catalog visibility is not runtime verification. The observed HTTP 500 errors do not establish a global lack of agent support. The published [retirement date](https://learn.microsoft.com/azure/foundry/openai/concepts/model-retirement-schedule) (<https://learn.microsoft.com/azure/foundry/openai/concepts/model-retirement-schedule>) for gpt-4.1-mini is **2027-04-14**; new subscriptions can face deprecation restrictions. Operators must verify actual model access before class.

The rehearsal identified and addressed these issues:

- **TaskAdherence is binary 0/1 Pass/Fail, not a 1–5 score.** The first generic-threshold UI check was retained as a configuration pilot. The actual comparison started with Relevance 4 and TaskAdherence 1. The pilot was not overwritten or relabeled as the comparison baseline.
- **Portal Add run item-schema error:** The UI returned Unable to create data source configuration from item schema. The supplied scripts/add_foundry_eval_run.py added the candidate to the **same Foundry evaluation definition and dataset** using the official SDK. This is actual managed evaluation, not local scoring.
- **Honest uncertainty can receive a low Relevance score:** atlas-dev-011 passed TaskAdherence but received Relevance 3 for declining to assert an unverified feature status. Do not invent policies or capabilities to increase a Judge score; inspect its reasoning against your own business criteria.
- **Separate resource-group policy failure:** The lab ARM deployment succeeded. A separate organization-policy deployment failed because its organization-managed log target was absent. Shared policy was not changed and the failure was not hidden.

Retention, deletion, and publication boundaries

The excluded training exercise's two dedicated jobs, trained model/checkpoint, model deployment, four uploaded files, and two result files were deleted and their absence verified. Its 202 local execution/log files were also removed. Shared models, the Foundry project, Search, and the current evaluation evidence remain. Deletion does not reverse charges already incurred.

Raw evaluation responses and the deletion audit remain private. Public deliverables exclude authentication state, cookies, tokens, approval files, and signed download URLs. No claim is made that Git history or Azure's general activity/billing records were erased.

Sources

Contoso Atlas Cloud policies, questions, and reference responses are synthetic—not actual customer data or provider terms. Earlier design references are preserved at [the archived repository's pinned commit](https://github.com/junwoojeong100/foundry-evaluation-labs-v0.9/tree/93bc07e31373c4cfc278a2dc3757785946404cf2) (<https://github.com/junwoojeong100/foundry-evaluation-labs-v0.9/tree/93bc07e31373c4cfc278a2dc3757785946404cf2>).

The English capture manifest (file: `web/assets/portal/en/captures.json`) records actual screen provenance, redactions, dimensions, and hashes. NOTICE (file: `web/assets/NOTICE.txt`) records usage boundaries for Microsoft artwork and portal captures. Only cropping and privacy masks are applied; scores, states, and data are not rewritten.

Official procedures: [Evaluate Foundry agents](https://learn.microsoft.com/azure/foundry/observability/how-to/evaluate-agent) (<https://learn.microsoft.com/azure/foundry/observability/how-to/evaluate-agent>) · [Agent Optimizer](https://learn.microsoft.com/azure/foundry/agents/quickstarts/quickstart-optimize-prompt-agent) (<https://learn.microsoft.com/azure/foundry/agents/quickstarts/quickstart-optimize-prompt-agent>) · [Evaluator-specific scales](https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/agent-evaluators) (<https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/agent-evaluators>).

Recheck locally

```
python scripts/build_datasets.py --language ko --check
python scripts/build_datasets.py --language en --check
python scripts/build_guide.py
python scripts/build_guide.py --check
python -m unittest discover -s tests -q
```

Commands explained: Verify both source corpora without modification, generate both HTML editions, check generated-file freshness, and run local tests. These commands do not start Azure evaluations. Refresh PDFs and the distribution ZIP separately, then record only checks actually performed in the latest JSON.

Synthetic English evaluation data

[Participant dataset step](#) · [Operator handoff](#) · Korean data guide (file: docs/ko/data-guide.html)

Use the workshop file unchanged

The English workshop uses **data/en/optimizer/dev.jsonl** (file: data/en/optimizer/dev.jsonl) for direct Microsoft Foundry managed Evaluation, Agent Optimizer, and separate direct reevaluation. It contains **exactly 12 JSONL records** with `query`, `context`, and `ground_truth`. Despite the directory name, it is the shared dataset for the entire evaluation loop.

`contoso-eval-en-dev12` **version 1 is registered through the native Evaluation wizard** from that original 12-row file. Use it with `contoso-eval-en` **version 1**; do not create a duplicate. Retain the dataset version and original file **SHA-256** across baseline, optimization, and reevaluation. Registration is not an evaluation result.

The Korean path uses operator-prepared `contoso-eval-ko` and `data/optimizer/dev.jsonl`, with separate `contoso-eval-ko-dev12` registration to confirm. **Korean was not re-executed in this rehearsal and has no new measured results here.** English responses must be English; Korean responses must be Korean. Switching a guide's language does not translate the selected dataset or agent. Cross-language runs are not paired evidence of instruction improvement.

The earlier `contoso-en-baseline-luna-judge` remains a **CONFIGURATION PILOT**. Canonical `contoso-en-learning-loop` is **Completed: 12 rows, 10 passed / 2 failed / 0 errored; Relevance 10/12, TaskAdherence 12/12 (binary 1)**, on the same pinned agent/data v1 and `gpt-6-luna` Judge. Actual IDs are in [step 03](#). These are baseline-only results, not candidate improvement.

Instruction-only job `opt_e44bcf5701a348deb62a1cd4f9cb3910` succeeded with **one candidate**, same dataset, `gpt-5.5` generation / `gpt-6-luna` Judge, no model comparison. UI **0.635 → 0.646, UI-reported +0.010, 264,260 reported tokens** are Optimizer results, not direct candidate scores. Promote created lab-only v2 with exact candidate instructions and unchanged model/knowledge tool. **Latest active version affects all channels: only isolated unpublished lab agents, never production.**

Purpose and limits

Contoso Atlas Cloud is **wholly fictional**. Its eight policies, questions, reference answers, and workspace examples are synthetic. They contain no actual private customer conversations, accounts, credentials, or production approval and do not represent a real provider's terms.

The method is **evaluate → learn from failure → improve → reevaluate**. For an application, representative company-owned tasks, policy edge cases, and known failures are more useful than a public benchmark score alone. This corpus demonstrates how to inspect those decisions without using a company's private data.

Synthetic authorship is not completed human review. Counts and hashes measure data integrity, not model quality. Twelve development cases cannot certify production behavior, generalization, policy compliance, or safety.

Source composition and workshop boundary

The source data/en/cases.jsonl (file: data/en/cases.jsonl) contains **100 cases** with fixed source labels:

Source subset	Rows	Use in this workshop
train	56	Unused/reserved source subset
validation	12	Unused/reserved source subset
dev	12	The only selected cases; exported to data/en/optimizer/dev.jsonl
test	20	Unused/reserved source subset

Do not upload the 100-case source in place of dev12, merge subsets, select only easy questions, or duplicate rows. Preserve the original source assets; no other subset is a participant exercise.

The English policy documents (file: data/en/knowledge/documents.json) supply stable ATLAS-* IDs and original policy excerpts. The prepared agent's policy access is an operator prerequisite, not a separate dataset or knowledge-construction lab.

Upload and response contracts

Each line is one JSON object with exactly these three columns:

Column	Type	Boundary
query	String	Only this field is sent to the agent
context	String	Source-policy excerpts for supported evaluators and reviewer reference; not appended to generation

Column	Type	Boundary
ground_truth	JSON string, not a nested object	Reference answer for supported evaluators and review; never a generation input

There is **no prefilled** response. Foundry runs the selected agent version to obtain an actual response. Do not invent response, rename columns, wrap records in a JSON array, or add an alternative job schema.

ground_truth decodes to an object with exactly four keys. The agent uses the same response contract:

Key	Expected value
answer	Nonempty English text; a short contract example is "Please confirm the purchase date."
citations	Supporting stable policy-ID array, such as ["ATLAS-REF-001"]; no fabricated IDs
route	answer for guidance, clarify for needed information, escalate for required human judgment, or refuse for prohibited requests
needs_human	Boolean; true only when route is escalate

These are schema examples, not actual generated answers. An **escalate** response does not mean a person has already been contacted. The agent has no business-execution tools and must not invent approvals, ticket IDs, refunds, deletions, or completed access changes.

Source-case IDs, route labels, and other metadata stay in source records; they are not extra generation inputs. Match evaluation rows by their actual question and stable row position/source case when reviewing results.

Policy excerpts are reference material, not proof of what an agent tool actually retrieved. Do not turn a quoted instruction in a question or document into a higher-priority command. Compare meaning and policy decisions, not exact answer-text equality.

Register once and select the same version

For a fresh agent still at v1, use **Foundry New experience → Build → Evaluations → Create → Create new evaluation → Agent**. Pin before changes, reselect the cleared checkbox and confirm **1 target before Next**; choose Individual turns / One time / Existing dataset. Recorded contoso-eval-en retains v2 but has restored active v1. Reuse the explicit saved v1 baseline/data rather than assuming "latest" means active.

Do not upload the registered English file again. Only in a fresh authorized project without that registration, use **Upload new dataset → name → Choose file** for `data/en/optimizer/dev.jsonl` → **Upload**. Confirm the selected dataset and preview. The preview shows **only the first five rows**; retain **all 12 records** and confirm the source file's SHA-256.

In the observed English wizard, the matching schema **auto-resolved Field mapping**, proceeding directly to **Configure agents**. Leave **custom prompt override unset**, keep agent input to query **only**, then continue to **Criteria**. If a mapping screen appears for the selected schema, use **query → query**.

Remove extras from the **23 suggestions**, retaining exactly **Relevance + TaskAdherence** with `lab-judge-luna-dea3cec5`. **Relevance is 1–5, threshold 4. TaskAdherence is Binary Pass/Fail, raw 0/1, pass 1**; if its generic control appears, set **Threshold 1**, then Update. The [official definitions](https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/agent-evaluators) (<https://learn.microsoft.com/azure/foundry/concepts/evaluation-evaluators/agent-evaluators>) confirm the binary type. The recorded **Evaluation name** `contoso-en-learning-loop` → **Submit** already succeeded with HTTP 201; open that run instead of submitting again.

Inspect, do not override, generated mappings. The corrected submission confirmed **Relevance** `response={{sample.output_text}}` and **TaskAdherence** `response={{sample.output_items}}`. UI defaults also showed `query={{item.query}}` and TaskAdherence's `tool_definitions={{sample.tool_definitions}}`. Read **Raw JSON**; do not restore TaskAdherence's earlier UI output-text value. The upload still has no prefilled response, and generation receives only query.

Version limit: Catalog links showed `relevance v14` and `task_adherence v17`, but actual service criteria had `evaluator_version empty/default`. Record evaluator names/configuration, Judge and dataset versions, and that private service rubric versions are not proven pinned.

Agent Optimizer uses **Next → Select dataset and criteria** and the **same registered dev12 version**, not Generate data or an edited upload. If Criteria says **No custom evaluators available**, switch **Custom only OFF** or click **View built-in evaluators**. Selecting a row opens **Configure...**: set **Relevance 4 / TaskAdherence 1**, then **Apply** for each. Keep binary pass 1 and the same Judge; no custom evaluator is required. See the [Optimizer step](#).

Compare without changing the evidence

Keep identical	Record from each actual run
All 12 questions, references, language, dataset version, file SHA-256	Same evaluation definition ID, distinct run IDs, explicit agent versions, final status
Same corrected definition; Relevance threshold 4; TaskAdherence binary pass 1; recorded configuration	Relevance 1–5 scores and TaskAdherence 0/1 results, per-metric pass counts/coverage/errors, and service-version limits

Keep identical	Record from each actual run
Judge deployment/model version, agent model, tools and connections	Actual answers and reasons, instruction diff, and per-case regressions

Read a failing/lowest-Relevance case and a good case when available. **n = 12** includes missing/errors.
Relevance's 1–5 score is not percent accuracy; TaskAdherence's 0/1 is Fail/Pass, not a five-point score.
A missing score is not 0, and Completed is not a quality pass.

For an actual answer, open the row's **conversation_id → User view**. That is a question/JSON-response view, not an inline Judge-reason panel. Return to **Detailed metrics result** for Relevance. reason / TaskAdherence. reason and compare the reference policy; never change policy facts just to raise a score.

Portal candidate submission is blocked: Add run → Pin v2 / Individual turns required Config required → Add custom prompt / User prompt; even `{{item.query}}` failed client-side with `Unable to create data source configuration from item schema`. No candidate remote run was submitted by that attempt. Do not alter the dataset, add a fake response, or create a new evaluation definition to work around it.

After [venv/requirements.lock](#) and `az login` [account verification](#), use the [single SDK command](#) in `scripts/add_foundry_eval_run.py`. The operator supplies endpoint/subscription; copy evaluation/baseline-run IDs from the baseline portal URL or Raw JSON. The Azure AI Projects/OpenAI Evals helper clones baseline `data_source/existing criteria`, changes only target version to **2**, verifies the same model/tools, and submits a **real Foundry run in the same evalID**, not local Judge scoring.

Verified native result: the helper submitted `evalrun_f3b710fc835d444fb8aa0d2bb7797bdf` for `contoso-eval-en v2` under the same `eval_94feef6f6f644fabb22a5680f5f24fb1`. It completed **12 rows: 11 passed / 1 failed / 0 errored**, versus baseline **10 passed / 2 failed / 0 errored**. Dataset/evaluators/Judge/model/tools matched; only instructions/version changed.

Keep the original receipt and collect only the existing run. Select both rows in `contoso-en-learning-loop` → **Evaluation runs → Compare runs**, then **choose the original** `contoso-eval-en` in the **Baseline dropdown**. It defaults to the first selected row, which was `candidate-v2` here; do not read before/after directions until the baseline is correct.

Final v1→v2 observations: **Relevance 10/12→11/12 passes but mean 4.4167→4.3333; TaskAdherence 12/12 both, binary mean 1.0**; zero errors in both runs. Relevance row 1 improved **3→4**, rows 2/6 declined **5→4**, and row 11 stayed **3** for honest uncertainty rather than an unsupported definitive answer. Keep policy facts unchanged.

Latency **p50 5,891.09→7,287.52 ms, p95 8,817.33→16,038.35 ms**, and Agent tokens **35,187→43,751** increased. Native **PairedTTest is Inconclusive for both metrics. HOLD adoption; keep pinned v1 pending further review/new representative cases**. Pass count alone or Optimizer +0.010 is not proven overall improvement.

This same-dev12 observation does **not establish statistical significance, independent generalization, or production approval**. The actual SDK run proves remote execution; manifests and local files alone do not. Keep the original data assets unchanged.

Final lab state: adoption HOLD; active version restored to **1**, with candidate v2 and evaluation/receipt evidence retained—not deleted. No production Publish or production channels/traffic were configured. Foundry still automatically exposes **RBAC-only Responses/preview endpoints without Publish**; these are not evidence of production deployment.